

A SEMI-SUPERVISED CLUSTERING APPROACH FOR GRAPH LEARNING WITH NEURAL NETWORKS

Chester Tan,* Anatol Eugen Wegner, Ingo Scholtes
Machine Learning for Complex Networks
Center for Artificial Intelligence and Data Science
Julius-Maximilians-Universität Würzburg
Würzburg, Germany

ABSTRACT

We propose a semi-supervised approach that combines any unsupervised clustering objective and supervised objective for end-to-end training any neural networks to improve node classification in attributed graphs, particularly when training labels are sparse. Our framework formulates node classification as semi-supervised inference of neural network models of attributed graphs with cluster structure. We use this framework to understand how neural networks for graph clustering can jointly cluster node attributes and graph structure, despite graph clustering objectives explicitly considering only graph structure and cluster assignments. Our framework also enables neural network architectures such as transformers and multilayer perceptrons to learn on graphs without positional encodings and without spectral or message passing layers found in graph neural networks. We evaluate our framework on six real-world attributed graph datasets.

1 INTRODUCTION

In graph learning, a common task is to identify groups of similar nodes, known as *classes* or *clusters* (Wu et al., 2022; Liu et al., 2023; Daneshfar et al., 2024). In a social network, for example, where connections between people are represented as edges in a graph, node classes or clusters can represent groups of people that are members of the same community (Bedi & Sharma, 2016). Graphs can also have attributes associated with each node, such as the age and dietary preferences of each person in a social network. Many different notions of communities in the same social network can be considered meaningful, and how the both attributes of individuals and their connections in the network give rise to these communities can be complex. Analogously, many different clusterings of the same attributed graph can be considered meaningful, and how both the attributes of nodes and their adjacencies give rise to these clusterings can be complex. Graph neural networks (GNNs) (Wu et al., 2022; Corso et al., 2024) are artificial neural networks designed to learn complex functions of node attributes and adjacencies. Where sufficient training labels are available, GNNs can be trained supervised and *end-to-end* for node classification (Kipf & Welling, 2016a). When no training labels are available, GNNs have recently been adapted for end-to-end unsupervised clustering (Shchur & Günnemann, 2019; Tsitsulin et al., 2023; Blöcker et al., 2024). However, unsupervised clustering may yield cluster assignments that do not align with the node labels of interest in real-world graphs. While it can be prohibitively expensive to obtain sufficient training labels for supervised node classification, it can be practical to label just a handful of examples for each label. In this work we propose a semi-supervised approach that combines unsupervised clustering objectives and supervised objectives for training neural networks to improve node classification in attributed graphs, particularly when training labels are sparse.

*correspondence to: chester.tan@uni-wuerzburg.de

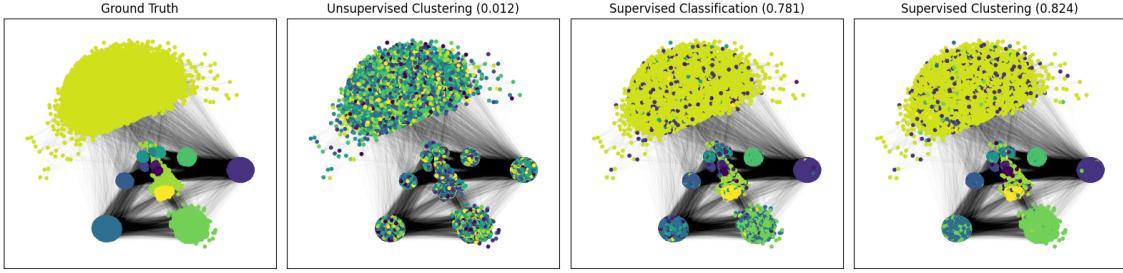


Figure 1: Visualization of node labellings comparing an unsupervised clustering, a semi-supervised node classification, and a semi-supervised clustering of the Coauthor CS dataset (Shchur et al., 2018), selected among 10 runs of random splits of 2 training nodes per cluster, 50 validation nodes, and 1000 test nodes. The unsupervised clustering labelling was selected from the run and configuration with the lowest training loss. The supervised classification and supervised clustering labellings were selected from the runs that yielded the highest validation set Matthews correlation coefficient (MCC). MCCs for each node labelling is displayed in parentheses.

1.1 RESEARCH GAPS AND CONTRIBUTIONS

We summarize the following research gaps we identify and our contributions to each.

Research Gap 1: While self-supervised deep graph representation learning methods (Veličković et al., 2019; Hassani & Khasahmadi, 2020; Zhu et al., 2020; Devvrit et al., 2022; Thakoor et al., 2022) have enabled deep graph clustering (Wang et al., 2023; Liu et al., 2023) with a two-step approach, Tsitsulin et al. (2023) recently identified a research gap in graph neural networks for *end-to-end* unsupervised graph clustering. They proposed a method to bridge the gap between traditional graph clustering objectives and graph neural networks, upon which other methods (Blöcker et al., 2024; Hansen & Bianchi, 2023) have been proposed. We identify a similar research gap, also identified by a recent review (Daneshfar et al., 2024), in graph neural networks for *end-to-end semi-supervised* graph clustering.

Contribution 1: We propose a reformulation of semi-supervised node classification on attributed graphs in the very sparse label setting as semi-supervised attributed graph clustering. Our experiments show that training with semi-supervised objectives consistently outperform a purely supervised objective. Figure 1 visualizes our approach on a real-world attributed graph dataset.

Research Gap 2: In unsupervised clustering on real-world attributed graphs, the loss landscape of clusters can be complex with many candidate solutions that can be considered meaningful (Peixoto, 2021; Blöcker & Scholtes, 2024).

Contribution 2: We propose a semi-supervised framework for evaluating existing end-to-end unsupervised deep graph clustering methods such as NOCD (Shchur & Günnemann, 2019), DMoN (Tsitsulin et al., 2023), Neuromap (Blöcker et al., 2024), and contribute an initial evaluation of these methods on six real-world attributed graph datasets. We additionally contribute a fully neural stochastic block model which we evaluate against.

Research Gap 3: We identify a contradiction where attributed graph clustering methods such as NOCD (Shchur & Günnemann, 2019), DMoN (Tsitsulin et al., 2023), Neuromap (Blöcker et al., 2024) optimize objectives that are functions of cluster assignments and graph structure, but not of node attributes. Constructing Bayesian network representations of their conditional dependencies, we identify a further contradiction that the explicit or *implicit* generative models that these methods infer are only generative models of graph struc-

ture and not of node attributes. On the other hand, we observe that MLPs have been shown effectively optimize unsupervised graph clustering objectives end-to-end (Shchur & Günnemann, 2019; Blöcker et al., 2024).

Contribution 3: We propose an alternative answer to the question posed in Shchur & Günnemann (2019): “Do we really need a graph neural network?” by proposing two simple ways existing methods can be adapted to learn explicit or *implicit* generative models both of graph structure and node attributes. If graph neural networks are used to optimize graph clustering objectives, an additional neural network module can be integrated to generate node attributes from cluster assignments. Alternatively, optimizing graph clustering objectives with neural networks of *only* node attributes learns a generative model akin to the *neural-prior stochastic block model* (Duranton & Ždeborová, 2023b), a recently proposed generative model of graphs with cluster structure from node attributes.

Research Gap 4: Adapting other neural networks such as multilayer perceptrons and transformers (Vaswani et al., 2017) typically involve incorporating a combination of the spectral and/or message passing layers found in graph neural networks (Chen et al., 2021; Rampasek et al., 2022; Qarkaxhija et al., 2024; Buterez et al., 2024) and/or additional encodings informative of graph structure (Ying et al., 2021; Kim et al., 2022; Ma et al., 2023; Huang et al., 2024).

Contribution 4: We propose an alternative approach for neural networks such as transformers and multi-layer perceptrons to learn on attributed graphs without positional or structural encodings and without spectral or message passing layers found in graph neural networks. We evaluate this ability of the framework on real-world attributed graph datasets and demonstrate that it improves learning on architectures such as transformers and MLPs.

2 SEMI-SUPERVISED DEEP GRAPH CLUSTERING END-TO-END

Setting: Consider a node-attributed graph $G := (V, E)$ of n nodes with node set V and edge set E . Nodes have attributes captured in a node attribute matrix \mathbf{X} and edges (equivalently, node adjacencies) are represented by an adjacency matrix \mathbf{A} . Each node is additionally assigned to one or multiple clusters from a set S of cardinality d_S of which $(d_S)_{\text{obs}} \leq d_S$ are observed. The node cluster assignments are represented with a label matrix \mathbf{S} of shape $n \times d_S$. Nodes and clusters are indexed with zero-based integers. Without loss of generality it is assumed in this work that the labels of the first s nodes belonging to the first $(d_S)_{\text{obs}}$ clusters are observed, which can be used for training. \mathbf{S} is defined in generality to allow for various forms of cluster assignments, such as hard or soft assignments (Yu et al., 2005), and overlapping or non-overlapping clusters (Shchur & Günnemann, 2019). The primary learning task of interest is to infer the node label matrix \mathbf{S} . We primarily consider the setting of sparse training labels, i.e. $s \ll n$, though our proposed framework proposed does not rely on this assumption.

The formulations in this section primarily consider unweighted and undirected graphs without edge and graph attributes for both brevity and to match the setting of the experiments on real-world attributed graph datasets that are also unweighted and undirected. In appendix B.1 we discuss extensions to weighted and directed graphs with edge and graph attributes.

2.1 GRAPH NEURAL NETWORKS WITHOUT *Attributed* GRAPH GENERATIVE MODELS

Our first contribution is to identify the generative models that graph neural networks for semi-supervised node classification and unsupervised graph clustering infer. By constructing their Bayesian network representations, we show how these graph neural networks should not learn generative models of (node-)attributed graphs.

2.1.1 EXPLICIT AND *Implicit* GENERATIVE MODELS

Graph clustering functions such as NOCD (Shchur & Günnemann, 2019), DMoN (Tsitsulin et al., 2023), Neuromap (Blöcker et al., 2024) and graph pooling functions such as MinCutPool (Bianchi et al., 2020) and DiffPool (Ying et al., 2018) optimize objectives that can generally be expressed

$$l_{\text{clustering}} \approx c(\hat{\mathbf{S}}, \mathbf{A}) \quad (1)$$

where $l_{\text{clustering}}$ is the resulting training loss and c is the clustering objective function. For the clustering objective $c(\hat{\mathbf{S}}, \mathbf{A})$ of NOCD, DMoN, Neuromap, and similar methods, we propose that there exists a function $\hat{\mathbf{A}} \sim f(\hat{\mathbf{S}})$ that enables the existence a graph reconstruction objective $L_E(\hat{\mathbf{A}}, \mathbf{A})$ such that

$$\arg \min_{\hat{\mathbf{A}}} (L_E(f(\hat{\mathbf{S}}), \mathbf{A})) \approx \arg \min_{\hat{\mathbf{S}}} (c(\hat{\mathbf{S}}, \mathbf{A})) \quad (2)$$

We refer to this function $\hat{\mathbf{A}} \sim f(\hat{\mathbf{S}})$ as the explicit or *implicit* generative model of the graph structure of the clustering objective. For DiffPool and NOCD, the generative model is explicit (see background section A). For DMoN, Neuromap, and similar methods, this generative model is *implicit*. The implicit generative models of modularity maximization and Infomap optimizing the map equation (Smiljanić et al., 2023) – the discrete combinatorial optimization objectives relaxed by DMoN and Neuromap, respectively – and their equivalence to stochastic block models have been shown by Peixoto & Kirkley (2022). We hypothesise that this can be extended to the graph clustering functions of DMoN, Neuromap, and similar methods, and empirically evaluate this in our experiments.

2.1.2 BAYESIAN NETWORKS REVEALING GENERATIVE MODELS

Using this interpretation of generative models of graph clustering functions, we construct and analyse Bayesian network representations of graph neural networks for (semi-)supervised node classification and unsupervised graph clustering in figure 2. This reveals the generative models inferred in each case:

- Graph neural networks $\text{GNN}_{(\mathbf{X}, \mathbf{A}) \rightarrow \hat{\mathbf{S}}}$ for (semi-)supervised node classification learn generative models of labels *from* node attributes and graph structure.
- Graph neural networks for unsupervised graph clustering GAE_o use $\text{GNN}_{(\mathbf{X}, \mathbf{A}) \rightarrow \hat{\mathbf{S}}}$ encoders to infer generative models $\hat{\mathbf{A}} \sim f(\hat{\mathbf{S}})$ of graph structure from cluster assignments, but *not* of node attributes.

This also reveals an equivalence between graph neural networks for unsupervised graph clustering and graph autoencoders (Kipf & Welling, 2016b).

2.2 CONDITIONAL DEPENDENCIES OF ATTRIBUTED GRAPH GENERATIVE MODELS WITH CLUSTER STRUCTURE

To construct models that learn generative models of *attributed* graphs with cluster structure, we make the contribution of identifying the following requirements of node-attributed generative models of graphs with cluster structure and their representations as Bayesian networks:

- $\hat{\mathbf{A}}$ depends on $\hat{\mathbf{S}}$, represented by directed edge from $\hat{\mathbf{S}}$ to $\hat{\mathbf{A}}$ in a Bayesian network.
- Either $\hat{\mathbf{X}}$ depends on $\hat{\mathbf{S}}$ or $\hat{\mathbf{S}}$ depends on $\hat{\mathbf{X}}$, represented by the existence of a directed edge between them in a Bayesian network.

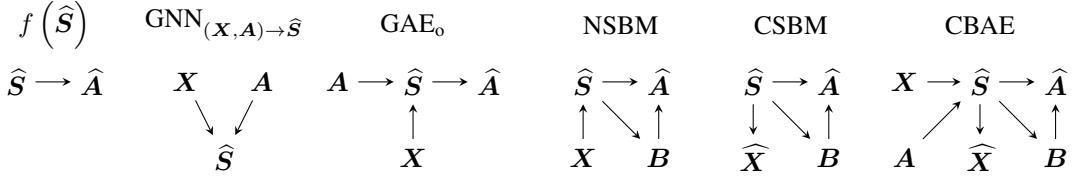


Figure 2: Bayesian network representations of a graph neural network for (semi-)supervised node classification ($\text{GNN}_{(\mathbf{X}, \mathbf{A}) \rightarrow \hat{\mathbf{S}}}$), a class of graph neural network autoencoders for unsupervised graph clustering (GAE_o) that only optimize graph reconstruction objectives or unsupervised clustering objectives, and our proposed NSBM, CSBM, and CBAE architectures.

- Conditional independence between \mathbf{X} (or $\widehat{\mathbf{X}}$) and $\widehat{\mathbf{A}}$ given $\widehat{\mathbf{S}}$.

Requiring conditional independence between \mathbf{X} (or $\widehat{\mathbf{X}}$) and $\widehat{\mathbf{A}}$ given $\widehat{\mathbf{S}}$ ensures that the cluster structure of the attributed graph contains all the necessary information to infer both node attributes and graph structure.

We identify two generative models of attributed graphs with cluster structure that satisfy the conditions proposed above – the neural-prior stochastic block model (Duranton & Zdeborová, 2023b) and the contextual stochastic block model (Deshpande et al., 2018) – as evidenced by their Bayesian network representations in figure 3 in appendix A. In the following sections, we propose fully neural network variants of these models, the NSBM and the CSBM, respectively.

2.3 TRANSFORMERS AND MLPs LEARN GRAPHS WITH NEURAL(-PRIOR) SBMs

The first fully end-to-end neural network generative model of attributed graphs with cluster structure that we propose is the NSBM, constructed as follows:

$$\widehat{\mathbf{A}} \sim \text{NSBM}(\mathbf{X}) = \text{SBM}_f(\widehat{\mathbf{S}}) \quad (3)$$

where

$$\widehat{\mathbf{S}} := \text{NN}_{\mathbf{X} \rightarrow \widehat{\mathbf{S}}}(\mathbf{X})$$

As we constructed in section 2.1.1, inspired by Peixoto & Kirkley (2022), we propose that $\text{SBM}_f = f$ to be any explicit or *implicit* stochastic block model of $\widehat{\mathbf{S}}$, such that $L_E(\widehat{\mathbf{A}}, \mathbf{A})$ can represent the graph clustering functions of DMoN (Ying et al., 2018), MinCutPool (Bianchi et al., 2020), DMoN (Tsitsulin et al., 2023), Neuromap (Blöcker et al., 2024), and similar approaches.

We also contribute a construction of a fully neural stochastic block model SBM_{NN} , constructed as follows:

$$\widehat{\mathbf{A}} \sim \text{SBM}_{\text{NN}}(\widehat{\mathbf{S}}) = \text{SBM}_B(\widehat{\mathbf{S}}, \mathbf{B}) \quad (4)$$

where SBM_B is a Bernoulli SBM (see equation 12, appendix A) and

$$\begin{aligned} \mathbf{B} &:= \phi(\mathbf{Z}^T \mathbf{Z}) \\ \mathbf{Z} &:= \text{NN}_{\widehat{\mathbf{S}} \rightarrow (\mathbf{Z})}(\widehat{\mathbf{S}}) \end{aligned}$$

ϕ is any function that ensures that the values of \mathbf{B} are valid probabilities. Its Bayesian network representation is shown in figure 2. Learning a block matrix \mathbf{B} can enable more diverse cluster structures (Peixoto, 2019; Pei et al., 2019; Scholkemper & Schaub, 2023) to align with observed node labels. The neural networks denoted NN can be any neural network. In particular, $\text{NN}_{\mathbf{X} \rightarrow (\hat{\mathbf{S}}, \mathbf{B})}$ in the NSBM can be any neural network function of node attributes. We focus on the cases where $\text{NN}_{\mathbf{X} \rightarrow (\hat{\mathbf{S}}, \mathbf{B})}$ is a transformer (Vaswani et al., 2017) or a multilayer perceptron (MLP).

2.3.1 INFERRING ATTRIBUTED STOCHASTIC BLOCK MODELS BY TRAINING NEURAL NETWORKS END-TO-END

Using the NSBM, we contribute a framework for neural networks such as transformers and MLPs to learn on attributed graphs without positional or structural encodings (Ying et al., 2021; Kim et al., 2022; Ma et al., 2023; Huang et al., 2024) and without spectral or message passing layers found in graph neural networks (Chen et al., 2021; Rampasek et al., 2022; Qarkaxhija et al., 2024; Buterez et al., 2024).

Where $\text{NN}_{\mathbf{X} \rightarrow \hat{\mathbf{S}}}(\mathbf{X})$ in equation 3 is a transformer or MLP, the NSBM can be trained end-to-end by optimizing reconstruction losses of its predictions of \mathbf{S} and \mathbf{A} from their observed values.

$$l_{\text{NSBM}} := L_S(\hat{\mathbf{S}}_{0:s,:(d_S)_{\text{obs}}}, \mathbf{S}_{0:s,:(d_S)_{\text{obs}}}) + L_E(\hat{\mathbf{A}}, \mathbf{A}) + L_{\text{regularization}} \quad (5)$$

$L_{\text{regularization}}$ is an optional regularization term such as in DiffPool, MinCutPool, and DMoN (see appendix A). As we constructed in section 2.1.1 $L_E(\hat{\mathbf{A}}, \mathbf{A})$ can be replaced by the clustering objectives of DMoN, Neuromap, and similar methods. As $\hat{\mathbf{A}}$ is reconstructed from $\hat{\mathbf{S}}$, we can choose $\hat{\mathbf{S}}$ to be of size $n \times (d_S)_{\text{max}}$ where $(d_S)_{\text{max}} > (d_S)_{\text{obs}}$ particularly when $(d_S)_{\text{obs}}$ is small to enable a less lossy reconstruction of \mathbf{A} , and to allow for unseen clusters to be inferred (see appendix B.2.1 and A).

Using a SBM_f function for the NSBM such as such as DiffPool, NOCD, MinCutPool, DMoN, Neuromap, or SBM_{NN} , a transformer or MLP-based NSBM makes *no changes* to the architecture of a transformer or MLP, respectively, apart from its training loss, while enabling it to learn on attributed graphs without positional or structural encodings, spectral or message passing layers.

2.4 GRAPH NEURAL NETWORKS WITH ATTRIBUTE RECONSTRUCTION LEARN CONTEXTUAL SBMs

The second fully end-to-end neural network generative model of attributed graphs with cluster structure that we propose is the CSBM, constructed as follows:

$$\hat{\mathbf{X}}, \hat{\mathbf{A}} \sim \text{CSBM}(\hat{\mathbf{S}}) := \text{NN}_{\hat{\mathbf{S}} \rightarrow \hat{\mathbf{X}}}(\hat{\mathbf{S}}), \text{SBM}_f(\hat{\mathbf{S}}) \quad (6)$$

Unlike NSBM, CSBM has $\hat{\mathbf{S}}$ as an independent variable and hence requires a graph neural network encoder. We therefore construct the following CBAE autoencoder which infers the CSBM.

$$\hat{\mathbf{X}}, \hat{\mathbf{A}} \sim \text{CBAE} := \text{CSBM}\left(\text{GNN}_{(\mathbf{X}, \mathbf{A}) \rightarrow \hat{\mathbf{S}}}(\mathbf{X}, \mathbf{A})\right) \quad (7)$$

Its Bayesian network representation is shown in figure 2.

2.4.1 RECONSTRUCTING ATTRIBUTES LEARNS GENERATIVE MODELS OF ATTRIBUTED GRAPHS

Like the NSBM, the CBAE is also trained end-to-end by optimizing the reconstruction losses of its predictions.

$$l_{\text{CBAE}} := L_E(\hat{\mathbf{A}}, \mathbf{A}) + L_V(\hat{\mathbf{X}}, \mathbf{X}) + L_{\text{regularization}} \quad (8)$$

Depending on the attributes observed, the attribute reconstruction objective L_V can a combination of regression objectives such as mean squared error or mean absolute error, and classification objectives such as cross-entropy. As we identified in section 2.1.1, without the attribute reconstruction we propose for the CBAE, graph neural networks for unsupervised graph clustering should learn generative models only of graph structure. This highlights the importance of attribute reconstruction for learning on graphs with attributes with graph neural networks.

2.5 RECONSTRUCTING SEMI-SUPERVISED NODE CLASSIFICATION

By reinterpreting the reconstruction objectives of the NSBM (equation 5) and CBAE (equation 8) as semi-supervised clustering objectives, we contribute a framework where any neural network or graph neural network can utilize any unsupervised clustering objective, reconstruction objective, or self-supervised objective to learn on attributed graphs with cluster structure in a semi-supervised manner.

$$l_{\text{NSBM}} := L_{\text{supervised}} \left(\hat{\mathbf{S}}_{0:s,:(d_S)_{\text{obs}}}, \mathbf{S}_{0:s,:(d_S)_{\text{obs}}} \right) + (L_{\text{unsupervised}})_{\text{NSBM}} \quad (9)$$

$$l_{\text{CBAE}} := L_{\text{supervised}} \left(\hat{\mathbf{S}}_{0:s,:(d_S)_{\text{obs}}}, \mathbf{S}_{0:s,:(d_S)_{\text{obs}}} \right) + (L_{\text{unsupervised}})_{\text{CBAE}} \quad (10)$$

where $L_S = L_{\text{supervised}}$ is a supervised classification objective such as cross-entropy (see appendix A), and

$$(L_{\text{unsupervised}})_{\text{NSBM}} = L_E + L_{\text{regularization}}$$

$$(L_{\text{unsupervised}})_{\text{CBAE}} = L_E + L_V + L_{\text{regularization}}$$

are unsupervised objectives. As we constructed in section 2.1.1, L_E can be replaced with any unsupervised clustering objective such NOCD, DMoN, Neuromap, and similar methods, and $L_{\text{regularization}}$ is an optional regularization term such as in DiffPool, MinCutPool, and DMoN (see appendix A and E.4.1).

While we formulate semi-supervised node classification and semi-supervised graph clustering equivalently in setting of our proposed methodology, we propose two distinctions between the two: by the methods used to solve the problem, and by the sparsity of the labels. For semi-supervised node classification training losses are purely supervised while for semi-supervised graph clustering training loss combine supervised and unsupervised objectives. In our experiments (section 3), we explore the two label sparsity settings, one with labels designed to be sufficient for supervised training, and one with labels designed to be sparse for semi-supervised training where semi-supervised graph clustering objectives should outperform purely supervised objectives. By constructing this semi-supervised framework we contribute (1) a way for any unsupervised clustering objective, reconstruction objective, or self-supervised objective to be used alongside any supervised objective to improve the performance of any neural network or graph neural network on node classification, particularly when labels are sparse, and (2) a way to evaluate the performance of graph clustering objectives on real-world attributed graph datasets in the setting of semi-supervised node classification.

3 EXPERIMENTS

In our experiments we contribute (1) an evaluation of our proposed NSBM and CBAE in the setting of semi-supervised attributed graph clustering to improve the performance of any neural network or graph neural network on multi-class node classification. (2) a framework for evaluating the performance of attributed graph clustering objectives in the setting of semi-supervised node classification and an initial benchmark of existing attributed graph clustering methods adapted to end-to-end semi-supervised attributed graph clustering.

Our experiments focus on adapting multi-class classification to semi-supervised graph clustering on 6 real-world attributed graph datasets – Coauthor CS, Coauthor Physics, Amazon Computers, and Amazon Photo

from Shchur et al. (2018), and `roman-empire` and `amazon-ratings` from Platonov et al. (2023b). The cardinality of the set of labels d_S is known and a non-zero number of training nodes per label is given.

To adapt these datasets to a semi-supervised clustering, we sparsify the labels by randomly selecting a subset of labels to be observed and the rest to be unobserved. Specifically, for each prescribed training, validation, and test split for each dataset, we randomly select 2 training nodes per label from the training split, and 50 validation nodes from the validation split. The test split remains the same as the original split.

Following the proposed approach in equation 9, semi-supervised clustering objectives combine the supervised cross-entropy objective L_S in equation 16 with the unsupervised objectives L_E , $L_{\text{regularization}}$, and L_V in equation 17. We evaluate the performance of **4** unsupervised graph clustering objectives: NOCD (Shchur & Günnemann, 2019), DMoN (Tsitsulin et al., 2023), Neuromap (Blöcker et al., 2024), and our proposed SBM_{NN} , and **2** unsupervised attribute reconstruction objectives: DMoN (Tsitsulin et al., 2023) and (unit) L2 regularization. Shchur & Günnemann (2019) showed that L2 regularization was beneficial, and we evaluate the case of unit L2 regularization which corresponds to an exact variational and description length objective (Graves, 2011), motivated by the interpretation of these model as generative models. L_V is mean squared error loss. The performance of semi-supervised objectives over purely supervised objectives is evaluated on **4** neural network architectures, two graph neural networks – GCN (Kipf & Welling, 2016a) and GraphSAGE (Hamilton et al., 2017) – and two “non-graph” neural networks – MLP and Transformer.

Our main results are summarized in table 1 for label sparsified versions of the datasets, and in appendix F.2 for the default splits of the datasets. Semi-supervised objectives consistently outperform a purely supervised objective on label sparsified versions of the datasets and the performance of semi-supervised objectives over purely supervised objectives is significantly more pronounced when labels are sparse. Figure 1 and the additional visualizations in appendix F.1 support these findings. The best performing graph neural network on 3 of 6 datasets use the attribute reconstruction in our CBAE approach, also significantly more pronounced when labels are sparse (see appendix F.3 for comparing attribute reconstruction on default splits of the datasets).

Though graph neural networks such as GCN and GraphSAGE consistently produce the best performance on these datasets, transformers and MLPs trained with semi-supervised objectives under our NSBM framework both outperform a purely supervised objective on most datasets and perform competitively with graph neural networks on some datasets such as the Coauthor CS dataset. MLPs perform consistently better than transformers on these datasets. Results from Tsitsulin et al. (2023) and Blöcker et al. (2024) support these findings for unsupervised graph clustering showing that on datasets where k-means node clustering (only) on features outperforms SBM inference or Infomap (only) on the graph structure, MLPs outperformed graph neural networks. It has also been empirically observed and studied theoretically how graph neural networks use graphs when detrimentally in cases where graph structure should be ignored in favour of attributes (Bechler-Speicher et al., 2024).

Across architectures, the unsupervised graph clustering and regularization objectives of DMoN produced the best performance on 4 of 6 datasets, while transformers consistently benefited most from unit L2 regularization. Additional results in appendix F.4 echo that the SBM_{NN} function we designed and NOCD did not perform as well as implicit objectives such as DMoN. While we highlight the importance of designing architectures by uncovering and analysing their conditional dependencies, implicit models such as DMoN could provide a more expressive for learning on graphs. Notably we show that unsupervised graph clustering objectives with regularization outperform regularization alone, ablating the effect of regularization. Despite the small number of validation nodes in the sparse label setting, early stopping and model selection on the validation set still produces the best performance on the test set.

In the appendix, we include additional experimental details and results and discuss possible extensions of our framework to settings such as an unknown number of labels as well as unbalanced and unseen labels.

Table 1: Summary results comparing semi-supervised graph clustering (where L_E is not “None”) compared with semi-supervised node classification (where L_E is “None”) for 4 neural network architectures – Transformers, MultiLayer Perceptrons (MLPs), and Graph neural networks GCN and GraphSAGE – evaluated on 6 label sparsified real-world attributed graph datasets. The better result for each architecture is highlighted in **bold** and the best result across all architectures is underlined. The L_E and $L_{\text{regularization}}$ columns indicate which unsupervised clustering and regularization losses where used in addition to a supervised cross-entropy loss for training. A \times or \checkmark in the L_V column indicates if attributes were reconstructed, and a \times or \checkmark in the L_S column indicates if a supervised cross-entropy loss was used for training. The “ES” column indicates if training loss or validation MCC was used as an early stopping criterion.

Model	f	$L_{\text{regularization}}$	L_V	$NN_{S \rightarrow X}$	ES	MCC	Accuracy
Amazon Computers							
GCN	None	None	\times	None	Loss	0.502 ± 0.096	0.569 ± 0.101
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.507 ± 0.065	0.574 ± 0.065
GraphSAGE	None	DMoN	\times	None	Loss	0.466 ± 0.105	0.522 ± 0.114
GraphSAGE	DMoN	None	\times	None	MCC	0.484 ± 0.110	0.550 ± 0.121
MLP	None	DMoN	N/A	None	MCC	0.222 ± 0.034	0.297 ± 0.041
MLP	DMoN	DMoN	N/A	None	MCC	0.207 ± 0.058	0.297 ± 0.066
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.075 ± 0.051
Transformer	NOCD	L2	N/A	None	Loss	0.106 ± 0.093	0.177 ± 0.105
Amazon Photo							
GCN	None	None	\times	None	MCC	0.620 ± 0.078	0.653 ± 0.087
GCN	DMoN	None	\times	None	MCC	0.618 ± 0.093	0.659 ± 0.098
GraphSAGE	None	DMoN	\times	None	MCC	0.626 ± 0.089	0.671 ± 0.086
GraphSAGE	DMoN	DMoN	\times	None	MCC	0.571 ± 0.102	0.618 ± 0.105
MLP	None	DMoN	N/A	None	MCC	0.399 ± 0.052	0.474 ± 0.044
MLP	DMoN	None	N/A	None	MCC	0.387 ± 0.019	0.461 ± 0.016
Transformer	None	L2	N/A	None	MCC	0.016 ± 0.024	0.130 ± 0.079
Transformer	NOCD	L2	N/A	None	Loss	0.155 ± 0.138	0.232 ± 0.137
Coauthor CS							
GCN	None	DMoN	\times	None	MCC	0.748 ± 0.052	0.770 ± 0.050
GCN	DMoN	None	\times	None	MCC	0.756 ± 0.045	0.778 ± 0.042
GraphSAGE	None	DMoN	\times	None	MCC	0.758 ± 0.042	0.778 ± 0.041
GraphSAGE	DMoN	DMoN	\times	None	MCC	0.764 ± 0.037	0.784 ± 0.037
MLP	None	None	N/A	None	MCC	0.602 ± 0.061	0.639 ± 0.059
MLP	NOCD	L2	N/A	None	Loss	0.710 ± 0.046	0.738 ± 0.044
Transformer	None	L2	N/A	None	MCC	0.476 ± 0.120	0.516 ± 0.130
Transformer	NOCD	L2	N/A	None	Loss	0.677 ± 0.241	0.699 ± 0.245
Coauthor Physics							
GCN	None	None	\times	None	MCC	0.816 ± 0.046	0.870 ± 0.039
GCN	DMoN	None	\times	None	MCC	0.830 ± 0.042	0.882 ± 0.030
GraphSAGE	None	None	\times	None	MCC	0.795 ± 0.061	0.856 ± 0.048
GraphSAGE	DMoN	DMoN	\checkmark	Transformer	MCC	0.791 ± 0.040	0.853 ± 0.026
MLP	None	None	N/A	None	MCC	0.493 ± 0.120	0.613 ± 0.146

Continued on next page

Table 1: Summary results on label-sparsified real-world attributed graph datasets.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
MLP	NOCD	L2	N/A	None	Loss	0.622 ± 0.090	0.725 ± 0.064
Transformer	None	L2	N/A	None	MCC	0.490 ± 0.107	0.592 ± 0.134
Transformer	NOCD	L2	N/A	None	MCC	0.495 ± 0.265	0.586 ± 0.310
amazon-ratings							
GCN	None	L2	✗	None	MCC	-0.003 ± 0.030	0.252 ± 0.093
GCN	NOCD	None	✓	MLP	Loss	0.031 ± 0.010	0.344 ± 0.006
GraphSAGE	None	L2	✗	None	MCC	0.002 ± 0.004	0.032 ± 0.076
GraphSAGE	NOCD	L2	✗	None	MCC	0.001 ± 0.005	0.031 ± 0.085
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	NOCD	L2	N/A	None	MCC	0.000 ± 0.001	0.063 ± 0.134
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.321 ± 0.062
roman-empire							
GCN	None	DMoN	✗	None	Loss	0.090 ± 0.005	0.196 ± 0.006
GCN	DMoN	DMoN	✗	None	Loss	0.090 ± 0.006	0.195 ± 0.006
GraphSAGE	None	DMoN	✗	None	Loss	0.214 ± 0.184	0.301 ± 0.139
GraphSAGE	DMoN	DMoN	✓	MLP	Loss	0.327 ± 0.019	0.388 ± 0.014
MLP	None	DMoN	N/A	None	MCC	0.222 ± 0.034	0.240 ± 0.035
MLP	DMoN	DMoN	N/A	None	MCC	0.234 ± 0.043	0.261 ± 0.041
Transformer	None	L2	N/A	None	MCC	-0.000 ± 0.026	0.022 ± 0.050
Transformer	NOCD	DMoN	N/A	None	MCC	0.004 ± 0.012	0.013 ± 0.041

4 CONCLUSION

In this work we propose and evaluate a semi-supervised clustering approach for graph learning with neural networks. Our framework combines enables any unsupervised graph clustering objective to improve the performance of any neural network or graph neural network on multi-class node classification in the setting of semi-supervised graph clustering. The perspective we introduce helps us to understand how current graph neural network methods for graph clustering jointly cluster node attributes and graph structure, despite clustering objectives explicitly considering only node attributes and graph structure. Our framework also introduces an alternative graph learning approach, enabling such as neural networks such as transformers and multi-layer perceptrons to learn on graphs without positional or structural encodings and without spectral or message passing layers found in graph neural networks. The results of our experiments demonstrate that our approach is complementary to existing semi-supervised node classification approaches on six real-world node-attributed graph datasets, improving learning when training labels are sparse.

We propose that the adaptation evaluation of our approach for link prediction and other edge-level tasks, extending the approach for edge and graph-level learning tasks, and further investigating the impact of attribute reconstruction for graph neural network methods for graph clustering could be interesting avenues for future work. Theoretical analysis of the expressivity of the alternative graph learning approach for non-graph neural networks is also an open question which we hope to explore with the community. We discuss potential future work further in appendix B.2.

5 REPRODUCIBILITY STATEMENT

All source code and documentation to reproduce results is made available in the supplementary zip file. All dataset sources are detailed in the appendix, and the source code includes the necessary code to fetch the datasets. Experiments were run on a cluster with virtual machines each with 16 vCPUs and 128 GB RAM, a single NVIDIA L40 GPU with 48 GB VRAM, and 100 GB disk space. The presented experiments require approximately 1 year on a single virtual machine of the cluster to run, and the source code is designed to be run on a cluster with multiple virtual machines in parallel.

REFERENCES

- Jason Ansel, Edward Yang, Horace He, Natalia Gimelshein, Animesh Jain, Michael Voznesensky, Bin Bao, Peter Bell, David Berard, Evgeni Burovski, Geeta Chauhan, Anjali Chourdia, Will Constable, Alban Desmaison, Zachary DeVito, Elias Ellison, Will Feng, Jiong Gong, Michael Gschwind, Brian Hirsh, Sherlock Huang, Kshitij Kalambarkar, Laurent Kirsch, Michael Lazos, Mario Lezcano, Yanbo Liang, Jason Liang, Yinghai Lu, CK Luk, Bert Maher, Yunjie Pan, Christian Puhrsch, Matthias Reso, Mark Saroufim, Marcos Yukio Siraichi, Helen Suk, Michael Suo, Phil Tillet, Eikan Wang, Xiaodong Wang, William Wen, Shunting Zhang, Xu Zhao, Keren Zhou, Richard Zou, Ajit Mathews, Gregory Chanan, Peng Wu, and Soumith Chintala. PyTorch 2: Faster Machine Learning Through Dynamic Python Bytecode Transformation and Graph Compilation. In *29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2 (ASPLOS '24)*. ACM, April 2024. doi: 10.1145/3620665.3640366. URL <https://pytorch.org/assets/pytorch2-2.pdf>.
- Avanti Athreya, Donniell E Fishkind, Minh Tang, Carey E Priebe, Youngser Park, Joshua T Vogelstein, Keith Levin, Vince Lyzinski, Yichen Qin, and Daniel L Sussman. Statistical inference on random dot product graphs: a survey. *Journal of Machine Learning Research*, 18(226):1–92, 2018.
- Yingbin Bai, Erkun Yang, Bo Han, Yanhua Yang, Jiatong Li, Yidian Mao, Gang Niu, and Tongliang Liu. Understanding and improving early stopping for learning with noisy labels. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan (eds.), *Advances in Neural Information Processing Systems*, 2021. URL <https://openreview.net/forum?id=KbV-UZRKb3g>.
- Maya Bechler-Speicher, Ido Amos, Ran Gilad-Bachrach, and Amir Globerson. Graph neural networks use graphs when they shouldn't. In *Forty-first International Conference on Machine Learning*, 2024. URL <https://openreview.net/forum?id=fSNHK7mu3j>.
- Punam Bedi and Chhavi Sharma. Community detection in social networks. *Wiley interdisciplinary reviews: Data mining and knowledge discovery*, 6(3):115–135, 2016.
- Filippo Maria Bianchi, Daniele Grattarola, and Cesare Alippi. Spectral clustering with graph neural networks for graph pooling. In *International conference on machine learning*, pp. 874–883. PMLR, 2020.
- Christopher Blöcker and Ingo Scholtes. Flow divergence: Comparing maps of flows with relative entropy. *arXiv preprint arXiv:2401.09052*, 2024.
- Christopher Blöcker, Chester Tan, and Ingo Scholtes. The map equation goes neural: Mapping network flows with graph neural networks. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. URL <https://openreview.net/forum?id=aFWx1N84Fe>.
- David Buterez, Jon Paul Janet, Dino Oglic, and Pietro Lio. Masked attention is all you need for graphs. *arXiv preprint arXiv:2402.10793*, 2024.

Dexiong Chen, Leslie O’Bray, and Karsten Borgwardt. Structure-aware transformer for graph representation learning. In *International Conference on Machine Learning*, pp. 3469–3489. PMLR, 2022.

Lei Chen, Zhengdao Chen, and Joan Bruna. On graph neural networks versus graph-augmented {mlp}s. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=tjqI7w64JG2>.

Wei-Lin Chiang, Xuanqing Liu, Si Si, Yang Li, Samy Bengio, and Cho-Jui Hsieh. Cluster-gcn: An efficient algorithm for training deep and large graph convolutional networks. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, KDD ’19, pp. 257–266, New York, NY, USA, 2019. Association for Computing Machinery. ISBN 9781450362016. doi: 10.1145/3292500.3330925. URL <https://doi.org/10.1145/3292500.3330925>.

Gabriele Corso, Hannes Stark, Stefanie Jegelka, Tommi Jaakkola, and Regina Barzilay. Graph neural networks. *Nature Reviews Methods Primers*, 4(1):17, 2024.

Enyan Dai, Charu Aggarwal, and Suhang Wang. Nrgnn: Learning a label noise resistant graph neural network on sparsely and noisily labeled graphs. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, KDD ’21, pp. 227–236, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 9781450383325. doi: 10.1145/3447548.3467364. URL <https://doi.org/10.1145/3447548.3467364>.

Fatemeh Daneshfar, Sayvan Soleymanbaigi, Pedram Yamini, and Mohammad Sadra Amini. A survey on semi-supervised graph clustering. *Engineering Applications of Artificial Intelligence*, 133:108215, 2024. ISSN 0952-1976. doi: <https://doi.org/10.1016/j.engappai.2024.108215>. URL <https://www.sciencedirect.com/science/article/pii/S0952197624003737>.

Yash Deshpande, Subhabrata Sen, Andrea Montanari, and Elchanan Mossel. Contextual stochastic block models. *Advances in Neural Information Processing Systems*, 31, 2018.

Fnu Devvrit, Aditya Sinha, Inderjit S Dhillon, and Prateek Jain. S3GC: Scalable self-supervised graph clustering. In Alice H. Oh, Alekh Agarwal, Danielle Belgrave, and Kyunghyun Cho (eds.), *Advances in Neural Information Processing Systems*, 2022. URL <https://openreview.net/forum?id=1d12V3vLZ5>.

Maximilien Dreveton, Felipe Schreiber Fernandes, and Daniel R. Figueiredo. Exact recovery and bregman hard clustering of node-attributed stochastic block model. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023. URL <https://openreview.net/forum?id=TjJJmcHw9p>.

O Duranthon and L Zdeborová. Optimal inference in contextual stochastic block models. *arXiv preprint arXiv:2306.07948*, 2023a.

Odilon Duranthon and Lenka Zdeborová. Neural-prior stochastic block model. *Machine Learning: Science and Technology*, 4(3):035017, 2023b.

Matthias Fey and Jan Eric Lenssen. Fast graph representation learning with pytorch geometric. *arXiv preprint arXiv:1903.02428*, 2019.

Ben Finkelshtein, İsmail İlkan Ceylan, Michael Bronstein, and Ron Levie. Learning on large graphs using intersecting communities, 2024. URL <https://arxiv.org/abs/2405.20724>.

Yarin Gal and Zoubin Ghahramani. Dropout as a bayesian approximation: Representing model uncertainty in deep learning. In *international conference on machine learning*, pp. 1050–1059. PMLR, 2016.

Amir Ghasemian, Homa HosseiniMardi, and Aaron Clauset. Evaluating overfit and underfit in models of network community structure. *IEEE Transactions on Knowledge and Data Engineering*, 32(9):1722–1735, 2019.

Martijn Gösgens, Anton Zhiyanov, Alexey Tikhonov, and Liudmila Prokhorenkova. Good classification measures and how to find them. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan (eds.), *Advances in Neural Information Processing Systems*, 2021. URL <https://openreview.net/forum?id=TLXpi2j6F7>.

Alex Graves. Practical variational inference for neural networks. *Advances in neural information processing systems*, 24, 2011.

Aditya Grover, Aaron Zweig, and Stefano Ermon. Graphite: Iterative generative modeling of graphs. In *International conference on machine learning*, pp. 2434–2444. PMLR, 2019.

Aric A. Hagberg, Daniel A. Schult, and Pieter J. Swart. Exploring network structure, dynamics, and function using networkx. In Gaël Varoquaux, Travis Vaught, and Jarrod Millman (eds.), *Proceedings of the 7th Python in Science Conference*, pp. 11 – 15, Pasadena, CA USA, 2008. URL <https://doi.org/10.25080/TCWV9851>.

Will Hamilton, Zhitao Ying, and Jure Leskovec. Inductive representation learning on large graphs. *Advances in neural information processing systems*, 30, 2017.

Jonas Berg Hansen and Filippo Maria Bianchi. Total variation graph neural networks. In *International Conference on Machine Learning*, pp. 12445–12468. PMLR, 2023.

Kaveh Hassani and Amir Hosein Khasahmadi. Contrastive multi-view representation learning on graphs. In *International conference on machine learning*, pp. 4116–4126. PMLR, 2020.

Yixuan He, Michael Perlmutter, Gesine Reinert, and Mihai Cucuringu. MSGNN: A spectral graph neural network based on a novel magnetic signed laplacian. In *The First Learning on Graphs Conference*, 2022a. URL <https://openreview.net/forum?id=KUGwmnSdPV3>.

Yixuan He, Gesine Reinert, and Mihai Cucuringu. DIGRAC: Digraph clustering based on flow imbalance. In *The First Learning on Graphs Conference*, 2022b. URL <https://openreview.net/forum?id=UqamDYtuh9>.

Yixuan He, Gesine Reinert, Songchao Wang, and Mihai Cucuringu. Sssnet: semi-supervised signed network clustering. In *Proceedings of the 2022 SIAM International Conference on Data Mining (SDM)*, pp. 244–252. SIAM, 2022c.

Yixuan He, Xitong Zhang, Junjie Huang, Benedek Rozemberczki, Mihai Cucuringu, and Gesine Reinert. Pytorch geometric signed directed: A software package on graph neural networks for signed and directed graphs. In *The Second Learning on Graphs Conference*, 2023. URL <https://openreview.net/forum?id=mni7vnYmvY>.

Zhenyu Hou, Xiao Liu, Yukuo Cen, Yuxiao Dong, Hongxia Yang, Chunjie Wang, and Jie Tang. Graphmae: Self-supervised masked graph autoencoders. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pp. 594–604, 2022.

Weihua Hu, Matthias Fey, Marinka Zitnik, Yuxiao Dong, Hongyu Ren, Bowen Liu, Michele Catasta, and Jure Leskovec. Open graph benchmark: Datasets for machine learning on graphs, 2021. URL <https://arxiv.org/abs/2005.00687>.

- Yinan Huang, Haoyu Wang, and Pan Li. What are good positional encodings for directed graphs? *arXiv preprint arXiv:2407.20912*, 2024.
- J. D. Hunter. Matplotlib: A 2d graphics environment. *Computing in Science & Engineering*, 9(3):90–95, 2007. doi: 10.1109/MCSE.2007.55.
- Mathieu Jacomy, Tommaso Venturini, Sébastien Heymann, and Mathieu Bastian. Forceatlas2, a continuous graph layout algorithm for handy network visualization designed for the gephi software. *PloS one*, 9(6): e98679, 2014.
- Ziwei Ji, Justin D. Li, and Matus Telgarsky. Early-stopped neural networks are consistent. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan (eds.), *Advances in Neural Information Processing Systems*, 2021. URL <https://openreview.net/forum?id=vPVTsuJtGky>.
- Jinwoo Kim, Dat Tien Nguyen, Seonwoo Min, Sungjun Cho, Moontae Lee, Honglak Lee, and Seunghoon Hong. Pure transformers are powerful graph learners. In Alice H. Oh, Alekh Agarwal, Danielle Belgrave, and Kyunghyun Cho (eds.), *Advances in Neural Information Processing Systems*, 2022. URL https://openreview.net/forum?id=um2BxfgkT2_.
- Thomas N Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. *arXiv preprint arXiv:1609.02907*, 2016a.
- Thomas N Kipf and Max Welling. Variational graph auto-encoders. *arXiv preprint arXiv:1611.07308*, 2016b.
- Neville Kenneth Kitson, Anthony C Constantinou, Zhigao Guo, Yang Liu, and Kiattikun Chobtham. A survey of bayesian network structure learning. *Artificial Intelligence Review*, 56(8):8721–8814, 2023.
- Gregory M Kurtzer, Vanessa Sochat, and Michael W Bauer. Singularity: Scientific containers for mobility of compute. *PloS one*, 12(5):e0177459, 2017.
- Renaud Lambiotte and Martin Rosvall. Ranking and clustering of nodes in networks with smart teleportation. *Physical Review E—Statistical, Nonlinear, and Soft Matter Physics*, 85(5):056107, 2012.
- Juanhui Li, Harry Shomer, Haitao Mao, Shenglai Zeng, Yao Ma, Neil Shah, Jiliang Tang, and Dawei Yin. Evaluating graph neural networks for link prediction: Current pitfalls and new benchmarking. In *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2023. URL <https://openreview.net/forum?id=YdjWXrd0Th>.
- Mingchen Li, Mahdi Soltanolkotabi, and Samet Oymak. Gradient descent with early stopping is provably robust to label noise for overparameterized neural networks. In *International conference on artificial intelligence and statistics*, pp. 4313–4324. PMLR, 2020.
- Qimai Li, Xiao-Ming Wu, Han Liu, Xiaotong Zhang, and Zhichao Guan. Label efficient semi-supervised learning via graph filtering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2019.
- Derek Lim, Felix Matthew Hohne, Xiuyu Li, Sijia Linda Huang, Vaishnavi Gupta, Omkar Prasad Bhalerao, and Ser-Nam Lim. Large scale learning on non-homophilous graphs: New benchmarks and strong simple methods. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan (eds.), *Advances in Neural Information Processing Systems*, 2021. URL <https://openreview.net/forum?id=DfGu8WwT0d>.

Wanyu Lin, Zhaolin Gao, and Baochun Li. Shoestring: Graph-based semi-supervised classification with severely limited labeled data. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2020.

Yue Liu, Jun Xia, Sihang Zhou, Xihong Yang, Ke Liang, Chenchen Fan, Yan Zhuang, Stan Z. Li, Xinwang Liu, and Kunlun He. A survey of deep graph clustering: Taxonomy, challenge, application, and open resource, 2023. URL <https://arxiv.org/abs/2211.12875>.

Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=Bkg6RiCqY7>.

Liheng Ma, Chen Lin, Derek Lim, Adriana Romero-Soriano, Puneet K Dokania, Mark Coates, Philip Torr, and Ser-Nam Lim. Graph inductive biases in transformers without message passing. In *International Conference on Machine Learning*, pp. 23321–23337. PMLR, 2023.

Xinyu Mao and Jiapeng Zhang. On the power of SVD in the stochastic block model. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023. URL <https://openreview.net/forum?id=tC0r8duG9z>.

Nikhil Mehta, Lawrence Carin Duke, and Piyush Rai. Stochastic blockmodels meet graph neural networks. In *International Conference on Machine Learning*, pp. 4466–4474. PMLR, 2019.

Rui Miao, Kaixiong Zhou, Yili Wang, Ninghao Liu, Ying Wang, and Xin Wang. Rethinking independent cross-entropy loss for graph-structured data. In *Forty-first International Conference on Machine Learning*, 2024. URL <https://openreview.net/forum?id=zrQIc9mQQN>.

Soumyasundar Pal, Saber Malekmohammadi, Florence Regol, Yingxue Zhang, Yishi Xu, and Mark Coates. Non parametric graph learning for bayesian graph neural networks. In *Conference on uncertainty in artificial intelligence*, pp. 1318–1327. PMLR, 2020.

Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, et al. Scikit-learn: Machine learning in python. *the Journal of machine Learning research*, 12:2825–2830, 2011.

Yulong Pei, Jianpeng Zhang, George Fletcher, and Mykola Pechenizkiy. Infinite motif stochastic blockmodel for role discovery in networks. In *Proceedings of the 2019 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining*, pp. 456–459, 2019.

Tiago P Peixoto. Nonparametric weighted stochastic block models. *Physical Review E*, 97(1):012306, 2018.

Tiago P Peixoto. Bayesian stochastic blockmodeling. *Advances in network clustering and blockmodeling*, pp. 289–332, 2019.

Tiago P Peixoto. Merge-split markov chain monte carlo for community detection. *Physical Review E*, 102(1):012305, 2020.

Tiago P Peixoto. Revealing consensus and dissensus between network partitions. *Physical Review X*, 11(2):021003, 2021.

Tiago P Peixoto. Ordered community detection in directed networks. *Physical Review E*, 106(2):024305, 2022.

Tiago P Peixoto and Alec Kirkley. Implicit models, latent compression, intrinsic biases, and cheap lunches in community detection. *arXiv preprint arXiv:2210.09186*, 2022.

- Oleg Platonov, Denis Kuznedelev, Artem Babenko, and Liudmila Prokhorenkova. Characterizing graph datasets for node classification: Homophily-heterophily dichotomy and beyond. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023a. URL <https://openreview.net/forum?id=m7PIJW0d1Y>.
- Oleg Platonov, Denis Kuznedelev, Michael Diskin, Artem Babenko, and Liudmila Prokhorenkova. A critical look at the evaluation of GNNs under heterophily: Are we really making progress? In *The Eleventh International Conference on Learning Representations*, 2023b. URL <https://openreview.net/forum?id=tJbbQfw-5wv>.
- Lisi Qarkaxhija, Anatol E Wegner, and Ingo Scholtes. Link prediction with untrained message passing layers. *arXiv preprint arXiv:2406.16687*, 2024.
- Ladislav Rampasek, Mikhail Galkin, Vijay Prakash Dwivedi, Anh Tuan Luu, Guy Wolf, and Dominique Beaini. Recipe for a general, powerful, scalable graph transformer. In Alice H. Oh, Alekh Agarwal, Danielle Belgrave, and Kyunghyun Cho (eds.), *Advances in Neural Information Processing Systems*, 2022. URL <https://openreview.net/forum?id=1MMaNf6oxKM>.
- Noam Razin, Asaf Maman, and Nadav Cohen. Implicit regularization in hierarchical tensor factorization and deep convolutional neural networks. In *International Conference on Machine Learning*, pp. 18422–18462. PMLR, 2022.
- Emanuele Rossi, Henry Kenlay, Maria I. Gorinova, Benjamin Paul Chamberlain, Xiaowen Dong, and Michael M. Bronstein. On the unreasonable effectiveness of feature propagation in learning on graphs with missing node features. In *Learning on Graphs Conference*, 2022. URL https://openreview.net/forum?id=qe_qOarxjg.
- Michael Scholkemper and Michael T Schaub. An optimization-based approach to node role discovery in networks: Approximating equitable partitions. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023. URL <https://openreview.net/forum?id=ztDxO15N7f>.
- Oleksandr Shchur and Stephan Günnemann. Overlapping community detection with graph neural networks. *arXiv preprint arXiv:1909.12201*, 2019.
- Oleksandr Shchur, Maximilian Mumme, Aleksandar Bojchevski, and Stephan Günnemann. Pitfalls of graph neural network evaluation. *Relational Representation Learning Workshop, NeurIPS 2018*, 2018.
- Amit K Shukla, Manvendra Janmaijaya, Ajith Abraham, and Pranab K Muhuri. Engineering applications of artificial intelligence: A bibliometric analysis of 30 years (1988–2018). *Engineering applications of artificial intelligence*, 85:517–532, 2019.
- Martin Simonovsky and Nikos Komodakis. Graphvae: Towards generation of small graphs using variational autoencoders. In *Artificial Neural Networks and Machine Learning–ICANN 2018: 27th International Conference on Artificial Neural Networks, Rhodes, Greece, October 4–7, 2018, Proceedings, Part I* 27, pp. 412–422. Springer, 2018.
- Jelena Smiljanic, Christopher Blöcker, Anton Holmgren, Daniel Edler, Magnus Neuman, and Martin Rosvall. Community detection with the map equation and infomap: Theory and applications. *arXiv preprint arXiv:2311.04036*, 2023.
- Jian Tang and Renjie Liao. Graph neural networks for node classification. *Graph Neural Networks: Foundations, Frontiers, and Applications*, pp. 41–61, 2022.
- RAPIDS Development Team. *RAPIDS: Libraries for End to End GPU Data Science*, 2023. URL <https://rapids.ai>.

Shantanu Thakoor, Corentin Tallec, Mohammad Gheshlaghi Azar, Mehdi Azabou, Eva L Dyer, Remi Munos, Petar Veličković, and Michal Valko. Large-scale representation learning on graphs via bootstrapping. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=0UXT6PpRpW>.

The pandas development team. pandas-dev/pandas: Pandas. URL <https://github.com/pandas-dev/pandas>.

Yu Tian, Long Zhao, Xi Peng, and Dimitris Metaxas. Rethinking kernel methods for node representation learning on graphs. *Advances in neural information processing systems*, 32, 2019.

Lawrence Tray and Ioannis Kontoyiannis. The feature-first block model. *arXiv preprint arXiv:2105.13762*, 2021.

Anton Tsitsulin, John Palowitch, Bryan Perozzi, and Emmanuel Müller. Graph clustering with graph neural networks. *Journal of Machine Learning Research*, 24(127):1–21, 2023.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.

Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks. In *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=rJXMpikCZ>.

Petar Veličković, William Fedus, William L. Hamilton, Pietro Liò, Yoshua Bengio, and R Devon Hjelm. Deep graph infomax. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=rklz9iAckQ>.

Sheng Wan, Yibing Zhan, Liu Liu, Baosheng Yu, Shirui Pan, and Chen Gong. Contrastive graph poisson networks: Semi-supervised learning with extremely limited labels. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan (eds.), *Advances in Neural Information Processing Systems*, 2021. URL <https://openreview.net/forum?id=ek0RuhPoGiD>.

Jie Wang, Jiye Liang, Kaixuan Yao, Jianqing Liang, and Dianhui Wang. Graph convolutional autoencoders with co-learning of graph structure and node attributes. *Pattern Recognition*, 121:108215, 2022.

Shiping Wang, Jinbin Yang, Jie Yao, Yang Bai, and William Zhu. An overview of advanced deep graph node clustering. *IEEE Transactions on Computational Social Systems*, 11(1):1302–1314, 2023.

Lingfei Wu, Peng Cui, Jian Pei, Liang Zhao, and Xiaojie Guo. Graph neural networks: Foundation, frontiers and applications. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, KDD ’22, pp. 4840–4841, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450393850. doi: 10.1145/3534678.3542609. URL <https://doi.org/10.1145/3534678.3542609>.

Zhu Xiaojin and Ghahramani Zoubin. Learning from labeled and unlabeled data with label propagation. In *Tech. Rep., Technical Report CMU-CALD-02-107*. Carnegie Mellon University, 2002.

Liang Yang, Fan Wu, Junhua Gu, Chuan Wang, Xiaochun Cao, Di Jin, and Yuanfang Guo. Graph attention topic modeling network. In *Proceedings of The Web Conference 2020*, WWW ’20, pp. 144–154, New York, NY, USA, 2020. Association for Computing Machinery. ISBN 9781450370233. doi: 10.1145/3366423.3380102. URL <https://doi.org/10.1145/3366423.3380102>.

- Zhilin Yang, William Cohen, and Ruslan Salakhudinov. Revisiting semi-supervised learning with graph embeddings. In *International conference on machine learning*, pp. 40–48. PMLR, 2016.
- Chengxuan Ying, Tianle Cai, Shengjie Luo, Shuxin Zheng, Guolin Ke, Di He, Yanming Shen, and Tie-Yan Liu. Do transformers really perform badly for graph representation? In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan (eds.), *Advances in Neural Information Processing Systems*, 2021. URL <https://openreview.net/forum?id=OeWooOxFwDa>.
- Zhitao Ying, Jiaxuan You, Christopher Morris, Xiang Ren, Will Hamilton, and Jure Leskovec. Hierarchical graph representation learning with differentiable pooling. *Advances in neural information processing systems*, 31, 2018.
- Andy B Yoo, Morris A Jette, and Mark Grondona. Slurm: Simple linux utility for resource management. In *Workshop on job scheduling strategies for parallel processing*, pp. 44–60. Springer, 2003.
- Kai Yu, Shipeng Yu, and Volker Tresp. Soft clustering on graphs. *Advances in neural information processing systems*, 18, 2005.
- Chulhee Yun, Srinadh Bhojanapalli, Ankit Singh Rawat, Sashank Reddi, and Sanjiv Kumar. Are transformers universal approximators of sequence-to-sequence functions? In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=ByxRM0Ntvr>.
- Muhan Zhang and Yixin Chen. Link prediction based on graph neural networks. *Advances in neural information processing systems*, 31, 2018.
- Jianan Zhao, Hesham Mostafa, Michael Galkin, Michael Bronstein, Zhaocheng Zhu, and Jian Tang. Graphany: A foundation model for node classification on any graph. *arXiv preprint arXiv:2405.20445*, 2024.
- Yanqiao Zhu, Yichen Xu, Feng Yu, Qiang Liu, Shu Wu, and Liang Wang. Deep graph contrastive representation learning, 2020. URL <https://arxiv.org/abs/2006.04131>.
- Yanqiao Zhu, Yuanqi Du, Yinkai Wang, Yichen Xu, Jieyu Zhang, Qiang Liu, and Shu Wu. A survey on deep graph generation: Methods and applications. In *The First Learning on Graphs Conference*, 2022. URL <https://openreview.net/forum?id=Im8G9R1boQi>.

A ADDITIONAL BACKGROUND

In this section we provide additional background on the neural-prior stochastic block model, the contextual stochastic block model, Bayesian networks, and end-to-end semi-supervised node classification and unsupervised graph clustering with graph neural networks.

A.1 NODE-ATTRIBUTED STOCHASTIC BLOCK MODELS

Attributed graphs with cluster structure can be generated by attributed variants of a class of random graph models called stochastic block models. Stochastic block models (SBMs) are generative models of random graphs with cluster structure. A prediction $\hat{\mathbf{A}}$ of \mathbf{A} is generated from predicted assignments of nodes to clusters $\hat{\mathbf{S}}$, and a block matrix \mathbf{B} , which can be fixed or inferred from data.

$$\hat{\mathbf{A}} \sim \text{SBM}(\hat{\mathbf{S}}, \mathbf{B}) \quad (11)$$

Each element $B_{a,b}$ of the block matrix defines the probability that a node in cluster a connects to a node in cluster b . A prediction of $\hat{\mathbf{S}}$ of \mathbf{S} and, if learned, \mathbf{B} can be inferred by minimizing the reconstruction error between $\hat{\mathbf{A}}$ and \mathbf{A} . \mathbf{B} and other variables that are considered never directed observed are always written in this work without a $\hat{\cdot}$.

One variant of the SBM can be constructed such that each edge $\hat{A}_{i,j} \sim \text{Bernoulli}(W_{i,j})$ is drawn independently element-wise from a Bernoulli distribution

$$\hat{\mathbf{A}} \sim \text{SBM}_\mathbf{B}(\hat{\mathbf{S}}, \mathbf{B}) := \text{Bernoulli}(\mathbf{W}) := \text{Bernoulli}\left(\sigma\left(\hat{\mathbf{S}}\mathbf{B}\hat{\mathbf{S}}^T\right)\right) \quad (12)$$

where σ is an activation function applied element-wise that ensures, when necessary, that \mathbf{W} is a matrix of valid edge probabilities e.g. sigmoid. For the case when \mathbf{S} is normalized such that cluster assignments for each node sums to 1, σ can simply be an identity map.

This work builds on two node-attributed variants of the SBM, the *neural-prior stochastic block model* (Duranthon & Zdeborová, 2023b) and the *contextual stochastic block model* (Deshpande et al., 2018).

A.1.1 NEURAL-PRIOR STOCHASTIC BLOCK MODEL

A generative model of $\hat{\mathbf{S}}$ and $\hat{\mathbf{A}}$ from \mathbf{X} has been proposed by Duranthon & Zdeborová (2023b), termed the neural-prior stochastic block model, which can be expressed as

$$\hat{\mathbf{A}} \sim \text{NSBM}_o(\mathbf{X}) := \text{SBM}(\text{NN}_{\mathbf{X} \rightarrow \hat{\mathbf{S}}}(\mathbf{X}), \mathbf{B}) \quad (13)$$

where $\hat{\mathbf{S}} = \text{NN}_{\mathbf{X} \rightarrow \hat{\mathbf{S}}}(\mathbf{X})$ and $\text{NN}_{\mathbf{X} \rightarrow \hat{\mathbf{S}}}$ can be any neural network function of \mathbf{X} that generates a matrix $\hat{\mathbf{S}}$, composed with a stochastic block model SBM. This neural-prior stochastic block model is denoted NSBM_o to distinguish it from its variant proposed later in this work.

A.1.2 CONTEXTUAL STOCHASTIC BLOCK MODEL

A generative model of $\hat{\mathbf{X}}$ and $\hat{\mathbf{A}}$ from $\hat{\mathbf{S}}$ has been proposed by Deshpande et al. (2018), termed the *contextual stochastic block model*, later studied and developed in works such as Dreveton et al. (2023); Duranthon & Zdeborová (2023a), and can be expressed as

$$\hat{\mathbf{X}}, \hat{\mathbf{A}} \sim \text{CSBM}_o(\hat{\mathbf{S}}, \mathbf{B}) := P_{\hat{\mathbf{S}} \rightarrow \hat{\mathbf{X}}}(\hat{\mathbf{S}}), \text{SBM}(\hat{\mathbf{S}}, \mathbf{B}) \quad (14)$$

such that $\hat{\mathbf{A}}$ is generated by an SBM as before, and $\hat{\mathbf{X}}$ is sampled from a distribution $P_{\hat{\mathbf{S}} \rightarrow \hat{\mathbf{X}}}(\hat{\mathbf{S}})$. This contextual-prior stochastic block model is denoted CSBM_o to distinguish it from its variant proposed later in this work.

A.2 GENERATIVE MODELS AS BAYESIAN NETWORKS

The dependence of variables in generative models such as the stochastic block model (SBM) and two of its node-attributed variants, NSBM_o and CSBM_o can be visualized as directed acyclic graphical models – Bayesian networks – as shown in figure 2. Bayesian networks (Kitson et al., 2023) can be helpful to visualize conditional dependencies and are used in section 2 to build upon the generative models introduced in this section. A directed edge from a variable a to a variable b in a Bayesian network represents a dependence of b on a . The Bayesian networks in figure 3 show that NSBM_o and CSBM_o are two possible Bayesian networks that generate graph structures $\hat{\mathbf{A}}$ from clusters $\hat{\mathbf{S}}$ and with \mathbf{A} and \mathbf{X} (or $\hat{\mathbf{X}}$) conditionally independent given $\hat{\mathbf{S}}$. There are, however, other generative models connecting \mathbf{X} and \mathbf{A} (or their predictions) and $\hat{\mathbf{S}}$, such as graph neural networks for (semi-supervised) node classification.

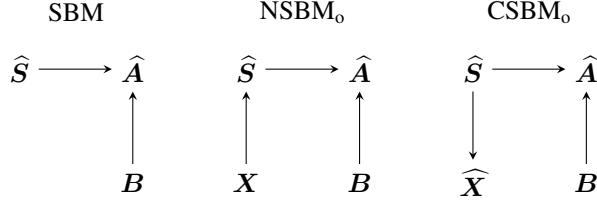


Figure 3: Bayesian network representations of the stochastic block model (SBM) and two of its node attributed variants, NSBM_o and CSBM_o .

The Bayesian network representations of graph neural networks with a purely supervised node classification objective illustrate that \mathbf{A} and \mathbf{X} are conditionally independent but made conditionally dependent given \mathbf{S} . This is the opposite case of the generative model of attributed graphs with cluster structure proposed, where $\hat{\mathbf{A}}$ and \mathbf{X} (or $\hat{\mathbf{X}}$) are conditionally dependent, but made conditionally independent given $\hat{\mathbf{S}}$.

We refer the reader to surveys of Bayesian networks (Kitson et al., 2023) for further background on Bayesian networks.

A.3 SEMI-SUPERVISED NODE CLASSIFICATION

In the case where \mathbf{S} is (partially) observed, a graph neural network for (semi-)supervised node classification can be interpreted as a generative model of labels from node attributes and adjacencies.

$$\hat{\mathbf{S}} := \text{GNN}_{(\mathbf{X}, \mathbf{A}) \rightarrow \hat{\mathbf{S}}} (\mathbf{X}, \mathbf{A}). \quad (15)$$

The cross entropy between observed (training) labels and predicted labels is commonly used as loss function for training (Kipf & Welling, 2016a; Tang & Liao, 2022; Zhao et al., 2024), which can be interpreted as a negative log-likelihood loss for the generative model (Miao et al., 2024).

$$l_{\text{supervised}} := L_S (\hat{\mathbf{S}}_{0:s,:}, \mathbf{S}_{0:s,:}) := L_{\text{cross-entropy}} (\hat{\mathbf{S}}_{0:s,:}, \mathbf{S}_{0:s,:}) \quad (16)$$

where $\hat{\mathbf{S}}$ and \mathbf{S} are of the same shape and $0 : s$ represents the indexed training split. It is assumed that $(d_S)_{\text{obs}} = d_S$, i.e. all clusters are observed in the training split so that $(d_S)_{\text{obs}} := (d_S)_{\text{obs}}$ and $\hat{\mathbf{S}}$ and \mathbf{S} are of the same shape. Its Bayesian network is visualized in figure 2. As \mathbf{A} is an independent variable, these graph neural networks are not generative models of graphs.

A.4 UNSUPERVISED CLUSTERING

In the case where \mathbf{S} is not observed, \mathbf{S} can be predicted by inferring a stochastic block model using graph neural networks. Graph neural networks for unsupervised clustering can be written as neural network functions identical to equation 15, while distinguished from node classification models by their training with unsupervised loss functions. This work proposes that unsupervised losses can be functions of $\hat{\mathbf{S}}$ and, optionally, any combination of $\hat{\mathbf{X}}$ and $\hat{\mathbf{A}}$ such that

$$l_{\text{unsupervised}} := L_{\text{clustering}} (\hat{\mathbf{S}}, \mathbf{X}, \mathbf{A}) := L_E (\hat{\mathbf{A}}, \mathbf{A}) + L_V (\hat{\mathbf{X}}, \mathbf{X}) + L_{\text{regularization}} \quad (17)$$

In architectures such as DiffPool (Ying et al., 2018), NOCD (Shchur & Günnemann, 2019), MinCut-Pool (Bianchi et al., 2020), DMoN (Tsitsulin et al., 2023), and TVGNN (Hansen & Bianchi, 2023), $L_V (\hat{\mathbf{X}}, \mathbf{X}) = 0$ and the unsupervised clustering loss only depends on \mathbf{S} and \mathbf{A} .

For example, DiffPool employs “entropy regularization” $L_{\text{regularization}} := -\frac{1}{n} \sum_{i=1}^n \sum_{a=1}^{d_S} S_{i,a} \log(S_{i,a})$ to encourage cluster assignments that are close to one-hot vectors, and reconstructs \mathbf{A} as $\hat{\mathbf{A}} := \hat{\mathbf{S}}\hat{\mathbf{S}}^T$, computing L_E as the Frobenius norm of the element-wise difference between $\hat{\mathbf{A}}$ and \mathbf{A} , which can be interpreted as a reconstruction negative log-likelihood of an SBM with weighted edges and an identity block matrix $\mathbf{B} = \mathbf{I}$.

Extending the perspective of Peixoto & Kirkley (2022), other clustering approaches including spectral clustering (MinCutPool (Bianchi et al., 2020)) and modularity (DMoN (Tsitsulin et al., 2023)) can be interpreted as also inferring an *implicit* stochastic block model.

NOCD (Shchur & Günemann, 2019) employs a “Bernoulli-Poisson” model, equivalent to variants of random dot product graphs and latent space models (Athreya et al., 2018), and a special case of the SBM_B formulation in equation 12 with a fixed identity block matrix $\mathbf{B} = \mathbf{I}$ and a custom activation

$$\mathbf{W} := \sigma_{\text{NOCD}}(\hat{\mathbf{S}}\hat{\mathbf{S}}^T) := 1 - \exp(\text{ReLU}(\hat{\mathbf{S}}\hat{\mathbf{S}}^T)) \quad (18)$$

Sampling a set of positive edges E_+ and a set of negative edges E_- of equal cardinality $|E_+| = |E_-| = m$ from \mathbf{A} , the negative log-likelihood L_E of reconstructing \mathbf{A} from $\hat{\mathbf{A}}$ can be approximated as the binary cross entropy between the set of sampled edges $E_{\text{sampled}} := E_+ \cup E_-$ and their corresponding probabilities in \mathbf{W} , a common balanced link prediction objective (Li et al., 2023).

$$L_E := -\frac{1}{m} \sum_{(i,j) \in E_{\text{sampled}}} (A_{i,j} \log(W_{i,j}) + (1 - A_{i,j}) \log(1 - W_{i,j})) \quad (19)$$

where $W_{i,j} = \sigma_{\text{NOCD}}((\mathbf{S})_{i,:} \odot \mathbf{S}_{j,:})$ and \odot represents the element-wise or Hadamard product. This approach to computing the reconstruction loss of \mathbf{A} can enable reduced computational costs by calculating only a subset of entries of \mathbf{W} . This work adapts this loss for generative models later proposed, as described in the section 2.

Methods such as DiffPool and NOCD that explicitly reconstruct a graph can be interpreted as graph autoencoders (GAE_o) (Kipf & Welling, 2016b), with an encoder generating $\hat{\mathbf{S}}$ according to equation 15, and a general decoder reconstructing \mathbf{A} as $\hat{\mathbf{A}} := \text{decoder}(\mathbf{S})$. A Bayesian network representation of these models in figure 2 shows how \mathbf{X} is treated as conditional evidence (Grover et al., 2019), and the ability of models to learn complex conditional dependencies between \mathbf{S} and \mathbf{A} and between \mathbf{S} and \mathbf{X} is limited.

B EXTENDING THE FRAMEWORK

In this section we describe extensions of our proposed framework to model other graph types and discuss potential extensions to other graph learning tasks.

B.1 GRAPH TYPES

Our proposed framework can be extended to model other graph types such as weighted graphs, directed graphs, and graphs with edge and graph attributes.

B.1.1 WEIGHTED GRAPHS

To model weighted graphs, a different distribution SBM $(\hat{\mathbf{S}}, \mathbf{B}) = P_{\mathbf{W} \rightarrow \hat{\mathbf{A}}}(\mathbf{W})$ (and corresponding reconstruction loss L_E) modelling continuous edge weights can be used (instead of the Bernoulli distribution in

equation 12). For example, for a Gaussian distribution with zero mean and unit variance, the reconstruction loss can be the mean squared error between $\hat{\mathbf{A}}$ and \mathbf{A} . For stochastic generation of $\hat{\mathbf{A}}$, $\hat{\mathbf{S}}$ (and, optionally, \mathbf{B}) can be generated with a stochastic neural network, e.g. with (Monte Carlo) dropout (Gal & Ghahramani, 2016) (see appendix B.2.3 for details).

B.1.2 DIRECTED GRAPHS

An SBM can be made directed if its block matrix is directed, which can be fixed or learned. The following subsection extending the framework to edge and graph attributes describes how a directed block matrix can be learned for the proposed NSBM and CSBM.

B.1.3 EDGE AND GRAPH ATTRIBUTES

To model node attributes \mathbf{X}_V , edge attributes \mathbf{X}_E , and graph attributes \mathbf{X}_G , the NSBM and CSBM we propose can be extended as follows. The graph neural network encoder of the autoencoders can be a neural network that is able to learn with edge attributes such as SAT (Chen et al., 2022), and a global pooling operation such as sum, mean, or max pooling can be incorporated to produce graph attributes.

\mathbf{U} can be interpreted as analogous to the node attributes in the NSBM and CSBM. $\hat{\mathbf{A}}$ and \mathbf{X} or $(\hat{\mathbf{X}})$ are conditionally independent given $\hat{\mathbf{S}}$ but conditionally dependent given $\hat{\mathbf{X}}_E$.

The NSBM can be extended to model weighted and directed graphs with edge and graph attributes as follows.

$$\begin{aligned}
 \mathbf{U}_V &:= \text{NN}_{\mathbf{X}_V \rightarrow \mathbf{U}_V} \left(\mathbf{X}_V \right) \\
 \mathbf{U}_G &:= \text{NN}_{\mathbf{X}_G \rightarrow \mathbf{U}_G} \left(\mathbf{X}_G \right) \\
 \mathbf{U} &:= \begin{bmatrix} \mathbf{U}_V \\ \mathbf{U}_G \end{bmatrix} \\
 \hat{\mathbf{S}} &:= \text{NN}_{\mathbf{U} \rightarrow \hat{\mathbf{S}}} (\mathbf{U}) \\
 \mathbf{Z}_\uparrow, \quad \mathbf{Z}_\downarrow &:= \text{NN}_{\hat{\mathbf{S}} \rightarrow (\mathbf{Z}_\uparrow, \mathbf{Z}_\downarrow)} (\hat{\mathbf{S}}) \\
 \mathbf{B} &:= \phi(\mathbf{Z}_\downarrow^T \mathbf{Z}_\uparrow) \\
 \hat{\mathbf{A}} &\sim \text{SBM} \left(\hat{\mathbf{S}}_{:n,:}, \mathbf{B} \right) \\
 \left(\hat{\mathbf{X}}_E \right)_{i,j,:} &:= \hat{A}_{ij} \text{NN}_{\mathbf{U} \rightarrow \hat{\mathbf{X}}_E} \left(\left(\mathbf{U} \right)_{i,:} \odot \left(\mathbf{U} \right)_{j,:} \right)
 \end{aligned}$$

where SBM broadly represents any suitable stochastic block model including weighted variants discussed earlier in this section.

In this approach, transformed graph attributes are treated as transformed attributes of a virtual node.




Figure 4: Bayesian networks of the NSBM and CSBM extended to generate graphs with node, edge, and graph attributes.

The CSBM can also be extended to model weighted and directed graphs with edge and graph attributes as follows.

$$\begin{aligned}
 \mathbf{U} &:= \text{NN}_{\hat{\mathbf{S}} \rightarrow \mathbf{U}} (\hat{\mathbf{S}}) \\
 \widehat{\mathbf{X}}_V &:= \text{NN}_{\mathbf{U} \rightarrow \widehat{\mathbf{X}}_V} (\mathbf{U}) \\
 \widehat{\mathbf{X}}_G &:= \text{NN}_{\mathbf{U} \rightarrow \widehat{\mathbf{X}}_G} \left(\sum_i^n (\mathbf{U})_{i,:} \right) \\
 \mathbf{Z}_\uparrow, \mathbf{Z}_\downarrow &:= \text{NN}_{\hat{\mathbf{S}} \rightarrow (\mathbf{Z}_\uparrow, \mathbf{Z}_\downarrow)} (\hat{\mathbf{S}}) \\
 \mathbf{B} &:= \phi (\mathbf{Z}_\downarrow^T \mathbf{Z}_\uparrow) \\
 \widehat{\mathbf{A}} &\sim \text{SBM} (\hat{\mathbf{S}}, \mathbf{B}) \\
 \left(\widehat{\mathbf{X}}_E \right)_{i,j,:} &:= \hat{A}_{ij} \text{NN}_{\mathbf{U} \rightarrow \widehat{\mathbf{X}}_E} \left(\left(\mathbf{U} \right)_{i,:} \underset{u}{\odot} \left(\mathbf{U} \right)_{j,:} \right)
 \end{aligned}$$

B.2 OTHER GRAPH LEARNING TASKS

Here we discuss potential extensions of our proposed framework to other graph learning tasks such as semi-supervised graph clustering with unseen and unbalanced clusters, edge and graph-level tasks, and stochastic generation of attributed graphs with cluster structure.

B.2.1 UNSEEN AND UNBALANCED CLUSTERS

When unknown, the cardinality d_S of the set of node labels S can be learned by setting an upper bound estimate $(d_S)_{\max} \geq (d_S)_{\text{obs}}$. One heuristic for choosing this is upper bound is \sqrt{n} , found from empirical studies of real-world graphs (Ghasemian et al., 2019). The first $(d_S)_{\text{obs}}$ rows of $\hat{\mathbf{S}}$ can be trained to align with supervised training labels, while the remaining $(d_S)_{\max} - (d_S)_{\text{obs}}$ learn unseen labels. Implicit regularization

of neural networks (Razin et al., 2022) and optional explicit regularization can learn compressed representations, learning a predicted $\widehat{d}_S < (d_S)_{\max}$ as demonstrated in the unsupervised clustering experiments in Blöcker et al. (2024).

Methods such as NOCD (Shchur & Günnemann, 2019), Neuromap (Blöcker et al., 2024), and the SBM_{NN} we propose avoid assuming a balanced clustering, and can be used in place of methods such as MinCutPool (Bianchi et al., 2020), DMoN (Tsitsulin et al., 2023), and ACCPool (Hansen & Bianchi, 2023), in the setting of unbalanced clusters.

B.2.2 EDGE AND GRAPH-LEVEL TASKS

The graph structure reconstruction objectives we study and propose in our framework is compatible with common link prediction objectives (see appendix A). The extension of our framework to link prediction is straightforward, as the reconstruction loss L_E can be replaced with a link prediction loss such as the binary cross entropy loss in equation 19.

The extensions to model edge and graph attributes we discussed in the previous section can be used adapt the framework to edge and graph-level tasks such as edge classification and regression, and graph classification and regression by optimizing their respective reconstruction losses.

B.2.3 STOCHASTIC GENERATION OF ATTRIBUTE GRAPHS WITH CLUSTER STRUCTURE

While we study and develop our framework using common neural network that are designed to be deterministic at test time, the proposed NSBM and CBAE generative models can be extended to sample from learned distributions at test time. As all neural networks in our experiments were trained with dropout, Monte Carlo dropout (Gal & Ghahramani, 2016) can be used to sample from the learned distribution at test time.

The extensions to edge and graph attributes we discussed in the previous section (appendix B.1) can be used to adapt the NSBM and CSBM to generate graphs with edge and graph attributes.

Our framework is also designed to be compatible with other deep generative models (Zhu et al., 2022), and the NSBM and CBAE with objectives such as variational learning, normalizing flows, and score matching can be used to learn and generate attributed graphs with cluster structure.

C SCALABILITY AND COMPUTATIONAL COMPLEXITY

Due to the generality of our proposed framework, the computational complexity of the proposed NSBM and CSBM and their respective autoencoders is not fixed, but depends on the neural network modules, clustering methods, and regularization methods used. For example, if prototypical transformers with quadratic complexity in the number of nodes are used as neural networks of the NSBM, the computational complexity of the NSBM will also be quadratic in the number of nodes. However, if an MLP is used, the computational complexity of the NSBM will be linear in the number of nodes, unless a clustering method of quadratic complexity such as DiffPool (Ying et al., 2018) is used.

While we primarily focus on the full-batch setting, which is unamenable to large-scale graphs, the proposed framework can in principle be extended to allow mini-batching by using methods such as Lim et al. (2021) and by combining any (graph) neural network and clustering method that allows for mini-batching. We leave the exploration of mini-batching for future work.

D OTHER RELATED WORK

D.1 SEMI-SUPERVISED NODE CLASSIFICATION

Semi-supervised node classification (Kipf & Welling, 2016a) studies the problem of classifying nodes in the setting of sparse training labels. The inductive bias provided by the observed graph can be effectively utilized by methods such as label propagation (Xiaojin & Zoubin, 2002; Li et al., 2019; Lin et al., 2020; Rossi et al., 2022).

Since the seminal work of semi-supervised node classification with graph convolutional networks (Kipf & Welling, 2016a), several later works study a setting of sparser training labels. Li et al. (2019) combines label propagation and graph filters proposing a “generalized label propagation” framework, while Lin et al. (2020) studied how metric learning can improve the performance for graph convolutional networks and label propagation frameworks. Several other later works adopt the perspective of Bayesian inference, from contrastive learning (Wan et al., 2021) to Bayesian graph neural networks (Pal et al., 2020), while others such as Dai et al. (2021) consider a similar problem of noisy labels.

This work provides an alternative and complementary approach to semi-supervised classification with sparse training labels, compatible with any end-to-end neural network architecture.

D.2 GENERATIVE MODELS OF ATTRIBUTED GRAPHS WITH CLUSTER STRUCTURE

Generative models of graphs with cluster structure and their inference methods can be broadly categorized into deep learning and non-deep learning methods. Numerous non-neural network-based variants of the stochastic block model have been developed to model different notions of cluster structures (Peixoto, 2018; 2022), along with various inference algorithms (Peixoto, 2020; Mao & Zhang, 2023).

Attributed stochastic block models include the contextual stochastic block model (Deshpande et al., 2018) and the neural-prior stochastic block model (Duranton & Zdeborová, 2023b). The contextual stochastic block model generates both node attributes and graph structure from node clusters (see appendix A.1.2), and prior works have proposed various distributions to model the generation of node attributes and adjacencies from node clusters including the exponential family of distributions (Dreveton et al., 2023).

Generative models of graphs with cluster structure that incorporate neural networks in their generative processes include the aforementioned neural-prior stochastic block model (see appendix A.1.1) and the feature-first block model (Tray & Kontoyiannis, 2021). The Feature-first Block Model (Tray & Kontoyiannis, 2021) is a similar model to the neural-prior stochastic block model (Duranton & Zdeborová, 2023b) proposed with an inference algorithm proposed consisting a two-part Markov chain Monte Carlo approach with Metropolis Hastings.

In this work we propose alternative fully neural variants of the contextual and neural stochastic block model and a corresponding deep learning-based inference framework that composes common neural network modules, providing both a perspective to understand current graph neural network approaches for graph clustering and an alternative deep learning approach for graph learning that enables graph learning with transformers and multilayer perceptrons.

D.3 SEMI-SUPERVISED DEEP GRAPH CLUSTERING

The application of the neural-prior stochastic block model and contextual stochastic block model to semi-supervised clustering been explored in synthetic data, inferred by an approximate message passing and belief propagation scheme (Duranton & Zdeborová, 2023b;a). This work proposes a fully neural version of the

contextual stochastic block model, and formulates the inference of the proposed NSBM and CSBM variant as neural network learning with gradient descent.

Graph neural network approaches for semi-supervised clustering have been explored in He et al. (2022c;b;a). Leveraging a supervised clustering objective based on graph kernel methods (Tian et al., 2019), He et al. (2022c) proposed a semi-supervised node clustering framework SSSNet with graph neural networks for signed graphs, later adapted for directed graphs He et al. (2022b;a)). SSSNet is trained with an unsupervised clustering loss together with two supervised losses proposed by Tian et al. (2019): a supervised cross-entropy classification loss and a third “triplet” graph kernel-based loss. DIGRAC (He et al., 2022b) and MSGNN He et al. (2022a) extended the framework for semi-supervised clustering on synthetic directed and signed directed graphs, respectively. DIGRAC studied on synthetic graphs generated from directed stochastic block models, without attributes as an ablation study, comparing self-supervised vs supervised methods. MSGNN evaluated various graph neural network architectures for directed and signed graphs using this three-part semi-supervised loss.

Our contributions are complementary, as we draw connections with other end-to-end graph clustering approaches, propose a positional encoding-free learning setting for transformer and multilayer perceptron without graph neural network layers, and position semi-supervised clustering as inference of (neural) attributed graph generative models. While the frameworks are compatible in principle, we did not include the unsupervised graph clustering losses in He et al. (2022c) and He et al. (2022b) for the following reasons: The unsupervised clustering loss proposed in He et al. (2022c) was designed for distinguishing friends and enemies and is unsuitable for unsigned graphs (He et al., 2022c). The experiments were unable to incorporate the open source implementations of the unsupervised probabilistic imbalance clustering loss in DIGRAC (He et al., 2022b; 2023) due to memory issues.

We refer the reader to Shukla et al. (2019) for a survey of semi-supervised graph clustering methods, including neural network-based methods. The survey highlights the graph neural networks for semi-supervised clustering as a promising direction for future research, and the proposed framework in this work is a step in this direction.

D.4 DEEP ATTRIBUTED GRAPH CLUSTERING

Methods for deep attributed graph clustering can be categorized into end-to-end deep attributed graph clustering methods and multi-step deep attributed graph clustering methods (Wang et al., 2023; Liu et al., 2023) End-to-end deep attributed graph clustering methods learn node embeddings and cluster them in a single step, while multi-step deep attributed graph clustering methods first learn node embeddings and then cluster them with a separate clustering method.

End-to-end deep attributed graph clustering methods include NOCD (Shchur & Günnemann, 2019), DMoN (Tsitsulin et al., 2023), and Neuromap (?).

Multi-step deep attributed graph clustering methods include Veličković et al. (2019); Zhu et al. (2020); Hassani & Khasahmadi (2020); Thakoor et al. (2022); Devvrit et al. (2022) that focus on self-supervised representation learning that enables a two-step classification or clustering approach by first learning node embeddings and then classifying these embeddings with a separate classifier or clustering them using non-neural network clustering methods such as k-means.

Our work focuses on end-to-end deep attributed graph clustering methods that learn node embeddings and cluster them in a single step. However, as our framework is compatible with any end-to-end unsupervised / self-supervised graph learning objective, it can in principle be used with end-to-end self-supervised graph representation learning methods. We leave the exploration of this for future work.

We refer the reader to Wang et al. (2023) and Liu et al. (2023) and for recent overviews and surveys of deep graph clustering methods.

D.5 LEARNING ON GRAPHS WITH “PURE” TRANSFORMERS AND MLPs

Transformers (Vaswani et al., 2017) are set-to-set and sequence-to-sequence (Yun et al., 2020) artificial neural networks that employ attention mechanisms to model relationships between input sets or sequences. Graph transformers (Rampasek et al., 2022) mix graph neural network layers with attention layers and use additional positional encodings informative of graph structure (Huang et al., 2024) to learn on graphs. Transformers without graph neural network layers can also learn on graphs using these positional encodings informative of graph structure (Ying et al., 2021; Kim et al., 2022; Ma et al., 2023). Conversely, Buterez et al. (2024) showed how Transformers without positional encodings can use adjacency matrices as attention masks to learn on graphs effectively, an approach that can be interpreted as using graph attention (Veličković et al., 2018). In this work we propose a learning method for Transformers and MLPs that neither uses graph neural network layers nor positional encodings.

Blöcker et al. (2024) and Shchur & Günnemann (2019) have evaluated the performance MLPs for unsupervised attributed graph clustering, equivalent to MLP-based NSBMs in our framework for the purely unsupervised setting. Although Shchur & Günnemann (2019) reported significantly poorer relative unsupervised clustering performance of MLPs relative to GCNs when optimizing the NOCD clustering objective, the more recent evaluation by Blöcker et al. (2024) on more recent datasets (including datasets used in our experiments) and methods demonstrated that MLPs perform similarly to GNNs with different clustering objectives. Notably the results by Blöcker et al. (2024) also demonstrated that while MLPs had significantly poorer performance versus GNNs optimizing the NOCD clustering objective on certain datasets, MLPs optimizing other clustering objectives on those datasets perform similarly to GNNs. Our work provides a complementary perspective to understand why MLPs should perform similarly to GNNs when optimizing graph clustering objectives for attributed graph clustering, and introduces a framework for evaluating end-to-end graph clustering objectives on real-world attributed graphs in the semi-supervised setting. We propose that the variability of MLP performance due to different clustering objectives could be an interesting avenue for future work.

D.6 GRAPH AUTOENCODERS WITH ATTRIBUTE RECONSTRUCTION

While prototypical graph autoencoder (Kipf & Welling, 2016b) were designed to only reconstruct graph structure, a notable exception developed in later works is the family of graph autoencoders for deep graph generation that includes GraphVAE (Simonovsky & Komodakis, 2018) that reconstruct both node attributes and graph structure (Zhu et al., 2022). Wang et al. (2022) proposed and evaluated a specific graph autoencoder architecture that reconstructs both node attributes and graph structure for attributed graph clustering and link prediction on the Planetoid datasets Cora, Citeseer, and Pubmed (Yang et al., 2016). Conversely, Hou et al. (2022) explored graph autoencoders that reconstruct attributes *instead* of graph structure for node classification and graph classification.

Our work introduces a perspective to understand the importance of reconstructing both node attributes and graph structure for learning on attributed graphs with graph neural networks, and introduces a general framework that can utilize any graph autoencoder that reconstructs both node attributes and graph structure for semi-supervised graph clustering.

D.7 CLUSTERING *for* GRAPH NEURAL NETWORKS

Graph neural networks can be designed to implicitly and/or explicitly utilize graph cluster structures in their learning. Implicit optimization of clustering objectives by graph neural networks have been studied in

works such as Hansen & Bianchi (2023) which studied how a local quadratic variation partitioning objective is minimized by message passing graph neural networks such as graph convolutional networks (Kipf & Welling, 2016a).

Explicit uses of clustering objectives include pooling (Ying et al., 2018; Bianchi et al., 2020; Hansen & Bianchi, 2023) and batch sampling (Chiang et al., 2019) for graph neural networks. Clustering-based pooling approaches in aforementioned works such as DiffPool (Ying et al., 2018), MinCutPool (Bianchi et al., 2020), and TVGNN (Hansen & Bianchi, 2023) use explicit clustering objectives to pool graphs for graph-level tasks such as graph classification and graph regression. Clustering-based subgraph methods include Chiang et al. (2019) that propose using low-memory complexity graph clustering algorithms suited to large graphs to sample and batch large graphs, reducing memory requirements for graph neural networks. Finkelshtein et al. (2024) studied a similar approach that leveraged community detection to sparsify large graphs.

Mehta et al. (2019) incorporated the stochastic block model into the design of variational graph autoencoders (Kipf & Welling, 2016b) to model sparse graphs. Yang et al. (2020) study a connection between graph attention networks (Veličković et al., 2018) and the stochastic block model, framing graph attention networks as semi-amortized variational inference of stochastic block models. Our work provides a complementary perspective to understand how neural networks can infer attributed stochastic block models. We propose end-to-end fully neural variants of attributed stochastic block models that can be used with any neural network architecture. Our framework describes how any graph neural network can infer these attributed stochastic block models with attribute reconstruction, and how non-graph neural networks can learn on attributed graphs without positional encodings or graph neural network layers such as message passing or spectral convolutions by inferring these attributed stochastic block models.

E ADDITIONAL EXPERIMENTAL DETAILS

In this section, we provide additional details on the datasets, architectures, and experimental setup used in the experiments in this work.

E.1 DATASET DETAILS

We sourced the real-world attributed graph multi-class node classification datasets in our experiments from the following code and data repositories using the following libraries.

- Coauthor CS, Coauthor Physics, Amazon Computer, and Amazon Photo from Shchur et al. (2018) using the geometric deep learning library PyTorch Geometric (Fey & Lenssen, 2019).
- roman-empire and amazon-ratings from Platonov et al. (2023b) using PyTorch Geometric.

Each undirected edge between nodes i and j in the undirected graph datasets is represented with two directed edges, one from node i to j and another from node j to i , following the approach in PyTorch Geometric (Fey & Lenssen, 2019). Node attributes are normalized so that the attributes for each node sum to 1.

10 splits for each dataset roman-empire and amazon-ratings from Platonov et al. (2023b) were made public. 1 split for each dataset ogbn-arxiv and ogbn-products from the Open Graph Benchmark (Hu et al., 2021) was made public. The splits for the Coauthor CS, Coauthor Physics, Amazon Computer, and Amazon Photo datasets from Shchur et al. (2018) were not made public, but the work proposes 10 random splits of 20 training nodes per class, 500 validation nodes, and 1000 test nodes.

We create a label-sparsified version of each dataset by selecting 2 training nodes per class and 50 validation nodes from the splits described above. We do not modify the test labels.

We run each experiment configuration 10 times on each dataset and its label-sparsified version and report the mean and standard deviation of the results. For the `roman-empire` and `amazon-ratings` datasets, we use a different public split for each run. For the Coauthor CS, Coauthor Physics, Amazon Computer, and Amazon Photo datasets we use a random split with the same number of training nodes per class, validation nodes, and test nodes for each run.

Table 2: Summary statistics of real-world attributed graph datasets used in experiments.

Dataset	Nodes	Edges	Node Attributes	Classes
Coauthor CS	18,333	163,788	6,805	15
Coauthor Physics	34,493	495,924	8,415	5
Amazon Computers	13,752	491,722	767	10
Amazon Photo	7,650	238,162	745	8
<code>roman-empire</code>	22,662	32,927	300	18
<code>amazon-ratings</code>	24,492	93,050	300	5
<code>ogbn-arxiv</code>	169,343	1,166,243	128	40
<code>ogbn-products</code>	2,449,029	61,859,140	100	47

E.2 ARCHITECTURE DETAILS

The transformer architecture used is a prototypical 6 layer encoder-only transformer (Vaswani et al., 2017), adopting its default PyTorch (Ansel et al., 2024) with 8 attention heads, 512 embedding dimensions and 2048 hidden dimensions in its feedforward block, with a bias in its linear and layernorm layers. MLP is equivalent to the feedforward block in the transformer. Both use a dropout rate of 0.1.

GCN and GraphSAGE are the graph neural network architectures used in experiments, adopting its implementations in PyTorch Geometric (Fey & Lenssen, 2019). They use 3 layers, 256 hidden dimensions, and a dropout rate of 0.5.

All models were trained with $s = 100$ cluster dimensions. This was chosen to be of the same order of magnitude as the square root of the number of nodes of the median dataset size, following empirical studies of the upper bound of the number of clusters in real-world graphs (Ghasemian et al., 2019). As graph structure \mathbf{A} is reconstructed from $\widehat{\mathbf{S}}$, the number of clusters s is chosen to be large enough so that $\widehat{\mathbf{A}} = \text{SBM}_f(\widehat{\mathbf{S}})$ is not too lossy a reconstruction of \mathbf{A} , while being small enough to limit the computational complexity of the model.

As evidenced by the additional experimental results in section F, having the number of clusters s larger than the number of classes does not significantly affect the performance of a purely supervised model.

E.3 HYPERPARAMETER SETTINGS

For a controlled experiment under a limited compute budget, hyperparameters such as learning rates or the number of layers were not tuned. Architectural hyperparameters are detailed above in the architecture details. We used the same learning rate of 0.01 for all parameters, for all experiments, optimized with AdamW (Loshchilov & Hutter, 2019). Where $L_{\text{regularization}}$ is “L2” weight decay is set to 1.0 and otherwise weight decay is set to 0.01.

E.4 UNSUPERVISED CLUSTERING OBJECTIVES

In this section we provide additional details on the unsupervised clustering objectives used in the experiments.

E.4.1 L_E OBJECTIVES

We evaluated 4 unsupervised clustering objectives – NOCD (Shchur & Günnemann, 2019), DMoN (Tsitsulin et al., 2023), Neuromap (Blöcker et al., 2024), and our proposed SBM_{NN} – as $c \approx L_E$ objectives (see section 2.1.1) in the experiments.

As NOCD can be interpreted as a special case of the SBM (see section A), we implement NOCD simply as the case where the SBM has a fixed identity block matrix $\mathbf{B} = \mathbf{I}$. Instead of the custom activation function in the NOCD implementation to generate valid edge probabilities, we use the sigmoid activation function which matches a typical binary cross entropy from logits loss used in link prediction objectives (Zhang & Chen, 2018; Li et al., 2023; Qarkaxhija et al., 2024). We argue that these implementations are equivalent.

“L2” is l_2 regularization on the weights of a neural network, motivated by the interpretation of these neural networks as generative models (Graves, 2011).

DMoN comprises two unsupervised objectives – a spectral modularity objective and a “collapse” regularization objective. We use its spectral modularity objective as an L_E objectives and its collapse regularization objective as an $L_{\text{regularization}}$ objective. We chose DMoN’s objectives as representative L_E and $L_{\text{regularization}}$ objectives as they were designed to jointly clustering node attributes and adjacencies (unsupervised) using graph neural networks.

DMoN’s spectral modularity L_E objective is calculated as

$$(L_E)_{\text{DMoN}} := -\frac{1}{2m} \text{Tr} \left(\mathbf{S}^T \left(\mathbf{A} - \frac{\mathbf{d} \otimes \mathbf{d}}{2m} \right) \mathbf{S} \right) \quad (20)$$

where $m = |E|$ is the number of edges in the graph and \mathbf{d} is a degree vector representing the sum of the rows of the adjacency matrix \mathbf{A} , and $\mathbf{d} \otimes \mathbf{d}$ is the outer product of \mathbf{d} with itself. The implementation of DMoN (Tsitsulin et al., 2023) was sourced from PyTorch Geometric (Fey & Lenssen, 2019).

Neurommap’s (Blöcker et al., 2024) codelength L_E objective is calculated as

$$(L_E)_{\text{Neuromap}} = q \log_2 q - (\mathbf{q}_m \log_2 \mathbf{q}_m)_s \mathbf{1} - (\mathbf{m}_{\text{exit}} \log_2 \mathbf{m}_{\text{exit}})_s \mathbf{1} - (\mathbf{p} \log_2 \mathbf{p})_n \mathbf{1} + (\mathbf{p}_m \log_2 \mathbf{p}_m)_s \mathbf{1} \quad (21)$$

$\mathbf{1}$ is the k -dimensional vector of ones, and logarithms are applied component-wise. Each of the terms is defined as follows.

$$q = 1 - \text{tr}(\mathbf{C}) \quad \mathbf{q}_m = \mathbf{C} \mathbf{1}_s - \text{diag}(\mathbf{C}) \quad \mathbf{m}_{\text{exit}} = (\mathbf{1}_s^\top \mathbf{C})^\top - \text{diag}(\mathbf{C}) \quad \mathbf{p}_m = \mathbf{q}_m + \mathbf{1}_s^\top \mathbf{C}$$

where

$$\mathbf{C} = \mathbf{S}^\top \mathbf{F} \mathbf{S} \quad (22)$$

and

$$\mathbf{F} = \frac{\alpha}{w_{\text{tot}}} \mathbf{A} + (1 - \alpha) \text{diag}(\mathbf{p}) \mathbf{T} \quad (23)$$

The transition matrix \mathbf{T} is calculated using the graph’s total weight $w_{\text{tot}} = \sum_{i \in V} \sum_{j \in V} w_{ij}$ and the vector of weighted node in-degrees \mathbf{d}^{in} (as Neuromap is defined for directed graphs) as

$$T_{ij} = \begin{cases} \frac{w_{ij}}{\sum_{j \in V} w_{ij}} & \text{if } \sum_{j \in V} w_{ij} > 0, \\ 0 & \text{otherwise,} \end{cases} \quad \mathbf{d}_j^{\text{in}} = \sum_{i \in V} w_{ij}.$$

and the vector \mathbf{p} of node visit rates is calculated using smart teleportation Lambiotte & Rosvall (2012) and the power iteration method which iteratively updates the visit rates $\mathbf{p}^{(t)}$ (until convergence) as

$$\mathbf{p}^{(t+1)} \leftarrow \frac{\alpha}{w_{\text{tot}}} \mathbf{d}^{\text{in}} + (1 - \alpha) \mathbf{p}^{(t)} \mathbf{T} \quad (24)$$

where $\mathbf{p}^{(0)} = \mathbf{d}^{\text{in}}$ and α is a parameter chosen proportionally to the nodes' in-degrees. The implementation of Neuromap (Blöcker et al., 2024) was sourced from its official repository at <https://github.com/chrisbloecker/neuromap>.

For the SBM_{NN} we propose, we evaluated two neural networks as $\text{NN}_{\widehat{\mathcal{S}} \rightarrow \widehat{\mathcal{Z}}}$ in equation 4 to generate \mathbf{B} – a transformer and an MLP with the same architectures and hyperparameter settings as the transformers and MLPs detailed above.

E.4.2 L_V OBJECTIVES

For the “GCN” and “GraphSAGE” models that refer to CBAE autoencoders inferring a CSBM according to equation 7 we evaluated two choices of $\text{NN}_{\widehat{\mathcal{S}} \rightarrow \widehat{\mathcal{X}}}$ – a transformer and an MLP. We set the transformer and MLP architectures and hyperparameter settings the same as the transformers and MLPs detailed above. L_V was then calculated as the mean squared error of the reconstruction $\widehat{\mathbf{X}}$ of \mathbf{X} .

E.4.3 REGULARIZATION OBJECTIVES

We evaluated two $L_{\text{regularization}}$ objectives – “L2” and “DMoN” (Tsitsulin et al., 2023).

The collapse regularization $L_{\text{regularization}}$ objective of DMoN aims to discourage the trivial partition of assigning all nodes to the same cluster and can be expressed as

$$(L_{\text{regularization}})_{\text{DMoN}} := \frac{\sqrt{s}}{n} \left\| \sum_i (\mathbf{S}_{:,i}) \right\|_2 - 1 \quad (25)$$

where $\|\cdot\|_2$ is the l_2 vector norm. The implementation of DMoN (Tsitsulin et al., 2023) was sourced from PyTorch Geometric (Fey & Lenssen, 2019).

L2 regularization was calculated as the l_2 norm of all weights in the neural network. Shchur & Günnemann (2019) showed that small L2 regularization was beneficial, and we evaluate the case of unit L2 regularization which corresponds to an exact variational and description length objective (Graves, 2011), motivated by the interpretation of these model as generative models (see section B.2.3).

E.5 EVALUATION METRICS

The experiments focused on semi-supervised *multi-class* node classification for real-world node-attributed graphs. Two evaluation metrics are used: accuracy and Matthews correlation coefficient (Gösgens et al., 2021; Platonov et al., 2023a). Both metrics are evaluated on the test set for each dataset and split. Accuracy alone is often standard as a sole evaluation metric for multi-class classification (as done in Yang et al. (2016); Kipf & Welling (2016a); Shchur et al. (2018); Platonov et al. (2023b); Hu et al. (2021)).

We also employed the Matthews correlation coefficient as an additional evaluation metric to provide a more comprehensive evaluation due to its theoretical and empirical (Gösgens et al., 2021; Platonov et al., 2023a) advantages in evaluating classifications, and its use in graph learning with encoding-free transformers without graph neural network layers such as Buterez et al. (2024). Its advantages are evident in the experiments, such as in evaluating performance on the amazon-ratings dataset (see table 1). Despite some models

achieving an accuracy of roughly 0.3 and other models achieving near zero accuracy, the Matthews correlation coefficient was zero across most models, indicating that these models were not predicting any better than random.

Other clustering evaluation metrics such as modularity or graph conductance (as used in Tsitsulin et al. (2023) for unsupervised clustering) were not used in the experiments as they are defined for non-attributed graphs.

E.6 EARLY STOPPING AND MODEL SELECTION WITHOUT A VALIDATION SET

While early stopping and model selection on the validation set is standard practice for supervised learning (Li et al., 2020; Bai et al., 2021; Ji et al., 2021), for unsupervised learning a validation set is unavailable, and early stopping and model selection on the training loss has been explored for unsupervised graph clustering (Shchur & Günnemann, 2019; Blöcker et al., 2024). As our framework positions semi-supervised clustering as from both supervised and unsupervised clustering perspectives, we evaluated early stopping and model selection on the training loss and on the validation Matthews correlation coefficient (MCC). Despite the small validation set size (50 nodes) on label sparsified splits, we found that early stopping and model selection on the validation MCC produced the best generalization performance on the test set for both default and label sparsified splits of the datasets.

E.7 SOFTWARE USED

Experiments leveraged PyTorch (Ansel et al., 2024) and PyTorch Geometric (PyG) (Fey & Lenssen, 2019) libraries for graph deep learning, and scikit learn for evaluation metrics (Pedregosa et al., 2011). Compiling results into tables utilized the library pandas (The pandas development team),, and visualizations utilized the libraries cuGraph (Team, 2023), NetworkX (Hagberg et al., 2008), and Matplotlib (Hunter, 2007). In particular, visualizations used the ForceAtlas2 (Jacomy et al., 2014) algorithm implementation by the GPU accelerated cuGraph library. All experiments were run on a SLURM (Yoo et al., 2003) cluster with Apptainer (formerly Singularity) (Kurtzer et al., 2017) using the PyG container NVIDIA NGC container registry which contains all of the software dependencies above.

F ADDITIONAL RESULTS

F.1 VISUALISATIONS OF GRAPH LABELLINGS

F.1.1 AMAZON-RATINGS DATASET, DEFAULT SPLIT WITH DEFAULT TRAIN NODES PER CLASS AND DEFAULT VALIDATION NODES.




Figure 5: Visualisations of graph labellings for the amazon-ratings dataset, default split with default train nodes per class and default validation nodes. Model selection based on training loss.




Figure 6: Visualisations of graph labellings for the amazon-ratings dataset, default split with default train nodes per class and default validation nodes. Model selection based on training set MCC.




Figure 7: Visualisations of graph labellings for the amazon-ratings dataset, default split with default train nodes per class and default validation nodes. Model selection based on validation set MCC.

F.1.2 AMAZON-RATINGS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

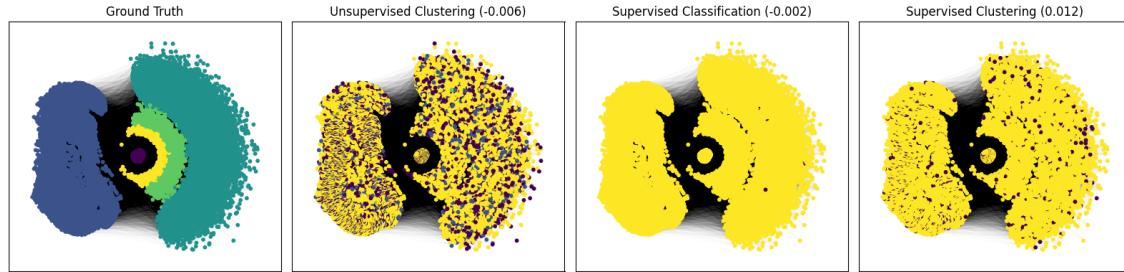


Figure 8: Visualisations of graph labellings for the `amazon-ratings` dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training loss.

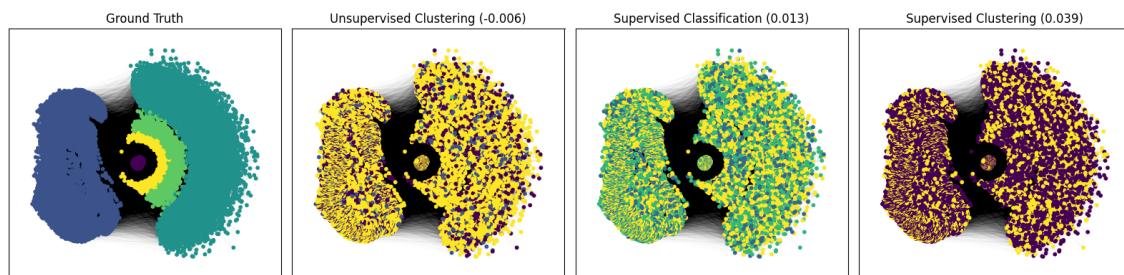


Figure 9: Visualisations of graph labellings for the `amazon-ratings` dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training set MCC.

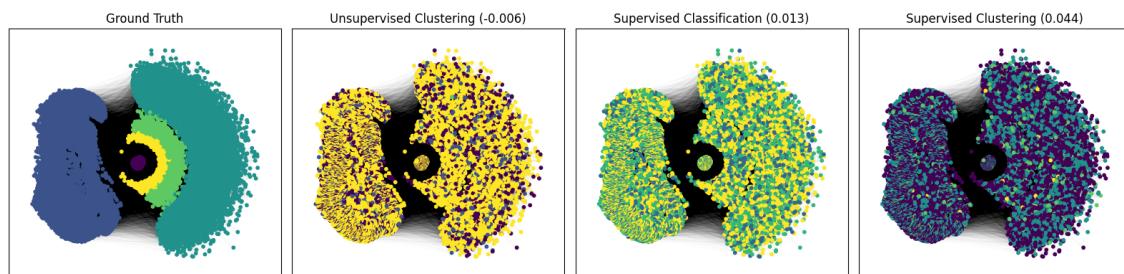


Figure 10: Visualisations of graph labellings for the `amazon-ratings` dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on validation set MCC.

F.1.3 COAUTHOR CS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

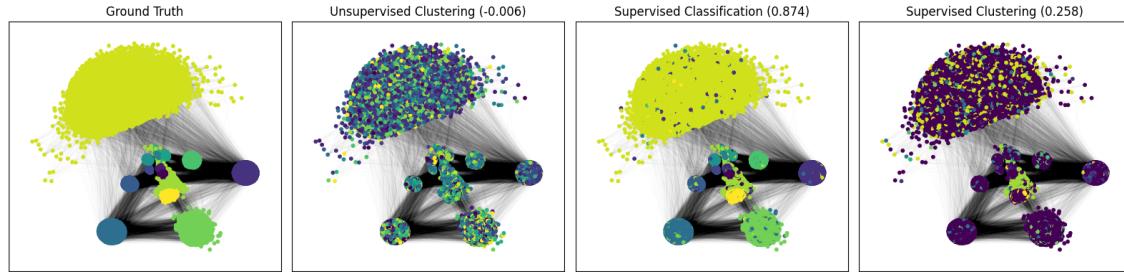


Figure 11: Visualisations of graph labellings for the Coauthor CS dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on training loss.




Figure 12: Visualisations of graph labellings for the Coauthor CS dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on training set MCC.

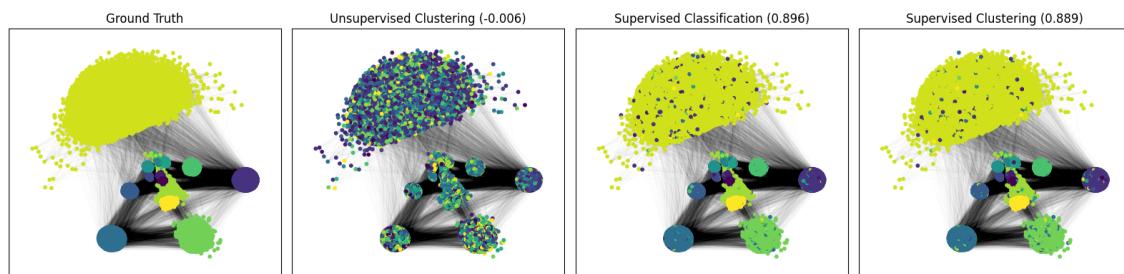


Figure 13: Visualisations of graph labellings for the Coauthor CS dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on validation set MCC.

F.1.4 COAUTHOR CS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.




Figure 14: Visualisations of graph labellings for the Coauthor CS dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training loss.




Figure 15: Visualisations of graph labellings for the Coauthor CS dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training set MCC.




Figure 16: Visualisations of graph labellings for the Coauthor CS dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on validation set MCC.

F.1.5 AMAZON COMPUTERS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

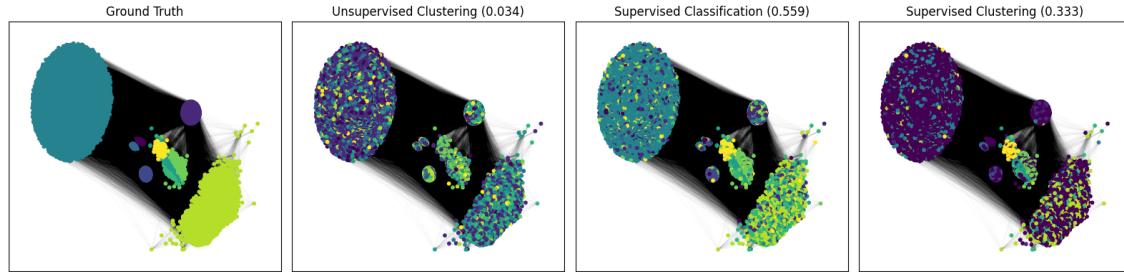


Figure 17: Visualisations of graph labellings for the Amazon Computers dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on training loss.




Figure 18: Visualisations of graph labellings for the Amazon Computers dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on training set MCC.




Figure 19: Visualisations of graph labellings for the Amazon Computers dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on validation set MCC.

F.1.6 AMAZON COMPUTERS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

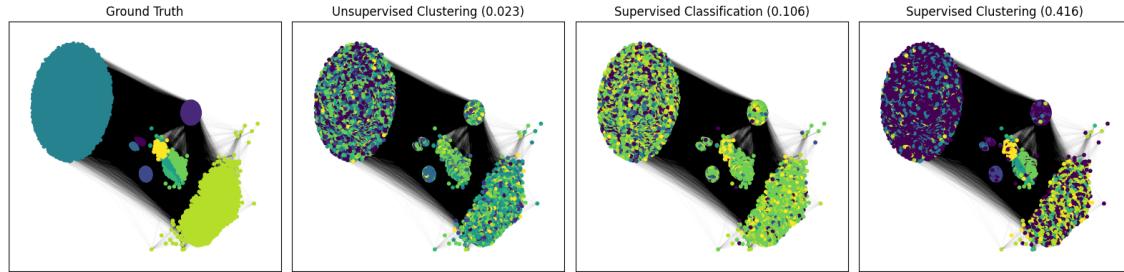


Figure 20: Visualisations of graph labellings for the Amazon Computers dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training loss.

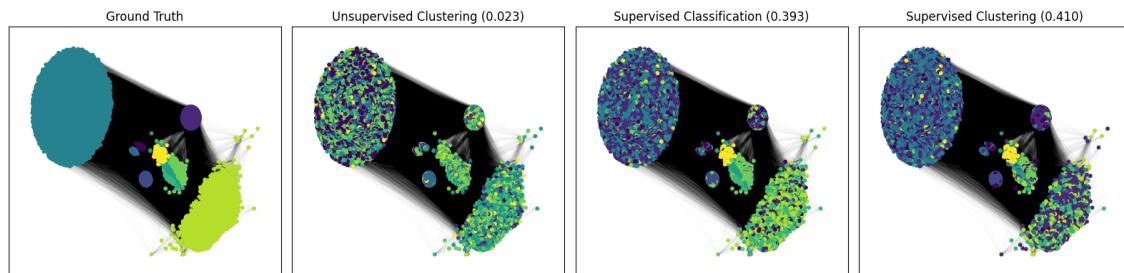


Figure 21: Visualisations of graph labellings for the Amazon Computers dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training set MCC.

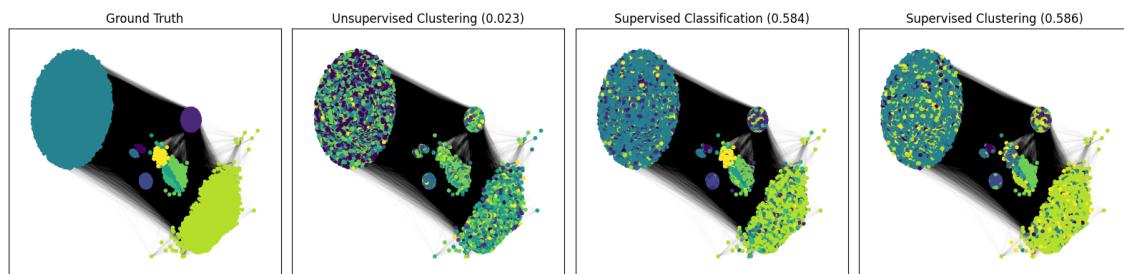


Figure 22: Visualisations of graph labellings for the Amazon Computers dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on validation set MCC.

F.1.7 AMAZON PHOTO DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

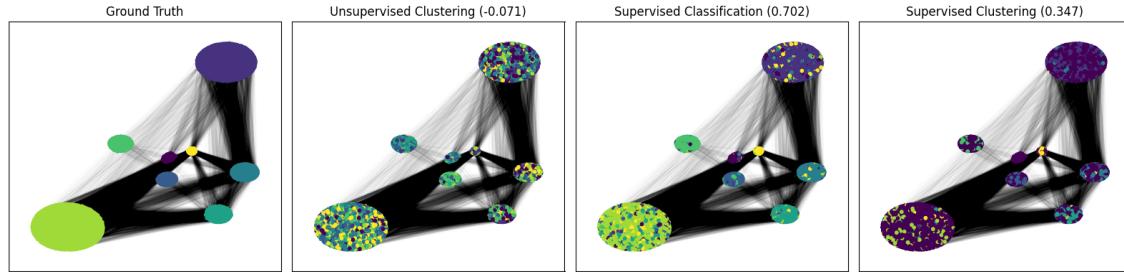


Figure 23: Visualisations of graph labellings for the Amazon Photo dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on training loss.

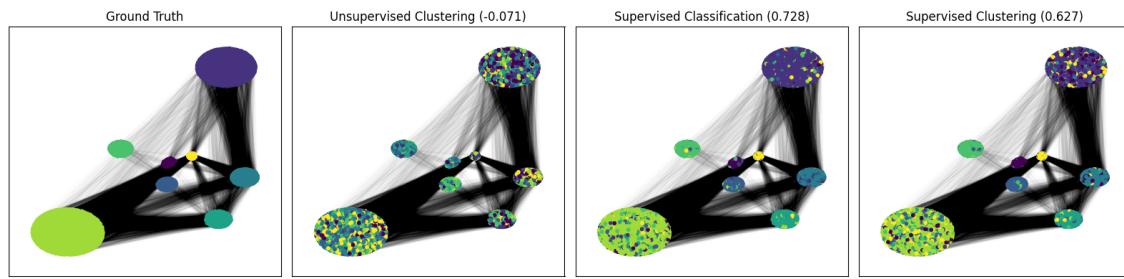


Figure 24: Visualisations of graph labellings for the Amazon Photo dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on training set MCC.




Figure 25: Visualisations of graph labellings for the Amazon Photo dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on validation set MCC.

F.1.8 AMAZON PHOTO DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.




Figure 26: Visualisations of graph labellings for the Amazon Photo dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training loss.




Figure 27: Visualisations of graph labellings for the Amazon Photo dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training set MCC.




Figure 28: Visualisations of graph labellings for the Amazon Photo dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on validation set MCC.

F.1.9 COAUTHOR PHYSICS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

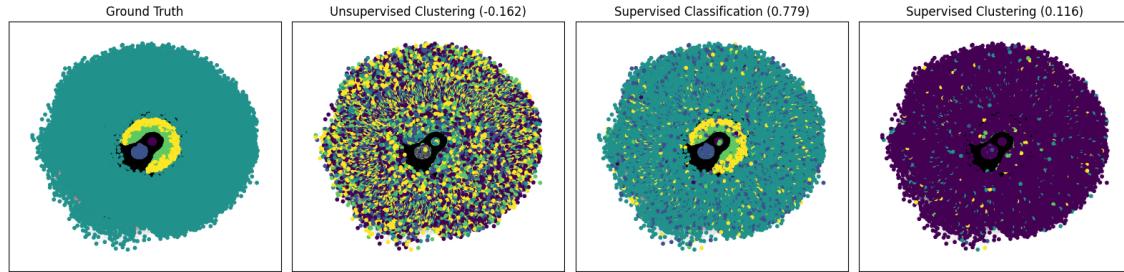


Figure 29: Visualisations of graph labellings for the Coauthor Physics dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on training loss.




Figure 30: Visualisations of graph labellings for the Coauthor Physics dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on training set MCC.




Figure 31: Visualisations of graph labellings for the Coauthor Physics dataset, default split with 20 train nodes per class and 500 validation nodes. Model selection based on validation set MCC.

F.1.10 COAUTHOR PHYSICS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

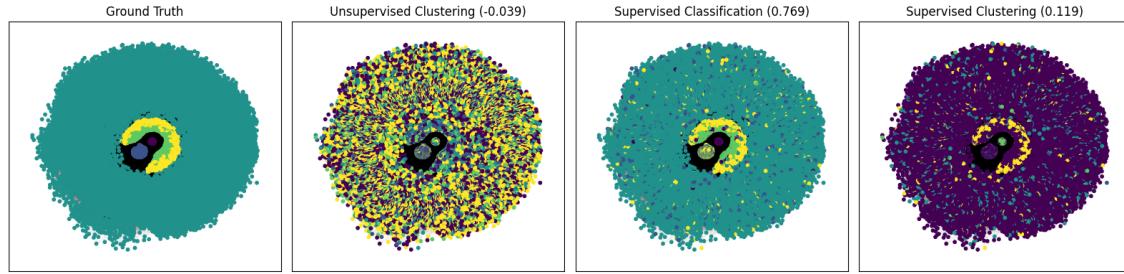


Figure 32: Visualisations of graph labellings for the Coauthor Physics dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training loss.

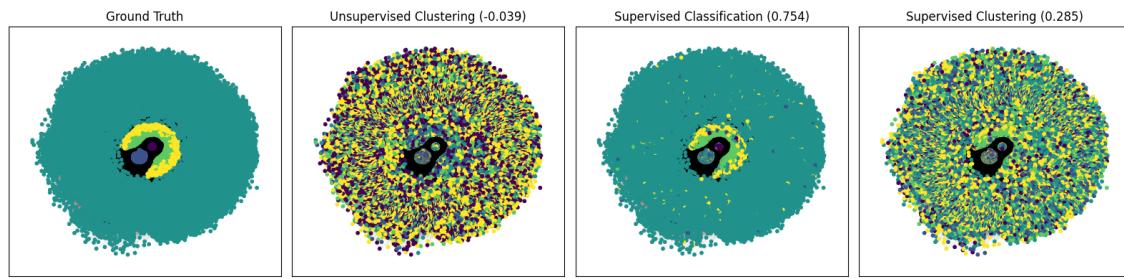


Figure 33: Visualisations of graph labellings for the Coauthor Physics dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training set MCC.

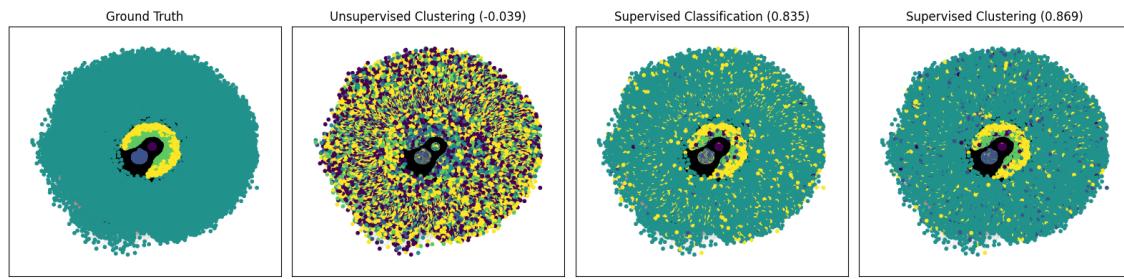


Figure 34: Visualisations of graph labellings for the Coauthor Physics dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on validation set MCC.

F.1.11 ROMAN-EMPIRE DATASET, DEFAULT SPLIT WITH DEFAULT TRAIN NODES PER CLASS AND DEFAULT VALIDATION NODES.




Figure 35: Visualisations of graph labellings for the `roman-empire` dataset, default split with default train nodes per class and default validation nodes. Model selection based on training loss.




Figure 36: Visualisations of graph labellings for the `roman-empire` dataset, default split with default train nodes per class and default validation nodes. Model selection based on training set MCC.




Figure 37: Visualisations of graph labellings for the `roman-empire` dataset, default split with default train nodes per class and default validation nodes. Model selection based on validation set MCC.

F.1.12 ROMAN-EMPIRE DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

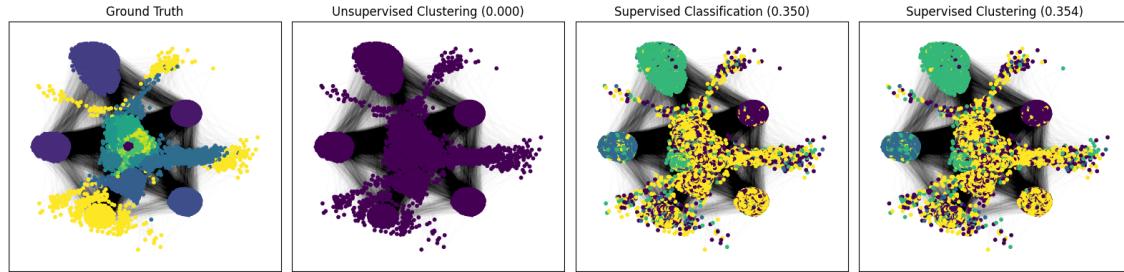


Figure 38: Visualisations of graph labellings for the `roman-empire` dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training loss.

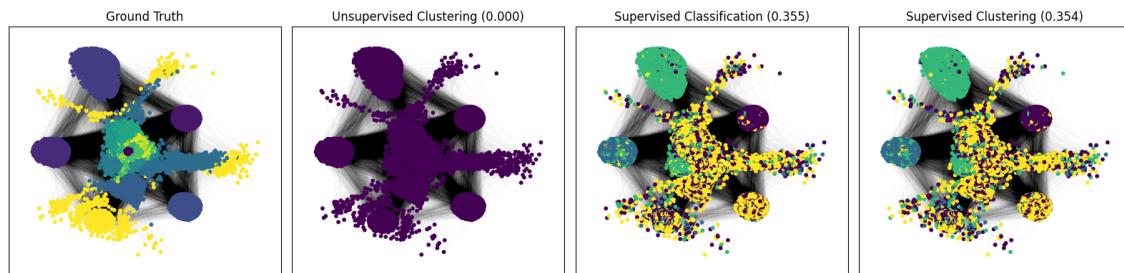


Figure 39: Visualisations of graph labellings for the `roman-empire` dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on training set MCC.

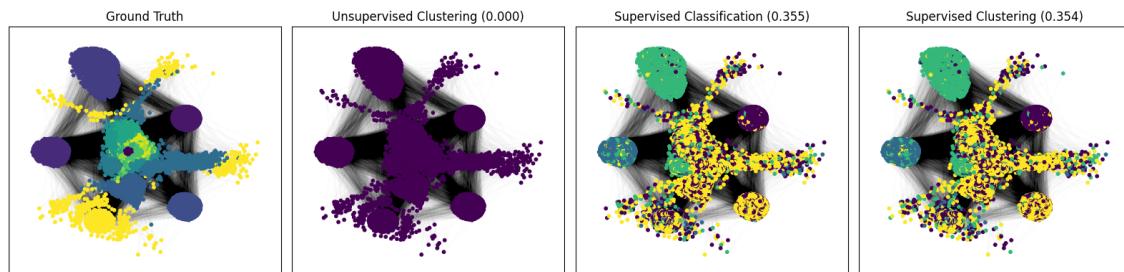


Figure 40: Visualisations of graph labellings for the `roman-empire` dataset, sparse split with 2 train nodes per class and 50 validation nodes. Model selection based on validation set MCC.

F.2 ADDITIONAL SUMMARY RESULTS

In this section we present additional summary results for both default and sparse label splits of each dataset. We compare early stopping and model selection based on training loss, training set MCC, and validation set MCC.

F.2.1 AMAZON-RATINGS DATASET, DEFAULT SPLIT WITH DEFAULT TRAIN NODES PER CLASS AND DEFAULT VALIDATION NODES.

Table 3: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.001 ± 0.004	0.368 ± 0.000
GCN	DMoN	None	\checkmark	MLP	Loss	0.001 ± 0.006	0.368 ± 0.000
GraphSAGE	None	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.354 ± 0.043
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000

Table 4: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	MCC	0.001 ± 0.003	0.027 ± 0.085
GCN	NOCD	None	\times	None	Loss	0.028 ± 0.009	0.331 ± 0.020
GraphSAGE	None	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	L2	\times	None	MCC	0.000 ± 0.001	0.023 ± 0.074
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.354 ± 0.043
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.308 ± 0.052
MLP	None	L2	N/A	None	MCC	0.002 ± 0.005	0.077 ± 0.124
MLP	Neuromap	L2	N/A	None	MCC	0.001 ± 0.004	0.090 ± 0.140

Table 5: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	MCC	0.000 ± 0.000	0.000 ± 0.000
GCN	NOCD	None	\times	None	Loss	0.028 ± 0.009	0.331 ± 0.020
GraphSAGE	None	L2	\times	None	MCC	-0.000 ± 0.001	0.027 ± 0.085
GraphSAGE	Neuromap	DMoN	\checkmark	MLP	MCC	-0.001 ± 0.002	0.005 ± 0.015
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.354 ± 0.043
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.308 ± 0.052
MLP	None	L2	N/A	None	MCC	0.002 ± 0.005	0.077 ± 0.124
MLP	DMoN	L2	N/A	None	MCC	0.000 ± 0.004	0.036 ± 0.075

F.2.2 AMAZON-RATINGS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 6: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	-0.001 ± 0.001	0.368 ± 0.000
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.004 ± 0.008	0.368 ± 0.000
GraphSAGE	None	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032

Table 7: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	L2	\times	None	MCC	-0.003 ± 0.030	0.252 ± 0.093
GCN	NOCD	None	\checkmark	MLP	Loss	0.031 ± 0.010	0.344 ± 0.006
GraphSAGE	None	L2	\times	None	MCC	0.002 ± 0.004	0.032 ± 0.076
GraphSAGE	SBM _{NN}	None	\times	None	MCC	-0.001 ± 0.002	0.026 ± 0.081
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.321 ± 0.062
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	NOCD	L2	N/A	None	MCC	0.000 ± 0.001	0.063 ± 0.134

F.2.3 COAUTHOR CS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 8: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.875 ± 0.014	0.889 ± 0.013
GCN	NOCD	DMoN	\checkmark	Transformer	Loss	0.319 ± 0.129	0.262 ± 0.159
GraphSAGE	None	None	\times	None	Loss	0.887 ± 0.007	0.899 ± 0.006
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.217 ± 0.027	0.153 ± 0.022
Transformer	None	None	N/A	None	Loss	0.038 ± 0.042	0.100 ± 0.089
Transformer	DMoN	DMoN	N/A	None	Loss	0.014 ± 0.025	0.027 ± 0.017
MLP	None	DMoN	N/A	None	Loss	0.348 ± 0.033	0.284 ± 0.055
MLP	DMoN	None	N/A	None	Loss	0.420 ± 0.050	0.365 ± 0.078

Table 9: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.875 ± 0.014	0.889 ± 0.013
GCN	DMoN	DMoN	\times	None	Loss	0.867 ± 0.013	0.881 ± 0.012
GraphSAGE	None	DMoN	\times	None	Loss	0.810 ± 0.024	0.824 ± 0.025
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.768 ± 0.020	0.782 ± 0.021
Transformer	None	L2	N/A	None	MCC	0.803 ± 0.035	0.823 ± 0.032
Transformer	DMoN	L2	N/A	None	MCC	0.783 ± 0.063	0.805 ± 0.058
MLP	None	None	N/A	None	Loss	0.859 ± 0.013	0.874 ± 0.012
MLP	DMoN	None	N/A	None	Loss	0.420 ± 0.050	0.365 ± 0.078

Table 10: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	MCC	0.876 ± 0.015	0.889 ± 0.014
GCN	DMoN	None	\checkmark	Transformer	MCC	0.875 ± 0.013	0.889 ± 0.012
GraphSAGE	None	None	\times	None	MCC	0.889 ± 0.008	0.901 ± 0.007
GraphSAGE	DMoN	None	\times	None	MCC	0.881 ± 0.015	0.894 ± 0.014
Transformer	None	L2	N/A	None	MCC	0.803 ± 0.035	0.823 ± 0.032
Transformer	DMoN	L2	N/A	None	MCC	0.783 ± 0.063	0.805 ± 0.058
MLP	None	None	N/A	None	MCC	0.862 ± 0.011	0.877 ± 0.011
MLP	DMoN	None	N/A	None	MCC	0.839 ± 0.013	0.856 ± 0.013

F.2.4 COAUTHOR CS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 11: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.742 ± 0.055	0.763 ± 0.054
GCN	NOCD	None	\times	None	Loss	0.424 ± 0.086	0.430 ± 0.089
GraphSAGE	None	None	\times	None	Loss	0.727 ± 0.069	0.748 ± 0.067
GraphSAGE	NOCD	DMoN	\checkmark	Transformer	Loss	0.296 ± 0.061	0.287 ± 0.059
Transformer	None	DMoN	N/A	None	Loss	0.002 ± 0.005	0.070 ± 0.065
Transformer	DMoN	DMoN	N/A	None	Loss	0.013 ± 0.031	0.075 ± 0.064
MLP	None	None	N/A	None	Loss	0.566 ± 0.084	0.589 ± 0.101
MLP	DMoN	DMoN	N/A	None	Loss	0.475 ± 0.079	0.497 ± 0.084

Table 12: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.742 ± 0.055	0.763 ± 0.054
GCN	DMoN	None	\checkmark	MLP	Loss	0.668 ± 0.051	0.690 ± 0.049
GraphSAGE	None	None	\times	None	Loss	0.727 ± 0.069	0.748 ± 0.067
GraphSAGE	DMoN	None	\times	None	Loss	0.637 ± 0.029	0.654 ± 0.027
Transformer	None	L2	N/A	None	MCC	0.476 ± 0.120	0.516 ± 0.130
Transformer	DMoN	L2	N/A	None	MCC	0.518 ± 0.104	0.552 ± 0.112
MLP	None	None	N/A	None	Loss	0.566 ± 0.084	0.589 ± 0.101
MLP	DMoN	None	N/A	None	Loss	0.467 ± 0.067	0.493 ± 0.077

F.2.5 AMAZON COMPUTERS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 13: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.716 ± 0.050	0.763 ± 0.045
GCN	NOCD	None	\checkmark	MLP	Loss	0.587 ± 0.069	0.626 ± 0.082
GraphSAGE	None	None	\times	None	Loss	0.740 ± 0.021	0.781 ± 0.020
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.290 ± 0.044	0.225 ± 0.086
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.087 ± 0.061
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.131 ± 0.140
MLP	None	None	N/A	None	Loss	0.547 ± 0.031	0.612 ± 0.032
MLP	DMoN	DMoN	N/A	None	Loss	0.431 ± 0.023	0.473 ± 0.029

Table 14: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	Loss	0.723 ± 0.036	0.770 ± 0.032
GCN	DMoN	DMoN	\times	None	Loss	0.709 ± 0.041	0.756 ± 0.038
GraphSAGE	None	None	\times	None	Loss	0.740 ± 0.021	0.781 ± 0.020
GraphSAGE	DMoN	None	\times	None	Loss	0.728 ± 0.022	0.772 ± 0.021
Transformer	None	L2	N/A	None	MCC	0.031 ± 0.050	0.234 ± 0.095
Transformer	DMoN	L2	N/A	None	MCC	0.048 ± 0.037	0.193 ± 0.106
MLP	None	None	N/A	None	Loss	0.547 ± 0.031	0.612 ± 0.032
MLP	DMoN	DMoN	N/A	None	Loss	0.431 ± 0.023	0.473 ± 0.029

Table 15: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	Loss	0.723 ± 0.036	0.770 ± 0.032
GCN	DMoN	DMoN	\times	None	Loss	0.709 ± 0.041	0.756 ± 0.038
GraphSAGE	None	DMoN	\times	None	MCC	0.751 ± 0.020	0.794 ± 0.018
GraphSAGE	DMoN	None	\times	None	MCC	0.732 ± 0.028	0.777 ± 0.026
Transformer	None	L2	N/A	None	MCC	0.031 ± 0.050	0.234 ± 0.095
Transformer	DMoN	L2	N/A	None	MCC	0.048 ± 0.037	0.193 ± 0.106
MLP	None	None	N/A	None	MCC	0.586 ± 0.026	0.656 ± 0.024
MLP	DMoN	None	N/A	None	MCC	0.553 ± 0.027	0.616 ± 0.026

F.2.6 AMAZON COMPUTERS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 16: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	Loss	0.521 ± 0.064	0.580 ± 0.069
GCN	NOCD	DMoN	✓	MLP	Loss	0.297 ± 0.070	0.271 ± 0.085
GraphSAGE	None	DMoN	\times	None	Loss	0.466 ± 0.105	0.522 ± 0.114
GraphSAGE	NOCD	DMoN	✓	Transformer	Loss	0.210 ± 0.075	0.162 ± 0.096
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.075 ± 0.051
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.133 ± 0.138
MLP	None	DMoN	N/A	None	Loss	0.192 ± 0.035	0.278 ± 0.048
MLP	DMoN	DMoN	N/A	None	Loss	0.189 ± 0.042	0.260 ± 0.070

Table 17: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	Loss	0.521 ± 0.064	0.580 ± 0.069
GCN	DMoN	DMoN	✓	MLP	Loss	0.507 ± 0.065	0.574 ± 0.065
GraphSAGE	None	DMoN	\times	None	Loss	0.466 ± 0.105	0.522 ± 0.114
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.453 ± 0.092	0.504 ± 0.100
Transformer	None	L2	N/A	None	MCC	-0.001 ± 0.030	0.172 ± 0.101
Transformer	NOCD	L2	N/A	None	Loss	0.106 ± 0.093	0.177 ± 0.105
MLP	None	DMoN	N/A	None	Loss	0.192 ± 0.035	0.278 ± 0.048
MLP	DMoN	DMoN	N/A	None	Loss	0.189 ± 0.042	0.260 ± 0.070

F.2.7 AMAZON PHOTO DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 18: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	Loss	0.799 ± 0.075	0.819 ± 0.081
GCN	NOCD	DMoN	✓	Transformer	Loss	0.706 ± 0.089	0.722 ± 0.105
GraphSAGE	None	DMoN	\times	None	Loss	0.816 ± 0.052	0.838 ± 0.055
GraphSAGE	NOCD	DMoN	✓	Transformer	Loss	0.314 ± 0.054	0.251 ± 0.058
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.149 ± 0.060
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.154 ± 0.072
MLP	None	None	N/A	None	Loss	0.680 ± 0.021	0.720 ± 0.020
MLP	DMoN	DMoN	N/A	None	Loss	0.585 ± 0.053	0.623 ± 0.054

Table 19: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	Loss	0.799 ± 0.075	0.819 ± 0.081
GCN	DMoN	DMoN	✓	Transformer	Loss	0.840 ± 0.031	0.862 ± 0.028
GraphSAGE	None	DMoN	\times	None	Loss	0.816 ± 0.052	0.838 ± 0.055
GraphSAGE	DMoN	DMoN	✓	Transformer	Loss	0.830 ± 0.026	0.853 ± 0.024
Transformer	None	L2	N/A	None	MCC	0.093 ± 0.032	0.198 ± 0.054
Transformer	NOCD	L2	N/A	None	Loss	0.071 ± 0.113	0.155 ± 0.110
MLP	None	None	N/A	None	Loss	0.680 ± 0.021	0.720 ± 0.020
MLP	DMoN	None	N/A	None	Loss	0.618 ± 0.048	0.652 ± 0.052

Table 20: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	MCC	0.804 ± 0.069	0.826 ± 0.071
GCN	DMoN	DMoN	✓	Transformer	Loss	0.840 ± 0.031	0.862 ± 0.028
GraphSAGE	None	None	\times	None	Loss	0.830 ± 0.059	0.850 ± 0.062
GraphSAGE	DMoN	None	✓	MLP	Loss	0.828 ± 0.030	0.851 ± 0.027
Transformer	None	L2	N/A	None	MCC	0.093 ± 0.032	0.198 ± 0.054
Transformer	DMoN	L2	N/A	None	MCC	0.085 ± 0.068	0.244 ± 0.058
MLP	None	None	N/A	None	MCC	0.698 ± 0.024	0.739 ± 0.023
MLP	DMoN	None	N/A	None	MCC	0.697 ± 0.031	0.740 ± 0.027

F.2.8 AMAZON PHOTO DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 21: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	✗	None	Loss	0.644 ± 0.088	0.679 ± 0.088
GCN	NOCD	DMoN	✓	MLP	Loss	0.374 ± 0.044	0.365 ± 0.039
GraphSAGE	None	None	✗	None	Loss	0.640 ± 0.068	0.675 ± 0.072
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.389 ± 0.107	0.402 ± 0.110
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.105 ± 0.078
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.170 ± 0.070
MLP	None	DMoN	N/A	None	Loss	0.235 ± 0.063	0.288 ± 0.062
MLP	DMoN	None	N/A	None	Loss	0.235 ± 0.049	0.301 ± 0.061

Table 22: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	✗	None	Loss	0.644 ± 0.088	0.679 ± 0.088
GCN	DMoN	DMoN	✗	None	Loss	0.700 ± 0.075	0.736 ± 0.068
GraphSAGE	None	None	✗	None	Loss	0.640 ± 0.068	0.675 ± 0.072
GraphSAGE	DMoN	None	✗	None	Loss	0.667 ± 0.068	0.702 ± 0.070
Transformer	None	L2	N/A	None	MCC	0.016 ± 0.024	0.130 ± 0.079
Transformer	NOCD	L2	N/A	None	Loss	0.155 ± 0.138	0.232 ± 0.137
MLP	None	DMoN	N/A	None	Loss	0.235 ± 0.063	0.288 ± 0.062
MLP	DMoN	None	N/A	None	Loss	0.235 ± 0.049	0.301 ± 0.061

F.2.9 COAUTHOR PHYSICS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 23: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.891 ± 0.015	0.925 ± 0.012
GCN	NOCD	DMoN	\checkmark	Transformer	Loss	0.298 ± 0.114	0.363 ± 0.138
GraphSAGE	None	None	\times	None	Loss	0.899 ± 0.011	0.931 ± 0.008
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.150 ± 0.053	0.212 ± 0.024
Transformer	None	None	N/A	None	Loss	0.032 ± 0.100	0.243 ± 0.200
Transformer	DMoN	None	N/A	None	Loss	0.034 ± 0.108	0.244 ± 0.194
MLP	None	None	N/A	None	Loss	0.820 ± 0.026	0.872 ± 0.023
MLP	DMoN	DMoN	N/A	None	Loss	0.230 ± 0.072	0.263 ± 0.076

Table 24: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.891 ± 0.015	0.925 ± 0.012
GCN	DMoN	None	\times	None	Loss	0.703 ± 0.031	0.750 ± 0.029
GraphSAGE	None	None	\times	None	Loss	0.899 ± 0.011	0.931 ± 0.008
GraphSAGE	DMoN	None	\times	None	Loss	0.565 ± 0.043	0.615 ± 0.030
Transformer	None	L2	N/A	None	MCC	0.751 ± 0.092	0.821 ± 0.074
Transformer	DMoN	L2	N/A	None	MCC	0.750 ± 0.073	0.829 ± 0.051
MLP	None	None	N/A	None	Loss	0.820 ± 0.026	0.872 ± 0.023
MLP	DMoN	None	N/A	None	Loss	0.228 ± 0.097	0.248 ± 0.047

Table 25: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	MCC	0.909 ± 0.012	0.938 ± 0.008
GCN	DMoN	None	\times	None	MCC	0.907 ± 0.008	0.936 ± 0.006
GraphSAGE	None	DMoN	\times	None	MCC	0.905 ± 0.006	0.935 ± 0.005
GraphSAGE	DMoN	None	\times	None	MCC	0.904 ± 0.007	0.934 ± 0.005
Transformer	None	L2	N/A	None	MCC	0.751 ± 0.092	0.821 ± 0.074
Transformer	DMoN	L2	N/A	None	MCC	0.750 ± 0.073	0.829 ± 0.051
MLP	None	None	N/A	None	MCC	0.819 ± 0.023	0.874 ± 0.016
MLP	NOCD	L2	N/A	None	Loss	0.744 ± 0.198	0.826 ± 0.103

F.2.10 COAUTHOR PHYSICS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 26: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.796 ± 0.058	0.859 ± 0.041
GCN	NOCD	None	\times	None	Loss	0.391 ± 0.104	0.434 ± 0.122
GraphSAGE	None	None	\times	None	Loss	0.802 ± 0.047	0.864 ± 0.034
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.213 ± 0.060	0.278 ± 0.030
Transformer	None	None	N/A	None	Loss	0.023 ± 0.081	0.159 ± 0.124
Transformer	DMoN	None	N/A	None	Loss	0.023 ± 0.074	0.199 ± 0.154
MLP	None	None	N/A	None	Loss	0.488 ± 0.060	0.604 ± 0.089
MLP	DMoN	DMoN	N/A	None	Loss	0.293 ± 0.106	0.408 ± 0.163

Table 27: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.796 ± 0.058	0.859 ± 0.041
GCN	DMoN	None	\times	None	Loss	0.487 ± 0.062	0.552 ± 0.060
GraphSAGE	None	DMoN	\times	None	MCC	0.757 ± 0.068	0.820 ± 0.064
GraphSAGE	DMoN	None	\times	None	Loss	0.209 ± 0.073	0.350 ± 0.054
Transformer	None	L2	N/A	None	Loss	0.356 ± 0.119	0.443 ± 0.164
Transformer	NOCD	L2	N/A	None	Loss	0.542 ± 0.124	0.632 ± 0.126
MLP	None	None	N/A	None	Loss	0.488 ± 0.060	0.604 ± 0.089
MLP	DMoN	None	N/A	None	Loss	0.302 ± 0.114	0.396 ± 0.145

F.2.11 ROMAN-EMPIRE DATASET, DEFAULT SPLIT WITH DEFAULT TRAIN NODES PER CLASS AND DEFAULT VALIDATION NODES.

Table 28: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.087 ± 0.005	0.191 ± 0.006
GCN	DMoN	DMoN	\times	None	Loss	0.088 ± 0.007	0.193 ± 0.007
GraphSAGE	None	None	\times	None	Loss	0.285 ± 0.150	0.355 ± 0.113
GraphSAGE	DMoN	None	\checkmark	MLP	Loss	0.318 ± 0.034	0.381 ± 0.027
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	None	DMoN	N/A	None	Loss	0.081 ± 0.102	0.183 ± 0.062
MLP	DMoN	None	N/A	None	Loss	0.061 ± 0.099	0.173 ± 0.065

Table 29: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.087 ± 0.005	0.191 ± 0.006
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.089 ± 0.002	0.199 ± 0.002
GraphSAGE	None	None	\times	None	Loss	0.285 ± 0.150	0.355 ± 0.113
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.323 ± 0.019	0.385 ± 0.014
Transformer	None	L2	N/A	None	MCC	0.012 ± 0.038	0.022 ± 0.068
Transformer	NOCD	L2	N/A	None	MCC	0.013 ± 0.039	0.030 ± 0.065
MLP	None	None	N/A	None	MCC	0.226 ± 0.037	0.251 ± 0.038
MLP	DMoN	DMoN	N/A	None	MCC	0.245 ± 0.043	0.262 ± 0.051

Table 30: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	Loss	0.087 ± 0.005	0.192 ± 0.005
GCN	DMoN	None	\times	None	Loss	0.089 ± 0.005	0.193 ± 0.004
GraphSAGE	None	None	\times	None	Loss	0.285 ± 0.150	0.355 ± 0.113
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.323 ± 0.019	0.385 ± 0.014
Transformer	None	L2	N/A	None	MCC	0.012 ± 0.038	0.022 ± 0.068
Transformer	NOCD	L2	N/A	None	MCC	0.013 ± 0.039	0.030 ± 0.065
MLP	None	None	N/A	None	MCC	0.226 ± 0.037	0.251 ± 0.038
MLP	DMoN	DMoN	N/A	None	MCC	0.245 ± 0.043	0.262 ± 0.051

F.2.12 ROMAN-EMPIRE DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 31: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.089 ± 0.003	0.195 ± 0.005
GCN	DMoN	DMoN	\times	None	Loss	0.090 ± 0.006	0.195 ± 0.006
GraphSAGE	None	DMoN	\times	None	Loss	0.214 ± 0.184	0.301 ± 0.139
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.327 ± 0.019	0.388 ± 0.014
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
MLP	None	DMoN	N/A	None	Loss	0.109 ± 0.098	0.196 ± 0.058
MLP	DMoN	None	N/A	None	Loss	0.112 ± 0.081	0.202 ± 0.049

Table 32: Comparing (semi-)supervised node classification with semi-supervised graph clustering. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	DMoN	\times	None	Loss	0.090 ± 0.005	0.196 ± 0.006
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.089 ± 0.003	0.198 ± 0.002
GraphSAGE	None	DMoN	\times	None	Loss	0.214 ± 0.184	0.301 ± 0.139
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.327 ± 0.019	0.388 ± 0.014
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	NOCD	DMoN	N/A	None	MCC	0.004 ± 0.012	0.013 ± 0.041
MLP	None	DMoN	N/A	None	MCC	0.222 ± 0.034	0.240 ± 0.035
MLP	DMoN	DMoN	N/A	None	MCC	0.234 ± 0.043	0.261 ± 0.041

F.3 NO ATTRIBUTE RECONSTRUCTION VS ATTRIBUTE RECONSTRUCTION

In this section we present results an ablation of attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE.

F.3.1 AMAZON-RATINGS DATASET, DEFAULT SPLIT WITH DEFAULT TRAIN NODES PER CLASS AND DEFAULT VALIDATION NODES.

Table 33: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	Loss	0.001 ± 0.002	0.368 ± 0.000
GCN	DMoN	None	\checkmark	MLP	Loss	0.001 ± 0.006	0.368 ± 0.000
GraphSAGE	DMoN	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	None	\checkmark	MLP	Loss	0.000 ± 0.000	0.368 ± 0.000

Table 34: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	None	\times	None	Loss	0.028 ± 0.009	0.331 ± 0.020
GCN	NOCD	DMoN	\checkmark	MLP	Loss	0.030 ± 0.007	0.336 ± 0.025
GraphSAGE	SBM _{NN}	L2	\times	None	MCC	0.000 ± 0.001	0.023 ± 0.074
GraphSAGE	Neuromap	DMoN	\checkmark	Transformer	MCC	-0.000 ± 0.000	0.023 ± 0.073

Table 35: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	None	\times	None	Loss	0.028 ± 0.009	0.331 ± 0.020
GCN	NOCD	None	\checkmark	MLP	Loss	0.032 ± 0.008	0.337 ± 0.020
GraphSAGE	SBM _{NN}	L2	\times	None	MCC	0.000 ± 0.001	0.023 ± 0.074
GraphSAGE	Neuromap	DMoN	\checkmark	MLP	MCC	-0.001 ± 0.002	0.005 ± 0.015

F.3.2 AMAZON-RATINGS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 36: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	Loss	-0.003 \pm 0.004	0.368 \pm 0.000
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.004 \pm 0.008	0.368 \pm 0.000
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.000 \pm 0.000	0.368 \pm 0.000
GraphSAGE	DMoN	None	\checkmark	MLP	Loss	0.000 \pm 0.000	0.368 \pm 0.000

Table 37: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	None	\times	None	Loss	0.030 \pm 0.008	0.339 \pm 0.005
GCN	NOCD	None	\checkmark	MLP	Loss	0.031 \pm 0.010	0.344 \pm 0.006
GraphSAGE	SBM _{NN}	None	\times	None	MCC	-0.001 \pm 0.002	0.026 \pm 0.081
GraphSAGE	Neuromap	None	\checkmark	MLP	MCC	0.000 \pm 0.000	0.037 \pm 0.116

Table 38: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	DMoN	\times	None	Loss	0.028 \pm 0.008	0.338 \pm 0.006
GCN	NOCD	None	\checkmark	MLP	Loss	0.031 \pm 0.010	0.344 \pm 0.006
GraphSAGE	NOCD	L2	\times	None	MCC	0.001 \pm 0.005	0.031 \pm 0.085
GraphSAGE	Neuromap	None	\checkmark	MLP	MCC	0.000 \pm 0.000	0.037 \pm 0.116

F.3.3 COAUTHOR CS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 39: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	DMoN	✗	None	Loss	0.252 ± 0.026	0.183 ± 0.025
GCN	NOCD	DMoN	✓	Transformer	Loss	0.319 ± 0.129	0.262 ± 0.159
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.217 ± 0.027	0.153 ± 0.022
GraphSAGE	NOCD	DMoN	✓	Transformer	Loss	0.215 ± 0.060	0.157 ± 0.054

Table 40: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	✗	None	Loss	0.867 ± 0.013	0.881 ± 0.012
GCN	DMoN	None	✓	Transformer	Loss	0.875 ± 0.012	0.888 ± 0.011
GraphSAGE	DMoN	None	✗	None	Loss	0.812 ± 0.035	0.825 ± 0.035
GraphSAGE	DMoN	DMoN	✓	MLP	Loss	0.768 ± 0.020	0.782 ± 0.021

Table 41: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	✗	None	Loss	0.867 ± 0.013	0.881 ± 0.012
GCN	DMoN	None	✓	Transformer	MCC	0.875 ± 0.013	0.889 ± 0.012
GraphSAGE	DMoN	None	✗	None	MCC	0.881 ± 0.015	0.894 ± 0.014
GraphSAGE	DMoN	DMoN	✓	Transformer	MCC	0.877 ± 0.013	0.890 ± 0.012

F.3.4 COAUTHOR CS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 42: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	None	\times	None	Loss	0.424 ± 0.086	0.430 ± 0.089
GCN	NOCD	None	\checkmark	Transformer	Loss	0.461 ± 0.080	0.468 ± 0.082
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.327 ± 0.042	0.307 ± 0.048
GraphSAGE	NOCD	DMoN	\checkmark	Transformer	Loss	0.296 ± 0.061	0.287 ± 0.059

Table 43: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.707 ± 0.066	0.729 ± 0.061
GCN	DMoN	None	\checkmark	MLP	Loss	0.668 ± 0.051	0.690 ± 0.049
GraphSAGE	DMoN	None	\times	None	Loss	0.637 ± 0.029	0.654 ± 0.027
GraphSAGE	DMoN	None	\checkmark	Transformer	Loss	0.675 ± 0.042	0.695 ± 0.043

Table 44: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	MCC	0.756 ± 0.045	0.778 ± 0.042
GCN	DMoN	DMoN	\checkmark	Transformer	MCC	0.754 ± 0.040	0.776 ± 0.040
GraphSAGE	DMoN	DMoN	\times	None	MCC	0.764 ± 0.037	0.784 ± 0.037
GraphSAGE	DMoN	DMoN	\checkmark	Transformer	MCC	0.735 ± 0.045	0.757 ± 0.043

F.3.5 AMAZON COMPUTERS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 45: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.709 ± 0.041	0.756 ± 0.038
GCN	NOCD	None	\checkmark	MLP	Loss	0.587 ± 0.069	0.626 ± 0.082
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.290 ± 0.044	0.225 ± 0.086
GraphSAGE	NOCD	DMoN	\checkmark	Transformer	Loss	0.250 ± 0.042	0.175 ± 0.039

Table 46: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.709 ± 0.041	0.756 ± 0.038
GCN	DMoN	None	\checkmark	Transformer	Loss	0.706 ± 0.037	0.754 ± 0.030
GraphSAGE	DMoN	None	\times	None	Loss	0.728 ± 0.022	0.772 ± 0.021
GraphSAGE	DMoN	DMoN	\checkmark	Transformer	Loss	0.714 ± 0.024	0.758 ± 0.024

Table 47: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.709 ± 0.041	0.756 ± 0.038
GCN	DMoN	None	\checkmark	Transformer	Loss	0.706 ± 0.037	0.754 ± 0.030
GraphSAGE	DMoN	None	\times	None	MCC	0.732 ± 0.028	0.777 ± 0.026
GraphSAGE	DMoN	DMoN	\checkmark	Transformer	Loss	0.714 ± 0.024	0.758 ± 0.024

F.3.6 AMAZON COMPUTERS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 48: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	None	\times	None	Loss	0.328 ± 0.094	0.302 ± 0.147
GCN	NOCD	DMoN	\checkmark	MLP	Loss	0.297 ± 0.070	0.271 ± 0.085
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.215 ± 0.102	0.190 ± 0.124
GraphSAGE	NOCD	DMoN	\checkmark	Transformer	Loss	0.210 ± 0.075	0.162 ± 0.096

Table 49: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	Loss	0.476 ± 0.086	0.529 ± 0.097
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.507 ± 0.065	0.574 ± 0.065
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.453 ± 0.092	0.504 ± 0.100
GraphSAGE	DMoN	None	\checkmark	Transformer	Loss	0.389 ± 0.061	0.446 ± 0.075

Table 50: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	Loss	0.476 ± 0.086	0.529 ± 0.097
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.507 ± 0.065	0.574 ± 0.065
GraphSAGE	DMoN	None	\times	None	MCC	0.484 ± 0.110	0.550 ± 0.121
GraphSAGE	NOCD	DMoN	\checkmark	Transformer	MCC	0.391 ± 0.122	0.455 ± 0.139

F.3.7 AMAZON PHOTO DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 51: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.812 ± 0.058	0.833 ± 0.062
GCN	NOCD	DMoN	\checkmark	Transformer	Loss	0.706 ± 0.089	0.722 ± 0.105
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.304 ± 0.062	0.241 ± 0.067
GraphSAGE	NOCD	DMoN	\checkmark	Transformer	Loss	0.314 ± 0.054	0.251 ± 0.058

Table 52: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.812 ± 0.058	0.833 ± 0.062
GCN	DMoN	DMoN	\checkmark	Transformer	Loss	0.840 ± 0.031	0.862 ± 0.028
GraphSAGE	DMoN	None	\times	None	Loss	0.794 ± 0.062	0.815 ± 0.067
GraphSAGE	DMoN	DMoN	\checkmark	Transformer	Loss	0.830 ± 0.026	0.853 ± 0.024

Table 53: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	MCC	0.790 ± 0.088	0.809 ± 0.093
GCN	DMoN	DMoN	\checkmark	Transformer	Loss	0.840 ± 0.031	0.862 ± 0.028
GraphSAGE	DMoN	DMoN	\times	None	MCC	0.794 ± 0.065	0.815 ± 0.070
GraphSAGE	DMoN	None	\checkmark	MLP	Loss	0.828 ± 0.030	0.851 ± 0.027

F.3.8 AMAZON PHOTO DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 54: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	DMoN	✗	None	Loss	0.333 ± 0.112	0.311 ± 0.127
GCN	NOCD	DMoN	✓	MLP	Loss	0.374 ± 0.044	0.365 ± 0.039
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.252 ± 0.082	0.205 ± 0.088
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.389 ± 0.107	0.402 ± 0.110

Table 55: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	✗	None	Loss	0.700 ± 0.075	0.736 ± 0.068
GCN	DMoN	DMoN	✓	Transformer	Loss	0.657 ± 0.068	0.698 ± 0.060
GraphSAGE	DMoN	None	✗	None	Loss	0.667 ± 0.068	0.702 ± 0.070
GraphSAGE	DMoN	None	✓	Transformer	Loss	0.594 ± 0.110	0.639 ± 0.099

Table 56: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	✗	None	MCC	0.618 ± 0.093	0.659 ± 0.098
GCN	DMoN	DMoN	✓	Transformer	Loss	0.657 ± 0.068	0.698 ± 0.060
GraphSAGE	DMoN	DMoN	✗	None	MCC	0.571 ± 0.102	0.618 ± 0.105
GraphSAGE	DMoN	None	✓	Transformer	Loss	0.594 ± 0.110	0.639 ± 0.099

F.3.9 COAUTHOR PHYSICS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 57: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	DMoN	\times	None	Loss	0.188 ± 0.068	0.249 ± 0.064
GCN	NOCD	DMoN	\checkmark	Transformer	Loss	0.298 ± 0.114	0.363 ± 0.138
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.150 ± 0.053	0.212 ± 0.024
GraphSAGE	NOCD	DMoN	\checkmark	Transformer	Loss	0.189 ± 0.064	0.243 ± 0.047

Table 58: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	Loss	0.703 ± 0.031	0.750 ± 0.029
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.686 ± 0.022	0.733 ± 0.022
GraphSAGE	DMoN	None	\times	None	Loss	0.565 ± 0.043	0.615 ± 0.030
GraphSAGE	DMoN	None	\checkmark	Transformer	Loss	0.646 ± 0.069	0.689 ± 0.068

Table 59: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	MCC	0.907 ± 0.008	0.936 ± 0.006
GCN	DMoN	DMoN	\checkmark	Transformer	MCC	0.902 ± 0.014	0.933 ± 0.010
GraphSAGE	DMoN	None	\times	None	MCC	0.904 ± 0.007	0.934 ± 0.005
GraphSAGE	DMoN	DMoN	\checkmark	MLP	MCC	0.899 ± 0.006	0.931 ± 0.004

F.3.10 COAUTHOR PHYSICS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 60: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	NOCD	None	\times	None	Loss	0.391 ± 0.104	0.434 ± 0.122
GCN	NOCD	None	\checkmark	Transformer	Loss	0.366 ± 0.156	0.452 ± 0.143
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.213 ± 0.060	0.278 ± 0.030
GraphSAGE	NOCD	DMoN	\checkmark	Transformer	Loss	0.261 ± 0.117	0.327 ± 0.100

Table 61: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	Loss	0.487 ± 0.062	0.552 ± 0.060
GCN	DMoN	None	\checkmark	Transformer	Loss	0.636 ± 0.081	0.704 ± 0.074
GraphSAGE	DMoN	None	\times	None	Loss	0.209 ± 0.073	0.350 ± 0.054
GraphSAGE	DMoN	None	\checkmark	Transformer	Loss	0.450 ± 0.095	0.541 ± 0.069

Table 62: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	MCC	0.830 ± 0.042	0.882 ± 0.030
GCN	DMoN	None	\checkmark	Transformer	MCC	0.761 ± 0.074	0.832 ± 0.057
GraphSAGE	DMoN	None	\times	None	MCC	0.792 ± 0.047	0.855 ± 0.034
GraphSAGE	DMoN	DMoN	\checkmark	Transformer	MCC	0.791 ± 0.040	0.853 ± 0.026

F.3.11 ROMAN-EMPIRE DATASET, DEFAULT SPLIT WITH DEFAULT TRAIN NODES PER CLASS AND DEFAULT VALIDATION NODES.

Table 63: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.088 ± 0.007	0.193 ± 0.007
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.089 ± 0.002	0.199 ± 0.002
GraphSAGE	DMoN	None	\times	None	Loss	0.177 ± 0.186	0.274 ± 0.141
GraphSAGE	DMoN	None	\checkmark	MLP	Loss	0.318 ± 0.034	0.381 ± 0.027

Table 64: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.088 ± 0.007	0.193 ± 0.007
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.089 ± 0.002	0.199 ± 0.002
GraphSAGE	NOCD	None	\times	None	Loss	0.213 ± 0.046	0.224 ± 0.047
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.323 ± 0.019	0.385 ± 0.014

Table 65: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	None	\times	None	Loss	0.089 ± 0.005	0.193 ± 0.004
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.089 ± 0.002	0.199 ± 0.002
GraphSAGE	NOCD	None	\times	None	Loss	0.213 ± 0.046	0.224 ± 0.047
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.323 ± 0.019	0.385 ± 0.014

F.3.12 ROMAN-EMPIRE DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 66: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.090 ± 0.006	0.195 ± 0.006
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.089 ± 0.003	0.198 ± 0.002
GraphSAGE	DMoN	None	\times	None	Loss	0.213 ± 0.183	0.301 ± 0.139
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.327 ± 0.019	0.388 ± 0.014

Table 67: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.090 ± 0.006	0.195 ± 0.006
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.089 ± 0.003	0.198 ± 0.002
GraphSAGE	DMoN	None	\times	None	Loss	0.213 ± 0.183	0.301 ± 0.139
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.327 ± 0.019	0.388 ± 0.014

Table 68: Ablating attribute reconstruction for semi-supervised graph clustering with graph neural networks GCN and GraphSAGE. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	DMoN	DMoN	\times	None	Loss	0.090 ± 0.006	0.195 ± 0.006
GCN	DMoN	None	\checkmark	MLP	Loss	0.089 ± 0.005	0.198 ± 0.004
GraphSAGE	DMoN	None	\times	None	Loss	0.213 ± 0.183	0.301 ± 0.139
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.327 ± 0.019	0.388 ± 0.014

F.4 COMPARISON OF POOLING METHODS AND REGULARISATION

In this section we present results comparing graph clustering and regularization objectives and the effect of ablating regularization objectives. We compare clustering with regularization, clustering without regularization, and regularization without clustering for each neural network.

Our results show that unsupervised graph clustering objectives with regularization outperform regularization alone, ablating the effect of regularization.

F.4.1 AMAZON-RATINGS DATASET, DEFAULT SPLIT WITH DEFAULT TRAIN NODES PER CLASS AND DEFAULT VALIDATION NODES.

Table 69: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.001 ± 0.004	0.368 ± 0.000
GCN	None	DMoN	\times	None	Loss	-0.001 ± 0.002	0.368 ± 0.000
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	DMoN	None	\checkmark	MLP	Loss	0.001 ± 0.006	0.368 ± 0.000
GCN	DMoN	DMoN	\checkmark	MLP	Loss	-0.000 ± 0.004	0.368 ± 0.000
GCN	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	NOCD	None	\times	None	Loss	0.028 ± 0.009	0.331 ± 0.020
GCN	NOCD	DMoN	\times	None	Loss	0.029 ± 0.010	0.333 ± 0.020
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	SBM _{NN}	None	\times	None	Loss	0.020 ± 0.014	0.346 ± 0.010
GCN	SBM _{NN}	DMoN	\times	None	Loss	0.021 ± 0.014	0.346 ± 0.006
GCN	SBM _{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	None	\checkmark	MLP	Loss	0.000 ± 0.000	0.354 ± 0.043
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	Neuromap	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	Neuromap	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	None	\checkmark	MLP	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.354 ± 0.043
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.308 ± 0.052
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.301 ± 0.060
Transformer	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.270 ± 0.122
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.241 ± 0.143
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000

Continued on next page

Table 69: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.284 ± 0.046
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.314 ± 0.058
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.301 ± 0.060
MLP	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.304 ± 0.056
MLP	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.306 ± 0.093
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.314 ± 0.058
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.288 ± 0.042
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.274 ± 0.035
MLP	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.304 ± 0.056

Table 70: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.001 ± 0.004	0.368 ± 0.000
GCN	None	DMoN	\times	None	MCC	0.001 ± 0.003	0.027 ± 0.085
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	DMoN	None	\checkmark	Transformer	MCC	0.010 ± 0.019	0.119 ± 0.130
GCN	DMoN	DMoN	\times	None	MCC	0.000 ± 0.000	0.000 ± 0.000
GCN	DMoN	L2	\times	None	MCC	-0.003 ± 0.034	0.269 ± 0.074
GCN	NOCD	None	\times	None	Loss	0.028 ± 0.009	0.331 ± 0.020
GCN	NOCD	DMoN	\times	None	Loss	0.029 ± 0.010	0.333 ± 0.020
GCN	NOCD	L2	\times	None	MCC	0.003 ± 0.005	0.108 ± 0.173
GCN	Neuromap	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	SBM_{NN}	None	\checkmark	Transformer	Loss	0.018 ± 0.017	0.355 ± 0.014
GCN	SBM_{NN}	DMoN	\times	None	Loss	0.021 ± 0.014	0.346 ± 0.006
GCN	SBM_{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000

Continued on next page

Table 70: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	Neuromap	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	Neuromap	DMoN	✓	Transformer	MCC	-0.000 ± 0.000	0.023 ± 0.073
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.000 ± 0.000	0.348 ± 0.042
GraphSAGE	SBM _{NN}	L2	✗	None	MCC	0.000 ± 0.001	0.023 ± 0.074
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.354 ± 0.043
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.308 ± 0.052
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.301 ± 0.060
Transformer	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.270 ± 0.122
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.241 ± 0.143
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.284 ± 0.046
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.314 ± 0.058
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	L2	N/A	None	MCC	0.002 ± 0.005	0.077 ± 0.124
MLP	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	L2	N/A	None	MCC	0.000 ± 0.004	0.036 ± 0.075
MLP	NOCD	None	N/A	None	MCC	0.000 ± 0.000	0.023 ± 0.074
MLP	NOCD	DMoN	N/A	None	MCC	-0.002 ± 0.005	0.061 ± 0.132
MLP	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.306 ± 0.093
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.314 ± 0.058
MLP	Neuromap	L2	N/A	None	MCC	0.001 ± 0.004	0.090 ± 0.140
MLP	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.288 ± 0.042
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.274 ± 0.035
MLP	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.304 ± 0.056

Table 71: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	MCC	0.000 ± 0.000	0.000 ± 0.000
GCN	None	DMoN	✗	None	MCC	0.001 ± 0.003	0.027 ± 0.085
GCN	None	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	DMoN	None	✓	Transformer	MCC	0.010 ± 0.019	0.119 ± 0.130
GCN	DMoN	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	DMoN	L2	✗	None	MCC	-0.003 ± 0.034	0.269 ± 0.074
GCN	NOCD	None	✗	None	Loss	0.028 ± 0.009	0.331 ± 0.020
GCN	NOCD	DMoN	✗	None	Loss	0.029 ± 0.010	0.333 ± 0.020
GCN	NOCD	L2	✗	None	MCC	0.003 ± 0.005	0.108 ± 0.173
GCN	Neuromap	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	DMoN	✓	Transformer	MCC	-0.002 ± 0.023	0.154 ± 0.164
GCN	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	SBM _{NN}	None	✗	None	Loss	0.020 ± 0.014	0.346 ± 0.010
GCN	SBM _{NN}	DMoN	✗	None	Loss	0.021 ± 0.014	0.346 ± 0.006
GCN	SBM _{NN}	L2	✗	None	MCC	-0.000 ± 0.007	0.050 ± 0.106
GraphSAGE	None	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	L2	✗	None	MCC	-0.000 ± 0.001	0.027 ± 0.085
GraphSAGE	DMoN	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	DMoN	✗	None	MCC	0.000 ± 0.000	0.027 ± 0.085
GraphSAGE	DMoN	L2	✓	MLP	MCC	0.000 ± 0.001	0.073 ± 0.154
GraphSAGE	NOCD	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	L2	✓	Transformer	MCC	-0.000 ± 0.000	0.030 ± 0.093
GraphSAGE	Neuromap	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	-0.001 ± 0.002	0.005 ± 0.015
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.000 ± 0.000	0.348 ± 0.042
GraphSAGE	SBM _{NN}	L2	✗	None	MCC	0.000 ± 0.001	0.023 ± 0.074
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.354 ± 0.043
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.308 ± 0.052
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.301 ± 0.060
Transformer	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.270 ± 0.122
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.241 ± 0.143
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000

Continued on next page

Table 71: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.284 ± 0.046
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.314 ± 0.058
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	L2	N/A	None	MCC	0.002 ± 0.005	0.077 ± 0.124
MLP	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	L2	N/A	None	MCC	0.000 ± 0.004	0.036 ± 0.075
MLP	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.301 ± 0.060
MLP	NOCD	DMoN	N/A	None	MCC	-0.002 ± 0.005	0.061 ± 0.132
MLP	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.306 ± 0.093
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.314 ± 0.058
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.288 ± 0.042
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.274 ± 0.035
MLP	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.304 ± 0.056

F.4.2 AMAZON-RATINGS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 72: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	-0.001 ± 0.001	0.368 ± 0.000
GCN	None	DMoN	\times	None	Loss	-0.002 ± 0.006	0.368 ± 0.000
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	DMoN	None	\checkmark	MLP	Loss	0.000 ± 0.004	0.368 ± 0.000
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.004 ± 0.008	0.368 ± 0.000
GCN	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	NOCD	None	\times	None	Loss	0.030 ± 0.008	0.339 ± 0.005
GCN	NOCD	DMoN	\times	None	Loss	0.028 ± 0.008	0.338 ± 0.006
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	SBM_{NN}	None	\times	None	Loss	0.021 ± 0.012	0.352 ± 0.007
GCN	SBM_{NN}	DMoN	\times	None	Loss	0.017 ± 0.015	0.349 ± 0.009
GCN	SBM_{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000

Continued on next page

Table 72: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	DMoN	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	Neuromap	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	Neuromap	DMoN	✗	None	Loss	0.000 ± 0.000	0.358 ± 0.032
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	None	✓	Transformer	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	DMoN	✓	Transformer	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	L2	✓	MLP	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.344 ± 0.051
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.321 ± 0.062
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.304 ± 0.056
Transformer	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.209 ± 0.150
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.233 ± 0.142
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.334 ± 0.055
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.321 ± 0.062
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.318 ± 0.053
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.348 ± 0.042
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
MLP	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.314 ± 0.058
MLP	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.298 ± 0.048
MLP	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.242 ± 0.119
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.296 ± 0.091
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.294 ± 0.052
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.288 ± 0.042
MLP	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.311 ± 0.062

Table 73: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	MCC	0.000 ± 0.000	0.000 ± 0.000
GCN	None	DMoN	✗	None	MCC	-0.001 ± 0.004	0.028 ± 0.088
GCN	None	L2	✗	None	MCC	-0.003 ± 0.030	0.252 ± 0.093
GCN	DMoN	None	✗	None	MCC	0.000 ± 0.001	0.027 ± 0.085
GCN	DMoN	DMoN	✓	Transformer	MCC	0.013 ± 0.026	0.246 ± 0.122
GCN	DMoN	L2	✗	None	MCC	0.002 ± 0.030	0.264 ± 0.084
GCN	NOCD	None	✓	MLP	Loss	0.031 ± 0.010	0.344 ± 0.006
GCN	NOCD	DMoN	✗	None	Loss	0.028 ± 0.008	0.338 ± 0.006
GCN	NOCD	L2	✓	MLP	MCC	-0.001 ± 0.008	0.064 ± 0.136
GCN	Neuromap	None	✓	Transformer	MCC	0.002 ± 0.014	0.137 ± 0.162
GCN	Neuromap	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	L2	✗	None	MCC	0.022 ± 0.026	0.268 ± 0.064
GCN	SBM _{NN}	None	✗	None	Loss	0.016 ± 0.022	0.348 ± 0.010
GCN	SBM _{NN}	DMoN	✓	MLP	Loss	0.021 ± 0.020	0.345 ± 0.006
GCN	SBM _{NN}	L2	✓	MLP	MCC	0.005 ± 0.015	0.148 ± 0.159
GraphSAGE	None	None	✗	None	MCC	0.001 ± 0.004	0.023 ± 0.073
GraphSAGE	None	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	L2	✗	None	MCC	0.002 ± 0.004	0.032 ± 0.076
GraphSAGE	DMoN	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	DMoN	✓	MLP	MCC	-0.001 ± 0.003	0.051 ± 0.107
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	L2	✗	None	MCC	0.001 ± 0.005	0.031 ± 0.085
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.000 ± 0.000	0.037 ± 0.116
GraphSAGE	Neuromap	DMoN	✗	None	Loss	0.000 ± 0.000	0.358 ± 0.032
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	SBM _{NN}	None	✗	None	MCC	-0.001 ± 0.002	0.026 ± 0.081
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	-0.000 ± 0.000	0.334 ± 0.055
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.344 ± 0.051
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.321 ± 0.062
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.304 ± 0.056
Transformer	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.209 ± 0.150
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.233 ± 0.142
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.334 ± 0.055

Continued on next page

Table 73: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.321 ± 0.062
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.318 ± 0.053
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.348 ± 0.042
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
MLP	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.314 ± 0.058
MLP	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.298 ± 0.048
MLP	NOCD	L2	N/A	None	MCC	0.000 ± 0.001	0.063 ± 0.134
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.242 ± 0.119
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.296 ± 0.091
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.294 ± 0.052
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.288 ± 0.042
MLP	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.311 ± 0.062

Table 74: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	MCC	0.000 ± 0.000	0.000 ± 0.000
GCN	None	DMoN	\times	None	Loss	-0.002 ± 0.006	0.368 ± 0.000
GCN	None	L2	\times	None	MCC	-0.003 ± 0.030	0.252 ± 0.093
GCN	DMoN	None	\times	None	MCC	0.000 ± 0.001	0.027 ± 0.085
GCN	DMoN	DMoN	\checkmark	Transformer	MCC	0.013 ± 0.026	0.246 ± 0.122
GCN	DMoN	L2	\times	None	MCC	0.002 ± 0.030	0.264 ± 0.084
GCN	NOCD	None	\checkmark	MLP	Loss	0.031 ± 0.010	0.344 ± 0.006
GCN	NOCD	DMoN	\checkmark	MLP	Loss	0.032 ± 0.008	0.331 ± 0.029
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	None	\checkmark	Transformer	MCC	0.002 ± 0.014	0.137 ± 0.162
GCN	Neuromap	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GCN	Neuromap	L2	\times	None	MCC	0.022 ± 0.026	0.268 ± 0.064
GCN	SBM_{NN}	None	\times	None	Loss	0.021 ± 0.012	0.352 ± 0.007
GCN	SBM_{NN}	DMoN	\times	None	Loss	0.017 ± 0.015	0.349 ± 0.009
GCN	SBM_{NN}	L2	\times	None	MCC	0.006 ± 0.014	0.064 ± 0.136
GraphSAGE	None	None	\times	None	MCC	0.001 ± 0.004	0.023 ± 0.073
GraphSAGE	None	DMoN	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	None	L2	\times	None	MCC	0.002 ± 0.004	0.032 ± 0.076
GraphSAGE	DMoN	None	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	DMoN	DMoN	\times	None	MCC	-0.002 ± 0.005	0.036 ± 0.114
GraphSAGE	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.368 ± 0.000

Continued on next page

Table 74: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
GraphSAGE	NOCD	L2	✗	None	MCC	0.001 ± 0.005	0.031 ± 0.085
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.000 ± 0.000	0.037 ± 0.116
GraphSAGE	Neuromap	DMoN	✗	None	Loss	0.000 ± 0.000	0.358 ± 0.032
GraphSAGE	Neuromap	L2	✗	None	MCC	0.000 ± 0.001	0.027 ± 0.084
GraphSAGE	SBM _{NN}	None	✗	None	MCC	-0.001 ± 0.002	0.026 ± 0.081
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	-0.000 ± 0.000	0.334 ± 0.055
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.344 ± 0.051
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.321 ± 0.062
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.304 ± 0.056
Transformer	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.209 ± 0.150
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.233 ± 0.142
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.334 ± 0.055
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.321 ± 0.062
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.318 ± 0.053
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.348 ± 0.042
MLP	None	None	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.358 ± 0.032
MLP	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.314 ± 0.058
MLP	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.298 ± 0.048
MLP	NOCD	L2	N/A	None	MCC	0.000 ± 0.001	0.063 ± 0.134
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.242 ± 0.119
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.296 ± 0.091
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.368 ± 0.000
MLP	SBM _{NN}	None	N/A	None	MCC	-0.001 ± 0.003	0.025 ± 0.079
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.288 ± 0.042
MLP	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.311 ± 0.062

F.4.3 COAUTHOR CS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 75: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	Loss	0.875 ± 0.014	0.889 ± 0.013
GCN	None	DMoN	✗	None	Loss	0.854 ± 0.020	0.869 ± 0.019
GCN	None	L2	✗	None	Loss	0.000 ± 0.000	0.038 ± 0.000
GCN	DMoN	None	✗	None	Loss	0.859 ± 0.011	0.874 ± 0.011
GCN	DMoN	DMoN	✗	None	Loss	0.867 ± 0.013	0.881 ± 0.012
GCN	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.038 ± 0.000
GCN	NOCD	None	✓	Transformer	Loss	0.422 ± 0.065	0.386 ± 0.081
GCN	NOCD	DMoN	✓	Transformer	Loss	0.319 ± 0.129	0.262 ± 0.159
GCN	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.082 ± 0.059
GCN	Neuromap	None	✓	Transformer	Loss	0.005 ± 0.011	0.070 ± 0.038
GCN	Neuromap	DMoN	✓	Transformer	Loss	0.011 ± 0.035	0.075 ± 0.064
GCN	Neuromap	L2	✓	MLP	Loss	0.000 ± 0.000	0.038 ± 0.000
GCN	SBM _{NN}	None	✓	Transformer	Loss	0.270 ± 0.040	0.207 ± 0.023
GCN	SBM _{NN}	DMoN	✓	Transformer	Loss	0.269 ± 0.031	0.205 ± 0.030
GCN	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.057 ± 0.042
GraphSAGE	None	None	✗	None	Loss	0.887 ± 0.007	0.899 ± 0.006
GraphSAGE	None	DMoN	✗	None	Loss	0.810 ± 0.024	0.824 ± 0.025
GraphSAGE	None	L2	✗	None	Loss	0.000 ± 0.000	0.051 ± 0.033
GraphSAGE	DMoN	None	✗	None	Loss	0.812 ± 0.035	0.825 ± 0.035
GraphSAGE	DMoN	DMoN	✓	MLP	Loss	0.768 ± 0.020	0.782 ± 0.021
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.083 ± 0.083
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.566 ± 0.099	0.576 ± 0.110
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.217 ± 0.027	0.153 ± 0.022
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.074 ± 0.065
GraphSAGE	Neuromap	None	✓	MLP	Loss	0.005 ± 0.017	0.008 ± 0.007
GraphSAGE	Neuromap	DMoN	✗	None	Loss	0.000 ± 0.000	0.035 ± 0.035
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.069 ± 0.065
GraphSAGE	SBM _{NN}	None	✓	MLP	Loss	0.159 ± 0.058	0.181 ± 0.046
GraphSAGE	SBM _{NN}	DMoN	✓	MLP	Loss	0.135 ± 0.057	0.156 ± 0.046
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.076 ± 0.060
Transformer	None	None	N/A	None	Loss	0.038 ± 0.042	0.100 ± 0.089
Transformer	None	DMoN	N/A	None	Loss	0.015 ± 0.030	0.053 ± 0.048
Transformer	None	L2	N/A	None	Loss	0.005 ± 0.009	0.053 ± 0.048
Transformer	DMoN	None	N/A	None	Loss	0.016 ± 0.034	0.065 ± 0.064
Transformer	DMoN	DMoN	N/A	None	Loss	0.014 ± 0.025	0.027 ± 0.017
Transformer	DMoN	L2	N/A	None	Loss	0.004 ± 0.011	0.040 ± 0.028
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.050 ± 0.026
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.036 ± 0.006
Transformer	NOCD	L2	N/A	None	Loss	0.089 ± 0.185	0.110 ± 0.189
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.049 ± 0.037
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.093 ± 0.082
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.053 ± 0.038

Continued on next page

Table 75: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.037 ± 0.005
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.043 ± 0.013
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.081 ± 0.072	0.107 ± 0.068
MLP	None	None	N/A	None	Loss	0.859 ± 0.013	0.874 ± 0.012
MLP	None	DMoN	N/A	None	Loss	0.348 ± 0.033	0.284 ± 0.055
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.063 ± 0.069
MLP	DMoN	None	N/A	None	Loss	0.420 ± 0.050	0.365 ± 0.078
MLP	DMoN	DMoN	N/A	None	Loss	0.423 ± 0.041	0.372 ± 0.062
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.090 ± 0.062
MLP	NOCD	None	N/A	None	Loss	0.144 ± 0.118	0.168 ± 0.115
MLP	NOCD	DMoN	N/A	None	Loss	0.139 ± 0.084	0.137 ± 0.065
MLP	NOCD	L2	N/A	None	Loss	0.647 ± 0.077	0.670 ± 0.079
MLP	Neuromap	None	N/A	None	Loss	0.018 ± 0.022	0.063 ± 0.071
MLP	Neuromap	DMoN	N/A	None	Loss	0.004 ± 0.013	0.091 ± 0.082
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.066 ± 0.062
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.038 ± 0.000
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.009 ± 0.020	0.039 ± 0.002
MLP	SBM_{NN}	L2	N/A	None	Loss	0.012 ± 0.035	0.069 ± 0.069

Table 76: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.875 ± 0.014	0.889 ± 0.013
GCN	None	DMoN	\times	None	Loss	0.854 ± 0.020	0.869 ± 0.019
GCN	None	L2	\times	None	MCC	0.120 ± 0.094	0.192 ± 0.079
GCN	DMoN	None	\times	None	Loss	0.859 ± 0.011	0.874 ± 0.011
GCN	DMoN	DMoN	\times	None	Loss	0.867 ± 0.013	0.881 ± 0.012
GCN	DMoN	L2	\times	None	MCC	0.029 ± 0.160	0.139 ± 0.101
GCN	NOCD	None	\checkmark	Transformer	Loss	0.422 ± 0.065	0.386 ± 0.081
GCN	NOCD	DMoN	\checkmark	Transformer	Loss	0.319 ± 0.129	0.262 ± 0.159
GCN	NOCD	L2	\times	None	MCC	0.146 ± 0.066	0.183 ± 0.082
GCN	Neuromap	None	\times	None	MCC	0.443 ± 0.120	0.447 ± 0.128
GCN	Neuromap	DMoN	\times	None	MCC	0.405 ± 0.098	0.375 ± 0.149
GCN	Neuromap	L2	\times	None	MCC	0.101 ± 0.161	0.171 ± 0.106
GCN	SBM_{NN}	None	\checkmark	Transformer	Loss	0.270 ± 0.040	0.207 ± 0.023
GCN	SBM_{NN}	DMoN	\checkmark	Transformer	Loss	0.269 ± 0.031	0.205 ± 0.030
GCN	SBM_{NN}	L2	\times	None	MCC	0.087 ± 0.058	0.115 ± 0.071
GraphSAGE	None	None	\times	None	Loss	0.887 ± 0.007	0.899 ± 0.006
GraphSAGE	None	DMoN	\times	None	Loss	0.810 ± 0.024	0.824 ± 0.025
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.051 ± 0.033
GraphSAGE	DMoN	None	\times	None	Loss	0.812 ± 0.035	0.825 ± 0.035
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.768 ± 0.020	0.782 ± 0.021
GraphSAGE	DMoN	L2	\checkmark	MLP	MCC	0.004 ± 0.013	0.024 ± 0.050

Continued on next page

Table 76: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.566 ± 0.099	0.576 ± 0.110
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.217 ± 0.027	0.153 ± 0.022
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.004 ± 0.015	0.037 ± 0.051
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.187 ± 0.142	0.174 ± 0.162
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.146 ± 0.068	0.089 ± 0.079
GraphSAGE	Neuromap	L2	✓	MLP	MCC	0.003 ± 0.011	0.017 ± 0.030
GraphSAGE	SBM _{NN}	None	✓	MLP	Loss	0.159 ± 0.058	0.181 ± 0.046
GraphSAGE	SBM _{NN}	DMoN	✓	Transformer	MCC	0.345 ± 0.134	0.301 ± 0.149
GraphSAGE	SBM _{NN}	L2	✓	MLP	MCC	0.007 ± 0.022	0.031 ± 0.065
Transformer	None	None	N/A	None	MCC	0.103 ± 0.077	0.117 ± 0.095
Transformer	None	DMoN	N/A	None	MCC	0.123 ± 0.109	0.133 ± 0.121
Transformer	None	L2	N/A	None	MCC	0.803 ± 0.035	0.823 ± 0.032
Transformer	DMoN	None	N/A	None	MCC	0.118 ± 0.147	0.137 ± 0.172
Transformer	DMoN	DMoN	N/A	None	MCC	0.102 ± 0.061	0.101 ± 0.080
Transformer	DMoN	L2	N/A	None	MCC	0.783 ± 0.063	0.805 ± 0.058
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.050 ± 0.026
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.036 ± 0.006
Transformer	NOCD	L2	N/A	None	MCC	0.248 ± 0.117	0.233 ± 0.137
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.049 ± 0.037
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.093 ± 0.082
Transformer	Neuromap	L2	N/A	None	MCC	0.005 ± 0.026	0.070 ± 0.066
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.037 ± 0.005
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.043 ± 0.013
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.081 ± 0.072	0.107 ± 0.068
MLP	None	None	N/A	None	Loss	0.859 ± 0.013	0.874 ± 0.012
MLP	None	DMoN	N/A	None	Loss	0.348 ± 0.033	0.284 ± 0.055
MLP	None	L2	N/A	None	MCC	0.067 ± 0.030	0.158 ± 0.045
MLP	DMoN	None	N/A	None	Loss	0.420 ± 0.050	0.365 ± 0.078
MLP	DMoN	DMoN	N/A	None	Loss	0.423 ± 0.041	0.372 ± 0.062
MLP	DMoN	L2	N/A	None	MCC	0.059 ± 0.025	0.152 ± 0.069
MLP	NOCD	None	N/A	None	Loss	0.144 ± 0.118	0.168 ± 0.115
MLP	NOCD	DMoN	N/A	None	Loss	0.139 ± 0.084	0.137 ± 0.065
MLP	NOCD	L2	N/A	None	Loss	0.647 ± 0.077	0.670 ± 0.079
MLP	Neuromap	None	N/A	None	Loss	0.018 ± 0.022	0.063 ± 0.071
MLP	Neuromap	DMoN	N/A	None	MCC	0.137 ± 0.097	0.122 ± 0.097
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.066 ± 0.062
MLP	SBM _{NN}	None	N/A	None	MCC	-0.001 ± 0.010	0.052 ± 0.048
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.009 ± 0.020	0.039 ± 0.002
MLP	SBM _{NN}	L2	N/A	None	Loss	0.012 ± 0.035	0.069 ± 0.069

Table 77: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	MCC	0.876 ± 0.015	0.889 ± 0.014
GCN	None	DMoN	✗	None	MCC	0.865 ± 0.023	0.879 ± 0.021
GCN	None	L2	✗	None	MCC	0.120 ± 0.094	0.192 ± 0.079
GCN	DMoN	None	✓	Transformer	MCC	0.875 ± 0.013	0.889 ± 0.012
GCN	DMoN	DMoN	✓	Transformer	MCC	0.861 ± 0.011	0.875 ± 0.011
GCN	DMoN	L2	✗	None	MCC	0.029 ± 0.160	0.139 ± 0.101
GCN	NOCD	None	✓	MLP	MCC	0.752 ± 0.081	0.768 ± 0.081
GCN	NOCD	DMoN	✓	MLP	MCC	0.718 ± 0.089	0.727 ± 0.099
GCN	NOCD	L2	✗	None	MCC	0.146 ± 0.066	0.183 ± 0.082
GCN	Neuromap	None	✗	None	MCC	0.443 ± 0.120	0.447 ± 0.128
GCN	Neuromap	DMoN	✗	None	MCC	0.405 ± 0.098	0.375 ± 0.149
GCN	Neuromap	L2	✗	None	MCC	0.101 ± 0.161	0.171 ± 0.106
GCN	SBM _{NN}	None	✓	MLP	MCC	0.733 ± 0.100	0.745 ± 0.110
GCN	SBM _{NN}	DMoN	✗	None	MCC	0.734 ± 0.070	0.750 ± 0.072
GCN	SBM _{NN}	L2	✗	None	MCC	0.087 ± 0.058	0.115 ± 0.071
GraphSAGE	None	None	✗	None	MCC	0.889 ± 0.008	0.901 ± 0.007
GraphSAGE	None	DMoN	✗	None	MCC	0.881 ± 0.017	0.893 ± 0.015
GraphSAGE	None	L2	✗	None	Loss	0.000 ± 0.000	0.051 ± 0.033
GraphSAGE	DMoN	None	✗	None	MCC	0.881 ± 0.015	0.894 ± 0.014
GraphSAGE	DMoN	DMoN	✓	Transformer	MCC	0.877 ± 0.013	0.890 ± 0.012
GraphSAGE	DMoN	L2	✓	Transformer	MCC	0.004 ± 0.012	0.052 ± 0.071
GraphSAGE	NOCD	None	✓	MLP	MCC	0.722 ± 0.111	0.736 ± 0.116
GraphSAGE	NOCD	DMoN	✓	MLP	MCC	0.688 ± 0.098	0.698 ± 0.116
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.004 ± 0.015	0.037 ± 0.051
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.187 ± 0.142	0.174 ± 0.162
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.146 ± 0.068	0.089 ± 0.079
GraphSAGE	Neuromap	L2	✓	Transformer	MCC	0.000 ± 0.001	0.001 ± 0.002
GraphSAGE	SBM _{NN}	None	✓	MLP	MCC	0.387 ± 0.154	0.349 ± 0.185
GraphSAGE	SBM _{NN}	DMoN	✗	None	MCC	0.342 ± 0.131	0.304 ± 0.142
GraphSAGE	SBM _{NN}	L2	✓	MLP	MCC	0.007 ± 0.022	0.031 ± 0.065
Transformer	None	None	N/A	None	MCC	0.103 ± 0.077	0.117 ± 0.095
Transformer	None	DMoN	N/A	None	MCC	0.123 ± 0.109	0.133 ± 0.121
Transformer	None	L2	N/A	None	MCC	0.803 ± 0.035	0.823 ± 0.032
Transformer	DMoN	None	N/A	None	MCC	0.118 ± 0.147	0.137 ± 0.172
Transformer	DMoN	DMoN	N/A	None	MCC	0.102 ± 0.061	0.101 ± 0.080
Transformer	DMoN	L2	N/A	None	MCC	0.783 ± 0.063	0.805 ± 0.058
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.050 ± 0.026
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.036 ± 0.006
Transformer	NOCD	L2	N/A	None	MCC	0.248 ± 0.117	0.233 ± 0.137
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.049 ± 0.037
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.093 ± 0.082
Transformer	Neuromap	L2	N/A	None	MCC	0.005 ± 0.026	0.070 ± 0.066

Continued on next page

Table 77: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.037 ± 0.005
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.043 ± 0.013
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.081 ± 0.072	0.107 ± 0.068
MLP	None	None	N/A	None	MCC	0.862 ± 0.011	0.877 ± 0.011
MLP	None	DMoN	N/A	None	MCC	0.838 ± 0.014	0.855 ± 0.013
MLP	None	L2	N/A	None	MCC	0.067 ± 0.030	0.158 ± 0.045
MLP	DMoN	None	N/A	None	MCC	0.839 ± 0.013	0.856 ± 0.013
MLP	DMoN	DMoN	N/A	None	MCC	0.837 ± 0.012	0.854 ± 0.011
MLP	DMoN	L2	N/A	None	MCC	0.059 ± 0.025	0.152 ± 0.069
MLP	NOCD	None	N/A	None	MCC	0.169 ± 0.097	0.137 ± 0.069
MLP	NOCD	DMoN	N/A	None	MCC	0.209 ± 0.121	0.161 ± 0.122
MLP	NOCD	L2	N/A	None	Loss	0.647 ± 0.077	0.670 ± 0.079
MLP	Neuromap	None	N/A	None	MCC	0.071 ± 0.092	0.059 ± 0.079
MLP	Neuromap	DMoN	N/A	None	MCC	0.137 ± 0.097	0.122 ± 0.097
MLP	Neuromap	L2	N/A	None	MCC	0.002 ± 0.015	0.069 ± 0.067
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.038 ± 0.000
MLP	SBM_{NN}	DMoN	N/A	None	MCC	0.016 ± 0.021	0.061 ± 0.087
MLP	SBM_{NN}	L2	N/A	None	Loss	0.012 ± 0.035	0.069 ± 0.069

F.4.4 COAUTHOR CS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 78: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.742 ± 0.055	0.763 ± 0.054
GCN	None	DMoN	\times	None	Loss	0.716 ± 0.053	0.736 ± 0.049
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.039 ± 0.000
GCN	DMoN	None	\checkmark	MLP	Loss	0.668 ± 0.051	0.690 ± 0.049
GCN	DMoN	DMoN	\times	None	Loss	0.707 ± 0.066	0.729 ± 0.061
GCN	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.039 ± 0.000
GCN	NOCD	None	\times	None	Loss	0.424 ± 0.086	0.430 ± 0.089
GCN	NOCD	DMoN	\checkmark	Transformer	Loss	0.372 ± 0.081	0.355 ± 0.079
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.063 ± 0.044
GCN	Neuromap	None	\checkmark	MLP	Loss	0.041 ± 0.063	0.049 ± 0.054
GCN	Neuromap	DMoN	\checkmark	MLP	Loss	0.019 ± 0.027	0.047 ± 0.069
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.039 ± 0.000
GCN	SBM_{NN}	None	\times	None	Loss	0.236 ± 0.068	0.260 ± 0.050
GCN	SBM_{NN}	DMoN	\checkmark	MLP	Loss	0.276 ± 0.083	0.291 ± 0.088
GCN	SBM_{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.104 ± 0.091
GraphSAGE	None	None	\times	None	Loss	0.727 ± 0.069	0.748 ± 0.067
GraphSAGE	None	DMoN	\times	None	Loss	0.593 ± 0.044	0.612 ± 0.043
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.044 ± 0.035

Continued on next page

Table 78: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	DMoN	None	✓	MLP	Loss	0.602 ± 0.034	0.621 ± 0.031
GraphSAGE	DMoN	DMoN	✓	MLP	Loss	0.597 ± 0.042	0.618 ± 0.039
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.043 ± 0.029
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.448 ± 0.130	0.462 ± 0.133
GraphSAGE	NOCD	DMoN	✓	Transformer	Loss	0.296 ± 0.061	0.287 ± 0.059
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.058 ± 0.065
GraphSAGE	Neuromap	None	✓	MLP	Loss	0.006 ± 0.018	0.032 ± 0.069
GraphSAGE	Neuromap	DMoN	✓	MLP	Loss	0.023 ± 0.033	0.025 ± 0.031
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.062 ± 0.062
GraphSAGE	SBM _{NN}	None	✓	MLP	Loss	0.174 ± 0.067	0.199 ± 0.059
GraphSAGE	SBM _{NN}	DMoN	✓	Transformer	Loss	0.131 ± 0.076	0.160 ± 0.080
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.062 ± 0.069
Transformer	None	None	N/A	None	Loss	0.009 ± 0.028	0.047 ± 0.049
Transformer	None	DMoN	N/A	None	Loss	0.002 ± 0.005	0.070 ± 0.065
Transformer	None	L2	N/A	None	Loss	0.095 ± 0.106	0.144 ± 0.114
Transformer	DMoN	None	N/A	None	Loss	0.008 ± 0.017	0.080 ± 0.043
Transformer	DMoN	DMoN	N/A	None	Loss	0.013 ± 0.031	0.075 ± 0.064
Transformer	DMoN	L2	N/A	None	Loss	0.049 ± 0.047	0.103 ± 0.075
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.040 ± 0.004
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.039 ± 0.000
Transformer	NOCD	L2	N/A	None	Loss	0.677 ± 0.241	0.699 ± 0.245
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.059 ± 0.036
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.057 ± 0.040
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.051 ± 0.037
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.041 ± 0.014
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.039 ± 0.000
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.054 ± 0.074	0.092 ± 0.079
MLP	None	None	N/A	None	Loss	0.566 ± 0.084	0.589 ± 0.101
MLP	None	DMoN	N/A	None	Loss	0.246 ± 0.034	0.260 ± 0.038
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.045 ± 0.040
MLP	DMoN	None	N/A	None	Loss	0.467 ± 0.067	0.493 ± 0.077
MLP	DMoN	DMoN	N/A	None	Loss	0.475 ± 0.079	0.497 ± 0.084
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.049 ± 0.041
MLP	NOCD	None	N/A	None	Loss	0.155 ± 0.155	0.186 ± 0.156
MLP	NOCD	DMoN	N/A	None	Loss	0.126 ± 0.077	0.146 ± 0.073
MLP	NOCD	L2	N/A	None	Loss	0.710 ± 0.046	0.738 ± 0.044
MLP	Neuromap	None	N/A	None	Loss	0.029 ± 0.052	0.079 ± 0.048
MLP	Neuromap	DMoN	N/A	None	Loss	0.008 ± 0.012	0.059 ± 0.071
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.038 ± 0.036
MLP	SBM _{NN}	None	N/A	None	Loss	0.003 ± 0.009	0.048 ± 0.026
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.003 ± 0.007	0.034 ± 0.007
MLP	SBM _{NN}	L2	N/A	None	Loss	0.092 ± 0.140	0.132 ± 0.135

Table 79: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.742 ± 0.055	0.763 ± 0.054
GCN	None	DMoN	\times	None	Loss	0.716 ± 0.053	0.736 ± 0.049
GCN	None	L2	\times	None	MCC	-0.014 ± 0.084	0.084 ± 0.057
GCN	DMoN	None	\checkmark	MLP	Loss	0.668 ± 0.051	0.690 ± 0.049
GCN	DMoN	DMoN	\times	None	Loss	0.707 ± 0.066	0.729 ± 0.061
GCN	DMoN	L2	\times	None	MCC	-0.004 ± 0.076	0.076 ± 0.062
GCN	NOCD	None	\times	None	Loss	0.424 ± 0.086	0.430 ± 0.089
GCN	NOCD	DMoN	\checkmark	Transformer	MCC	0.605 ± 0.097	0.628 ± 0.101
GCN	NOCD	L2	\checkmark	Transformer	MCC	0.060 ± 0.073	0.123 ± 0.059
GCN	Neuromap	None	\times	None	MCC	0.209 ± 0.174	0.206 ± 0.172
GCN	Neuromap	DMoN	\times	None	MCC	0.283 ± 0.157	0.282 ± 0.177
GCN	Neuromap	L2	\times	None	MCC	0.047 ± 0.099	0.126 ± 0.079
GCN	SBM _{NN}	None	\checkmark	MLP	Loss	0.247 ± 0.126	0.257 ± 0.123
GCN	SBM _{NN}	DMoN	\checkmark	Transformer	Loss	0.262 ± 0.104	0.270 ± 0.104
GCN	SBM _{NN}	L2	\checkmark	Transformer	MCC	0.043 ± 0.053	0.083 ± 0.045
GraphSAGE	None	None	\times	None	Loss	0.727 ± 0.069	0.748 ± 0.067
GraphSAGE	None	DMoN	\times	None	Loss	0.593 ± 0.044	0.612 ± 0.043
GraphSAGE	None	L2	\times	None	MCC	-0.000 ± 0.000	0.002 ± 0.007
GraphSAGE	DMoN	None	\times	None	Loss	0.637 ± 0.029	0.654 ± 0.027
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.623 ± 0.028	0.643 ± 0.026
GraphSAGE	DMoN	L2	\checkmark	MLP	MCC	-0.000 ± 0.000	0.002 ± 0.007
GraphSAGE	NOCD	None	\checkmark	Transformer	Loss	0.448 ± 0.130	0.462 ± 0.133
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.327 ± 0.042	0.307 ± 0.048
GraphSAGE	NOCD	L2	\checkmark	MLP	MCC	0.002 ± 0.007	0.006 ± 0.013
GraphSAGE	Neuromap	None	\checkmark	MLP	MCC	0.111 ± 0.102	0.098 ± 0.109
GraphSAGE	Neuromap	DMoN	\checkmark	MLP	MCC	0.097 ± 0.105	0.106 ± 0.112
GraphSAGE	Neuromap	L2	\times	None	MCC	0.003 ± 0.011	0.006 ± 0.012
GraphSAGE	SBM _{NN}	None	\checkmark	MLP	Loss	0.174 ± 0.067	0.199 ± 0.059
GraphSAGE	SBM _{NN}	DMoN	\checkmark	Transformer	Loss	0.131 ± 0.076	0.160 ± 0.080
GraphSAGE	SBM _{NN}	L2	\checkmark	Transformer	MCC	0.006 ± 0.010	0.037 ± 0.073
Transformer	None	None	N/A	None	MCC	0.042 ± 0.049	0.056 ± 0.081
Transformer	None	DMoN	N/A	None	MCC	0.039 ± 0.046	0.059 ± 0.060
Transformer	None	L2	N/A	None	MCC	0.476 ± 0.120	0.516 ± 0.130
Transformer	DMoN	None	N/A	None	Loss	0.008 ± 0.017	0.080 ± 0.043
Transformer	DMoN	DMoN	N/A	None	Loss	0.013 ± 0.031	0.075 ± 0.064
Transformer	DMoN	L2	N/A	None	MCC	0.518 ± 0.104	0.552 ± 0.112
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.040 ± 0.004
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.039 ± 0.000
Transformer	NOCD	L2	N/A	None	Loss	0.677 ± 0.241	0.699 ± 0.245
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.059 ± 0.036
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.057 ± 0.040
Transformer	Neuromap	L2	N/A	None	MCC	0.001 ± 0.011	0.049 ± 0.050

Continued on next page

Table 79: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.041 ± 0.014
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.039 ± 0.000
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.054 ± 0.074	0.092 ± 0.079
MLP	None	None	N/A	None	Loss	0.566 ± 0.084	0.589 ± 0.101
MLP	None	DMoN	N/A	None	Loss	0.246 ± 0.034	0.260 ± 0.038
MLP	None	L2	N/A	None	MCC	0.016 ± 0.020	0.092 ± 0.063
MLP	DMoN	None	N/A	None	Loss	0.467 ± 0.067	0.493 ± 0.077
MLP	DMoN	DMoN	N/A	None	Loss	0.475 ± 0.079	0.497 ± 0.084
MLP	DMoN	L2	N/A	None	MCC	0.024 ± 0.026	0.115 ± 0.089
MLP	NOCD	None	N/A	None	Loss	0.155 ± 0.155	0.186 ± 0.156
MLP	NOCD	DMoN	N/A	None	Loss	0.126 ± 0.077	0.146 ± 0.073
MLP	NOCD	L2	N/A	None	Loss	0.710 ± 0.046	0.738 ± 0.044
MLP	Neuromap	None	N/A	None	MCC	0.030 ± 0.045	0.041 ± 0.057
MLP	Neuromap	DMoN	N/A	None	MCC	0.044 ± 0.044	0.053 ± 0.047
MLP	Neuromap	L2	N/A	None	MCC	-0.002 ± 0.019	0.078 ± 0.084
MLP	SBM_{NN}	None	N/A	None	MCC	0.006 ± 0.015	0.025 ± 0.038
MLP	SBM_{NN}	DMoN	N/A	None	MCC	0.021 ± 0.035	0.040 ± 0.040
MLP	SBM_{NN}	L2	N/A	None	Loss	0.092 ± 0.140	0.132 ± 0.135

Table 80: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	MCC	0.757 ± 0.051	0.779 ± 0.052
GCN	None	DMoN	\times	None	MCC	0.748 ± 0.052	0.770 ± 0.050
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.039 ± 0.000
GCN	DMoN	None	\times	None	MCC	0.756 ± 0.045	0.778 ± 0.042
GCN	DMoN	DMoN	\checkmark	Transformer	MCC	0.754 ± 0.040	0.776 ± 0.040
GCN	DMoN	L2	\checkmark	MLP	MCC	0.007 ± 0.022	0.079 ± 0.060
GCN	NOCD	None	\times	None	MCC	0.551 ± 0.118	0.573 ± 0.117
GCN	NOCD	DMoN	\checkmark	Transformer	MCC	0.605 ± 0.097	0.628 ± 0.101
GCN	NOCD	L2	\times	None	MCC	0.052 ± 0.057	0.085 ± 0.071
GCN	Neuromap	None	\times	None	MCC	0.209 ± 0.174	0.206 ± 0.172
GCN	Neuromap	DMoN	\times	None	MCC	0.283 ± 0.157	0.282 ± 0.177
GCN	Neuromap	L2	\times	None	MCC	0.047 ± 0.099	0.126 ± 0.079
GCN	SBM_{NN}	None	\checkmark	Transformer	MCC	0.509 ± 0.119	0.528 ± 0.125
GCN	SBM_{NN}	DMoN	\checkmark	MLP	MCC	0.553 ± 0.142	0.568 ± 0.154
GCN	SBM_{NN}	L2	\checkmark	Transformer	MCC	0.043 ± 0.053	0.083 ± 0.045
GraphSAGE	None	None	\times	None	MCC	0.735 ± 0.069	0.757 ± 0.066
GraphSAGE	None	DMoN	\times	None	MCC	0.758 ± 0.042	0.778 ± 0.041
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.044 ± 0.035
GraphSAGE	DMoN	None	\times	None	MCC	0.758 ± 0.029	0.778 ± 0.028
GraphSAGE	DMoN	DMoN	\times	None	MCC	0.764 ± 0.037	0.784 ± 0.037
GraphSAGE	DMoN	L2	\checkmark	Transformer	MCC	0.003 ± 0.008	0.019 ± 0.059

Continued on next page

Table 80: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✓	MLP	MCC	0.514 ± 0.137	0.539 ± 0.143
GraphSAGE	NOCD	DMoN	✓	Transformer	MCC	0.589 ± 0.076	0.620 ± 0.070
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.002 ± 0.007	0.006 ± 0.013
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.111 ± 0.102	0.098 ± 0.109
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.097 ± 0.105	0.106 ± 0.112
GraphSAGE	Neuromap	L2	✓	MLP	MCC	0.006 ± 0.018	0.012 ± 0.039
GraphSAGE	SBM _{NN}	None	✗	None	MCC	0.289 ± 0.105	0.260 ± 0.106
GraphSAGE	SBM _{NN}	DMoN	✗	None	MCC	0.281 ± 0.091	0.257 ± 0.104
GraphSAGE	SBM _{NN}	L2	✗	None	MCC	0.002 ± 0.010	0.017 ± 0.029
Transformer	None	None	N/A	None	MCC	0.042 ± 0.049	0.056 ± 0.081
Transformer	None	DMoN	N/A	None	MCC	0.039 ± 0.046	0.059 ± 0.060
Transformer	None	L2	N/A	None	MCC	0.476 ± 0.120	0.516 ± 0.130
Transformer	DMoN	None	N/A	None	MCC	0.021 ± 0.047	0.050 ± 0.086
Transformer	DMoN	DMoN	N/A	None	MCC	0.045 ± 0.080	0.068 ± 0.099
Transformer	DMoN	L2	N/A	None	MCC	0.518 ± 0.104	0.552 ± 0.112
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.040 ± 0.004
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.039 ± 0.000
Transformer	NOCD	L2	N/A	None	Loss	0.677 ± 0.241	0.699 ± 0.245
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.059 ± 0.036
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.057 ± 0.040
Transformer	Neuromap	L2	N/A	None	MCC	0.001 ± 0.011	0.049 ± 0.050
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.041 ± 0.014
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.039 ± 0.000
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.054 ± 0.074	0.092 ± 0.079
MLP	None	None	N/A	None	MCC	0.602 ± 0.061	0.639 ± 0.059
MLP	None	DMoN	N/A	None	MCC	0.514 ± 0.061	0.548 ± 0.066
MLP	None	L2	N/A	None	MCC	0.016 ± 0.020	0.092 ± 0.063
MLP	DMoN	None	N/A	None	MCC	0.527 ± 0.048	0.554 ± 0.058
MLP	DMoN	DMoN	N/A	None	MCC	0.510 ± 0.030	0.544 ± 0.035
MLP	DMoN	L2	N/A	None	MCC	0.024 ± 0.026	0.115 ± 0.089
MLP	NOCD	None	N/A	None	MCC	0.166 ± 0.087	0.133 ± 0.077
MLP	NOCD	DMoN	N/A	None	MCC	0.146 ± 0.088	0.121 ± 0.095
MLP	NOCD	L2	N/A	None	Loss	0.710 ± 0.046	0.738 ± 0.044
MLP	Neuromap	None	N/A	None	MCC	0.030 ± 0.045	0.041 ± 0.057
MLP	Neuromap	DMoN	N/A	None	MCC	0.044 ± 0.044	0.053 ± 0.047
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.038 ± 0.036
MLP	SBM _{NN}	None	N/A	None	MCC	0.006 ± 0.015	0.025 ± 0.038
MLP	SBM _{NN}	DMoN	N/A	None	MCC	0.021 ± 0.035	0.040 ± 0.040
MLP	SBM _{NN}	L2	N/A	None	Loss	0.092 ± 0.140	0.132 ± 0.135

F.4.5 AMAZON COMPUTERS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 81: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.716 ± 0.050	0.763 ± 0.045
GCN	None	DMoN	\times	None	Loss	0.723 ± 0.036	0.770 ± 0.032
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.031 ± 0.000
GCN	DMoN	None	\checkmark	Transformer	Loss	0.706 ± 0.037	0.754 ± 0.030
GCN	DMoN	DMoN	\times	None	Loss	0.709 ± 0.041	0.756 ± 0.038
GCN	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.031 ± 0.000
GCN	NOCD	None	\checkmark	MLP	Loss	0.587 ± 0.069	0.626 ± 0.082
GCN	NOCD	DMoN	\checkmark	MLP	Loss	0.578 ± 0.102	0.606 ± 0.132
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.112 ± 0.106
GCN	Neuromap	None	\checkmark	Transformer	Loss	-0.002 ± 0.004	0.153 ± 0.161
GCN	Neuromap	DMoN	\checkmark	Transformer	Loss	-0.001 ± 0.002	0.056 ± 0.054
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.031 ± 0.000
GCN	SBM _{NN}	None	\times	None	Loss	0.469 ± 0.074	0.489 ± 0.098
GCN	SBM _{NN}	DMoN	\checkmark	MLP	Loss	0.406 ± 0.127	0.392 ± 0.177
GCN	SBM _{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.154 ± 0.161
GraphSAGE	None	None	\times	None	Loss	0.740 ± 0.021	0.781 ± 0.020
GraphSAGE	None	DMoN	\times	None	Loss	0.728 ± 0.027	0.771 ± 0.026
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.107 ± 0.108
GraphSAGE	DMoN	None	\times	None	Loss	0.728 ± 0.022	0.772 ± 0.021
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.734 ± 0.027	0.777 ± 0.025
GraphSAGE	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.095 ± 0.109
GraphSAGE	NOCD	None	\checkmark	Transformer	Loss	0.372 ± 0.086	0.346 ± 0.119
GraphSAGE	NOCD	DMoN	\times	None	Loss	0.290 ± 0.044	0.225 ± 0.086
GraphSAGE	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.103 ± 0.112
GraphSAGE	Neuromap	None	\checkmark	Transformer	Loss	0.003 ± 0.008	0.055 ± 0.056
GraphSAGE	Neuromap	DMoN	\times	None	Loss	0.002 ± 0.007	0.151 ± 0.135
GraphSAGE	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.065 ± 0.054
GraphSAGE	SBM _{NN}	None	\checkmark	Transformer	Loss	0.180 ± 0.098	0.177 ± 0.132
GraphSAGE	SBM _{NN}	DMoN	\checkmark	MLP	Loss	0.095 ± 0.082	0.094 ± 0.070
GraphSAGE	SBM _{NN}	L2	\checkmark	MLP	Loss	0.000 ± 0.000	0.111 ± 0.111
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.087 ± 0.061
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.186 ± 0.144
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.136 ± 0.137
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.138 ± 0.136
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.131 ± 0.140
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.127 ± 0.139
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.030 ± 0.005
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.055 ± 0.054
Transformer	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.121 ± 0.111
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.105 ± 0.147
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.145 ± 0.135
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.214 ± 0.176

Continued on next page

Table 81: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.031 ± 0.001
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.030 ± 0.003
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.002 ± 0.005	0.092 ± 0.110
MLP	None	None	N/A	None	Loss	0.547 ± 0.031	0.612 ± 0.032
MLP	None	DMoN	N/A	None	Loss	0.437 ± 0.028	0.482 ± 0.032
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.081 ± 0.114
MLP	DMoN	None	N/A	None	Loss	0.443 ± 0.030	0.497 ± 0.035
MLP	DMoN	DMoN	N/A	None	Loss	0.431 ± 0.023	0.473 ± 0.029
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.083 ± 0.049
MLP	NOCD	None	N/A	None	Loss	0.001 ± 0.002	0.029 ± 0.004
MLP	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.042 ± 0.040
MLP	NOCD	L2	N/A	None	Loss	0.017 ± 0.037	0.116 ± 0.107
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.001	0.136 ± 0.136
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.126 ± 0.103
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.079 ± 0.058
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.001	0.031 ± 0.000
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.031 ± 0.000
MLP	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.030 ± 0.003

Table 82: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.716 ± 0.050	0.763 ± 0.045
GCN	None	DMoN	\times	None	Loss	0.723 ± 0.036	0.770 ± 0.032
GCN	None	L2	\times	None	MCC	0.053 ± 0.137	0.198 ± 0.128
GCN	DMoN	None	\checkmark	Transformer	Loss	0.706 ± 0.037	0.754 ± 0.030
GCN	DMoN	DMoN	\times	None	Loss	0.709 ± 0.041	0.756 ± 0.038
GCN	DMoN	L2	\checkmark	Transformer	MCC	0.033 ± 0.073	0.138 ± 0.091
GCN	NOCD	None	\checkmark	MLP	Loss	0.587 ± 0.069	0.626 ± 0.082
GCN	NOCD	DMoN	\checkmark	MLP	Loss	0.578 ± 0.102	0.606 ± 0.132
GCN	NOCD	L2	\times	None	MCC	0.115 ± 0.039	0.195 ± 0.074
GCN	Neuromap	None	\times	None	MCC	0.086 ± 0.066	0.084 ± 0.058
GCN	Neuromap	DMoN	\times	None	MCC	0.113 ± 0.103	0.118 ± 0.140
GCN	Neuromap	L2	\times	None	MCC	-0.022 ± 0.149	0.162 ± 0.109
GCN	SBM_{NN}	None	\times	None	Loss	0.469 ± 0.074	0.489 ± 0.098
GCN	SBM_{NN}	DMoN	\checkmark	MLP	Loss	0.406 ± 0.127	0.392 ± 0.177
GCN	SBM_{NN}	L2	\times	None	MCC	0.051 ± 0.092	0.118 ± 0.114
GraphSAGE	None	None	\times	None	Loss	0.740 ± 0.021	0.781 ± 0.020
GraphSAGE	None	DMoN	\times	None	Loss	0.728 ± 0.027	0.771 ± 0.026
GraphSAGE	None	L2	\times	None	MCC	0.002 ± 0.014	0.074 ± 0.115
GraphSAGE	DMoN	None	\times	None	Loss	0.728 ± 0.022	0.772 ± 0.021
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.734 ± 0.027	0.777 ± 0.025
GraphSAGE	DMoN	L2	\checkmark	Transformer	MCC	0.002 ± 0.015	0.081 ± 0.111

Continued on next page

Table 82: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.372 ± 0.086	0.346 ± 0.119
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.290 ± 0.044	0.225 ± 0.086
GraphSAGE	NOCD	L2	✗	None	MCC	0.006 ± 0.012	0.052 ± 0.060
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.070 ± 0.081	0.109 ± 0.118
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.029 ± 0.058	0.060 ± 0.132
GraphSAGE	Neuromap	L2	✓	MLP	MCC	0.003 ± 0.007	0.040 ± 0.051
GraphSAGE	SBM _{NN}	None	✓	Transformer	Loss	0.180 ± 0.098	0.177 ± 0.132
GraphSAGE	SBM _{NN}	DMoN	✓	MLP	Loss	0.095 ± 0.082	0.094 ± 0.070
GraphSAGE	SBM _{NN}	L2	✗	None	MCC	0.010 ± 0.032	0.085 ± 0.112
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.087 ± 0.061
Transformer	None	DMoN	N/A	None	MCC	0.004 ± 0.013	0.016 ± 0.052
Transformer	None	L2	N/A	None	MCC	0.031 ± 0.050	0.234 ± 0.095
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.138 ± 0.136
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.131 ± 0.140
Transformer	DMoN	L2	N/A	None	MCC	0.048 ± 0.037	0.193 ± 0.106
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.030 ± 0.005
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.055 ± 0.054
Transformer	NOCD	L2	N/A	None	MCC	0.030 ± 0.035	0.133 ± 0.104
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.105 ± 0.147
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.145 ± 0.135
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.214 ± 0.176
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.031 ± 0.001
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.030 ± 0.003
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.002 ± 0.005	0.092 ± 0.110
MLP	None	None	N/A	None	Loss	0.547 ± 0.031	0.612 ± 0.032
MLP	None	DMoN	N/A	None	Loss	0.437 ± 0.028	0.482 ± 0.032
MLP	None	L2	N/A	None	MCC	-0.017 ± 0.048	0.188 ± 0.107
MLP	DMoN	None	N/A	None	Loss	0.443 ± 0.030	0.497 ± 0.035
MLP	DMoN	DMoN	N/A	None	Loss	0.431 ± 0.023	0.473 ± 0.029
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.083 ± 0.049
MLP	NOCD	None	N/A	None	MCC	0.002 ± 0.010	0.014 ± 0.026
MLP	NOCD	DMoN	N/A	None	MCC	0.001 ± 0.003	0.041 ± 0.051
MLP	NOCD	L2	N/A	None	Loss	0.017 ± 0.037	0.116 ± 0.107
MLP	Neuromap	None	N/A	None	MCC	0.004 ± 0.010	0.131 ± 0.147
MLP	Neuromap	DMoN	N/A	None	MCC	0.004 ± 0.007	0.042 ± 0.052
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.079 ± 0.058
MLP	SBM _{NN}	None	N/A	None	MCC	0.009 ± 0.022	0.089 ± 0.120
MLP	SBM _{NN}	DMoN	N/A	None	MCC	0.001 ± 0.005	0.052 ± 0.060
MLP	SBM _{NN}	L2	N/A	None	MCC	0.003 ± 0.004	0.072 ± 0.113

Table 83: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	Loss	0.716 ± 0.050	0.763 ± 0.045
GCN	None	DMoN	✗	None	Loss	0.723 ± 0.036	0.770 ± 0.032
GCN	None	L2	✗	None	MCC	0.053 ± 0.137	0.198 ± 0.128
GCN	DMoN	None	✓	Transformer	Loss	0.706 ± 0.037	0.754 ± 0.030
GCN	DMoN	DMoN	✗	None	Loss	0.709 ± 0.041	0.756 ± 0.038
GCN	DMoN	L2	✓	Transformer	MCC	0.033 ± 0.073	0.138 ± 0.091
GCN	NOCD	None	✓	MLP	Loss	0.587 ± 0.069	0.626 ± 0.082
GCN	NOCD	DMoN	✓	Transformer	MCC	0.576 ± 0.120	0.630 ± 0.127
GCN	NOCD	L2	✗	None	MCC	0.115 ± 0.039	0.195 ± 0.074
GCN	Neuromap	None	✗	None	MCC	0.086 ± 0.066	0.084 ± 0.058
GCN	Neuromap	DMoN	✗	None	MCC	0.113 ± 0.103	0.118 ± 0.140
GCN	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.031 ± 0.000
GCN	SBM _{NN}	None	✗	None	Loss	0.469 ± 0.074	0.489 ± 0.098
GCN	SBM _{NN}	DMoN	✓	Transformer	Loss	0.455 ± 0.168	0.471 ± 0.206
GCN	SBM _{NN}	L2	✗	None	MCC	0.072 ± 0.042	0.148 ± 0.102
GraphSAGE	None	None	✗	None	Loss	0.740 ± 0.021	0.781 ± 0.020
GraphSAGE	None	DMoN	✗	None	MCC	0.751 ± 0.020	0.794 ± 0.018
GraphSAGE	None	L2	✗	None	MCC	0.002 ± 0.014	0.074 ± 0.115
GraphSAGE	DMoN	None	✗	None	MCC	0.732 ± 0.028	0.777 ± 0.026
GraphSAGE	DMoN	DMoN	✗	None	MCC	0.726 ± 0.075	0.765 ± 0.086
GraphSAGE	DMoN	L2	✗	None	MCC	0.003 ± 0.020	0.079 ± 0.107
GraphSAGE	NOCD	None	✓	MLP	MCC	0.555 ± 0.148	0.574 ± 0.172
GraphSAGE	NOCD	DMoN	✓	Transformer	MCC	0.552 ± 0.132	0.584 ± 0.166
GraphSAGE	NOCD	L2	✗	None	MCC	0.006 ± 0.012	0.052 ± 0.060
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.070 ± 0.081	0.109 ± 0.118
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.029 ± 0.058	0.060 ± 0.132
GraphSAGE	Neuromap	L2	✗	None	MCC	0.000 ± 0.039	0.083 ± 0.095
GraphSAGE	SBM _{NN}	None	✓	MLP	MCC	0.212 ± 0.134	0.237 ± 0.122
GraphSAGE	SBM _{NN}	DMoN	✗	None	MCC	0.231 ± 0.136	0.229 ± 0.165
GraphSAGE	SBM _{NN}	L2	✗	None	MCC	0.010 ± 0.032	0.085 ± 0.112
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.087 ± 0.061
Transformer	None	DMoN	N/A	None	MCC	0.004 ± 0.013	0.016 ± 0.052
Transformer	None	L2	N/A	None	MCC	0.031 ± 0.050	0.234 ± 0.095
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.138 ± 0.136
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.131 ± 0.140
Transformer	DMoN	L2	N/A	None	MCC	0.048 ± 0.037	0.193 ± 0.106
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.030 ± 0.005
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.055 ± 0.054
Transformer	NOCD	L2	N/A	None	MCC	0.030 ± 0.035	0.133 ± 0.104
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.105 ± 0.147
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.145 ± 0.135
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.214 ± 0.176

Continued on next page

Table 83: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.031 ± 0.001
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.030 ± 0.003
Transformer	SBM_{NN}	L2	N/A	None	MCC	0.003 ± 0.007	0.030 ± 0.062
MLP	None	None	N/A	None	MCC	0.586 ± 0.026	0.656 ± 0.024
MLP	None	DMoN	N/A	None	MCC	0.552 ± 0.025	0.619 ± 0.024
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.081 ± 0.114
MLP	DMoN	None	N/A	None	MCC	0.553 ± 0.027	0.616 ± 0.026
MLP	DMoN	DMoN	N/A	None	MCC	0.550 ± 0.033	0.616 ± 0.039
MLP	DMoN	L2	N/A	None	MCC	0.001 ± 0.053	0.159 ± 0.093
MLP	NOCD	None	N/A	None	Loss	0.001 ± 0.002	0.029 ± 0.004
MLP	NOCD	DMoN	N/A	None	MCC	0.001 ± 0.003	0.041 ± 0.051
MLP	NOCD	L2	N/A	None	Loss	0.017 ± 0.037	0.116 ± 0.107
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.001	0.136 ± 0.136
MLP	Neuromap	DMoN	N/A	None	MCC	0.004 ± 0.007	0.042 ± 0.052
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.079 ± 0.058
MLP	SBM_{NN}	None	N/A	None	MCC	0.009 ± 0.022	0.089 ± 0.120
MLP	SBM_{NN}	DMoN	N/A	None	MCC	0.001 ± 0.005	0.052 ± 0.060
MLP	SBM_{NN}	L2	N/A	None	MCC	0.003 ± 0.004	0.072 ± 0.113

F.4.6 AMAZON COMPUTERS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 84: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.502 ± 0.096	0.569 ± 0.101
GCN	None	DMoN	\times	None	Loss	0.521 ± 0.064	0.580 ± 0.069
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.032 ± 0.000
GCN	DMoN	None	\times	None	Loss	0.476 ± 0.086	0.529 ± 0.097
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.507 ± 0.065	0.574 ± 0.065
GCN	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.032 ± 0.000
GCN	NOCD	None	\times	None	Loss	0.328 ± 0.094	0.302 ± 0.147
GCN	NOCD	DMoN	\checkmark	MLP	Loss	0.297 ± 0.070	0.271 ± 0.085
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.067 ± 0.054
GCN	Neuromap	None	\checkmark	MLP	Loss	0.000 ± 0.001	0.075 ± 0.052
GCN	Neuromap	DMoN	\checkmark	Transformer	Loss	-0.001 ± 0.002	0.101 ± 0.114
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.032 ± 0.000
GCN	SBM_{NN}	None	\checkmark	Transformer	Loss	0.294 ± 0.051	0.318 ± 0.082
GCN	SBM_{NN}	DMoN	\checkmark	MLP	Loss	0.271 ± 0.084	0.260 ± 0.110
GCN	SBM_{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.093 ± 0.068
GraphSAGE	None	None	\times	None	Loss	0.427 ± 0.105	0.470 ± 0.120
GraphSAGE	None	DMoN	\times	None	Loss	0.466 ± 0.105	0.522 ± 0.114
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.063 ± 0.052

Continued on next page

Table 84: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	DMoN	None	✗	None	Loss	0.389 ± 0.089	0.434 ± 0.103
GraphSAGE	DMoN	DMoN	✗	None	Loss	0.453 ± 0.092	0.504 ± 0.100
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.091 ± 0.113
GraphSAGE	NOCD	None	✗	None	Loss	0.260 ± 0.119	0.282 ± 0.127
GraphSAGE	NOCD	DMoN	✓	Transformer	Loss	0.210 ± 0.075	0.162 ± 0.096
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.073 ± 0.049
GraphSAGE	Neuromap	None	✓	Transformer	Loss	0.004 ± 0.020	0.125 ± 0.137
GraphSAGE	Neuromap	DMoN	✓	Transformer	Loss	-0.001 ± 0.004	0.109 ± 0.108
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.117 ± 0.143
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.082 ± 0.094	0.089 ± 0.098
GraphSAGE	SBM _{NN}	DMoN	✓	MLP	Loss	0.046 ± 0.063	0.054 ± 0.020
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.097 ± 0.107
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.075 ± 0.051
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.120 ± 0.107
Transformer	None	L2	N/A	None	Loss	-0.000 ± 0.000	0.078 ± 0.059
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.133 ± 0.138
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.090 ± 0.109
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.156 ± 0.157
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.032 ± 0.004
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.050 ± 0.044
Transformer	NOCD	L2	N/A	None	Loss	0.106 ± 0.093	0.177 ± 0.105
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.122 ± 0.136
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.073 ± 0.041
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.071 ± 0.054
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.031 ± 0.003
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.032 ± 0.001
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.006 ± 0.019	0.099 ± 0.108
MLP	None	None	N/A	None	Loss	0.150 ± 0.053	0.207 ± 0.057
MLP	None	DMoN	N/A	None	Loss	0.192 ± 0.035	0.278 ± 0.048
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.098 ± 0.111
MLP	DMoN	None	N/A	None	Loss	0.195 ± 0.065	0.257 ± 0.109
MLP	DMoN	DMoN	N/A	None	Loss	0.189 ± 0.042	0.260 ± 0.070
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.111 ± 0.109
MLP	NOCD	None	N/A	None	Loss	0.003 ± 0.005	0.040 ± 0.024
MLP	NOCD	DMoN	N/A	None	Loss	0.002 ± 0.005	0.040 ± 0.023
MLP	NOCD	L2	N/A	None	Loss	0.108 ± 0.076	0.166 ± 0.123
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.125 ± 0.105
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.137 ± 0.136
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.092 ± 0.108
MLP	SBM _{NN}	None	N/A	None	Loss	-0.000 ± 0.000	0.097 ± 0.110
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.032 ± 0.000
MLP	SBM _{NN}	L2	N/A	None	Loss	0.006 ± 0.017	0.057 ± 0.052

Table 85: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	Loss	0.502 ± 0.096	0.569 ± 0.101
GCN	None	DMoN	✗	None	Loss	0.521 ± 0.064	0.580 ± 0.069
GCN	None	L2	✗	None	MCC	0.030 ± 0.102	0.152 ± 0.098
GCN	DMoN	None	✗	None	Loss	0.476 ± 0.086	0.529 ± 0.097
GCN	DMoN	DMoN	✓	MLP	Loss	0.507 ± 0.065	0.574 ± 0.065
GCN	DMoN	L2	✓	MLP	MCC	0.029 ± 0.054	0.219 ± 0.108
GCN	NOCD	None	✗	None	Loss	0.328 ± 0.094	0.302 ± 0.147
GCN	NOCD	DMoN	✓	MLP	Loss	0.297 ± 0.070	0.271 ± 0.085
GCN	NOCD	L2	✓	MLP	MCC	0.057 ± 0.074	0.173 ± 0.126
GCN	Neuromap	None	✗	None	MCC	0.055 ± 0.066	0.086 ± 0.071
GCN	Neuromap	DMoN	✗	None	MCC	0.049 ± 0.063	0.061 ± 0.070
GCN	Neuromap	L2	✗	None	MCC	0.013 ± 0.099	0.183 ± 0.107
GCN	SBM _{NN}	None	✓	Transformer	Loss	0.294 ± 0.051	0.318 ± 0.082
GCN	SBM _{NN}	DMoN	✓	MLP	Loss	0.271 ± 0.084	0.260 ± 0.110
GCN	SBM _{NN}	L2	✓	MLP	MCC	0.039 ± 0.063	0.171 ± 0.114
GraphSAGE	None	None	✗	None	Loss	0.427 ± 0.105	0.470 ± 0.120
GraphSAGE	None	DMoN	✗	None	Loss	0.466 ± 0.105	0.522 ± 0.114
GraphSAGE	None	L2	✗	None	Loss	0.000 ± 0.000	0.063 ± 0.052
GraphSAGE	DMoN	None	✗	None	Loss	0.389 ± 0.089	0.434 ± 0.103
GraphSAGE	DMoN	DMoN	✗	None	Loss	0.453 ± 0.092	0.504 ± 0.100
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.091 ± 0.113
GraphSAGE	NOCD	None	✓	MLP	Loss	0.234 ± 0.134	0.245 ± 0.156
GraphSAGE	NOCD	DMoN	✓	Transformer	Loss	0.210 ± 0.075	0.162 ± 0.096
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.004 ± 0.022	0.087 ± 0.106
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.052 ± 0.072	0.111 ± 0.107
GraphSAGE	Neuromap	DMoN	✗	None	MCC	0.037 ± 0.059	0.143 ± 0.145
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.117 ± 0.143
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.082 ± 0.094	0.089 ± 0.098
GraphSAGE	SBM _{NN}	DMoN	✓	MLP	Loss	0.046 ± 0.063	0.054 ± 0.020
GraphSAGE	SBM _{NN}	L2	✓	MLP	MCC	0.004 ± 0.019	0.070 ± 0.118
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.075 ± 0.051
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.120 ± 0.107
Transformer	None	L2	N/A	None	MCC	-0.001 ± 0.030	0.172 ± 0.101
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.133 ± 0.138
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.090 ± 0.109
Transformer	DMoN	L2	N/A	None	MCC	0.014 ± 0.028	0.161 ± 0.110
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.032 ± 0.004
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.050 ± 0.044
Transformer	NOCD	L2	N/A	None	Loss	0.106 ± 0.093	0.177 ± 0.105
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.122 ± 0.136
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.073 ± 0.041
Transformer	Neuromap	L2	N/A	None	MCC	0.002 ± 0.011	0.075 ± 0.060

Continued on next page

Table 85: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.031 ± 0.003
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.032 ± 0.001
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.006 ± 0.019	0.099 ± 0.108
MLP	None	None	N/A	None	Loss	0.150 ± 0.053	0.207 ± 0.057
MLP	None	DMoN	N/A	None	Loss	0.192 ± 0.035	0.278 ± 0.048
MLP	None	L2	N/A	None	MCC	0.003 ± 0.024	0.125 ± 0.106
MLP	DMoN	None	N/A	None	Loss	0.195 ± 0.065	0.257 ± 0.109
MLP	DMoN	DMoN	N/A	None	Loss	0.189 ± 0.042	0.260 ± 0.070
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.111 ± 0.109
MLP	NOCD	None	N/A	None	Loss	0.003 ± 0.005	0.040 ± 0.024
MLP	NOCD	DMoN	N/A	None	Loss	0.002 ± 0.005	0.040 ± 0.023
MLP	NOCD	L2	N/A	None	Loss	0.108 ± 0.076	0.166 ± 0.123
MLP	Neuromap	None	N/A	None	MCC	0.005 ± 0.011	0.083 ± 0.119
MLP	Neuromap	DMoN	N/A	None	MCC	0.020 ± 0.029	0.056 ± 0.073
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.092 ± 0.108
MLP	SBM_{NN}	None	N/A	None	Loss	-0.000 ± 0.000	0.097 ± 0.110
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.032 ± 0.000
MLP	SBM_{NN}	L2	N/A	None	Loss	0.006 ± 0.017	0.057 ± 0.052

Table 86: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.502 ± 0.096	0.569 ± 0.101
GCN	None	DMoN	\times	None	Loss	0.521 ± 0.064	0.580 ± 0.069
GCN	None	L2	\times	None	MCC	0.030 ± 0.102	0.152 ± 0.098
GCN	DMoN	None	\times	None	Loss	0.476 ± 0.086	0.529 ± 0.097
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.507 ± 0.065	0.574 ± 0.065
GCN	DMoN	L2	\checkmark	MLP	MCC	0.029 ± 0.054	0.219 ± 0.108
GCN	NOCD	None	\checkmark	MLP	MCC	0.317 ± 0.182	0.367 ± 0.201
GCN	NOCD	DMoN	\checkmark	MLP	MCC	0.373 ± 0.143	0.414 ± 0.154
GCN	NOCD	L2	\checkmark	MLP	MCC	0.057 ± 0.074	0.173 ± 0.126
GCN	Neuromap	None	\times	None	MCC	0.055 ± 0.066	0.086 ± 0.071
GCN	Neuromap	DMoN	\checkmark	MLP	MCC	0.016 ± 0.030	0.095 ± 0.102
GCN	Neuromap	L2	\checkmark	MLP	MCC	0.015 ± 0.020	0.130 ± 0.099
GCN	SBM_{NN}	None	\checkmark	MLP	MCC	0.275 ± 0.124	0.326 ± 0.162
GCN	SBM_{NN}	DMoN	\checkmark	Transformer	MCC	0.326 ± 0.151	0.383 ± 0.180
GCN	SBM_{NN}	L2	\times	None	MCC	0.041 ± 0.058	0.174 ± 0.111
GraphSAGE	None	None	\times	None	MCC	0.373 ± 0.063	0.430 ± 0.087
GraphSAGE	None	DMoN	\times	None	Loss	0.466 ± 0.105	0.522 ± 0.114
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.063 ± 0.052
GraphSAGE	DMoN	None	\times	None	MCC	0.484 ± 0.110	0.550 ± 0.121
GraphSAGE	DMoN	DMoN	\times	None	MCC	0.446 ± 0.105	0.518 ± 0.099
GraphSAGE	DMoN	L2	\times	None	MCC	-0.003 ± 0.006	0.071 ± 0.121

Continued on next page

Table 86: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✓	MLP	MCC	0.338 ± 0.113	0.355 ± 0.140
GraphSAGE	NOCD	DMoN	✓	Transformer	MCC	0.391 ± 0.122	0.455 ± 0.139
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.073 ± 0.049
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.052 ± 0.072	0.111 ± 0.107
GraphSAGE	Neuromap	DMoN	✗	None	MCC	0.037 ± 0.059	0.143 ± 0.145
GraphSAGE	Neuromap	L2	✗	None	MCC	0.005 ± 0.019	0.079 ± 0.095
GraphSAGE	SBM _{NN}	None	✗	None	MCC	0.127 ± 0.144	0.178 ± 0.160
GraphSAGE	SBM _{NN}	DMoN	✓	MLP	MCC	0.094 ± 0.117	0.140 ± 0.138
GraphSAGE	SBM _{NN}	L2	✓	Transformer	MCC	0.002 ± 0.012	0.118 ± 0.137
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.075 ± 0.051
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.120 ± 0.107
Transformer	None	L2	N/A	None	Loss	-0.000 ± 0.000	0.078 ± 0.059
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.133 ± 0.138
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.090 ± 0.109
Transformer	DMoN	L2	N/A	None	MCC	0.014 ± 0.028	0.161 ± 0.110
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.032 ± 0.004
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.050 ± 0.044
Transformer	NOCD	L2	N/A	None	Loss	0.106 ± 0.093	0.177 ± 0.105
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.122 ± 0.136
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.073 ± 0.041
Transformer	Neuromap	L2	N/A	None	MCC	0.002 ± 0.011	0.075 ± 0.060
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.031 ± 0.003
Transformer	SBM _{NN}	DMoN	N/A	None	MCC	-0.001 ± 0.002	0.036 ± 0.115
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.006 ± 0.019	0.099 ± 0.108
MLP	None	None	N/A	None	MCC	0.180 ± 0.067	0.269 ± 0.099
MLP	None	DMoN	N/A	None	MCC	0.222 ± 0.034	0.297 ± 0.041
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.098 ± 0.111
MLP	DMoN	None	N/A	None	MCC	0.201 ± 0.037	0.266 ± 0.043
MLP	DMoN	DMoN	N/A	None	MCC	0.207 ± 0.058	0.297 ± 0.066
MLP	DMoN	L2	N/A	None	MCC	0.001 ± 0.024	0.134 ± 0.108
MLP	NOCD	None	N/A	None	Loss	0.003 ± 0.005	0.040 ± 0.024
MLP	NOCD	DMoN	N/A	None	Loss	0.002 ± 0.005	0.040 ± 0.023
MLP	NOCD	L2	N/A	None	Loss	0.108 ± 0.076	0.166 ± 0.123
MLP	Neuromap	None	N/A	None	MCC	0.005 ± 0.011	0.083 ± 0.119
MLP	Neuromap	DMoN	N/A	None	MCC	0.020 ± 0.029	0.056 ± 0.073
MLP	Neuromap	L2	N/A	None	MCC	0.002 ± 0.017	0.082 ± 0.072
MLP	SBM _{NN}	None	N/A	None	Loss	-0.000 ± 0.000	0.097 ± 0.110
MLP	SBM _{NN}	DMoN	N/A	None	MCC	-0.000 ± 0.002	0.009 ± 0.019
MLP	SBM _{NN}	L2	N/A	None	MCC	0.007 ± 0.015	0.068 ± 0.073

F.4.7 AMAZON PHOTO DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 87: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	Loss	0.785 ± 0.081	0.803 ± 0.088
GCN	None	DMoN	✗	None	Loss	0.799 ± 0.075	0.819 ± 0.081
GCN	None	L2	✗	None	Loss	0.000 ± 0.000	0.047 ± 0.001
GCN	DMoN	None	✓	MLP	Loss	0.796 ± 0.070	0.817 ± 0.077
GCN	DMoN	DMoN	✓	Transformer	Loss	0.840 ± 0.031	0.862 ± 0.028
GCN	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.047 ± 0.000
GCN	NOCD	None	✓	MLP	Loss	0.667 ± 0.113	0.683 ± 0.128
GCN	NOCD	DMoN	✓	Transformer	Loss	0.706 ± 0.089	0.722 ± 0.105
GCN	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.086 ± 0.031
GCN	Neuromap	None	✗	None	Loss	-0.002 ± 0.005	0.101 ± 0.059
GCN	Neuromap	DMoN	✓	Transformer	Loss	-0.001 ± 0.004	0.145 ± 0.084
GCN	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.047 ± 0.001
GCN	SBM _{NN}	None	✓	MLP	Loss	0.504 ± 0.075	0.482 ± 0.086
GCN	SBM _{NN}	DMoN	✓	MLP	Loss	0.534 ± 0.089	0.526 ± 0.093
GCN	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.109 ± 0.073
GraphSAGE	None	None	✗	None	Loss	0.830 ± 0.059	0.850 ± 0.062
GraphSAGE	None	DMoN	✗	None	Loss	0.816 ± 0.052	0.838 ± 0.055
GraphSAGE	None	L2	✗	None	Loss	0.000 ± 0.000	0.099 ± 0.053
GraphSAGE	DMoN	None	✓	MLP	Loss	0.828 ± 0.030	0.851 ± 0.027
GraphSAGE	DMoN	DMoN	✓	Transformer	Loss	0.830 ± 0.026	0.853 ± 0.024
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.102 ± 0.059
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.461 ± 0.065	0.439 ± 0.075
GraphSAGE	NOCD	DMoN	✓	Transformer	Loss	0.314 ± 0.054	0.251 ± 0.058
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.141 ± 0.075
GraphSAGE	Neuromap	None	✓	Transformer	Loss	0.004 ± 0.011	0.101 ± 0.052
GraphSAGE	Neuromap	DMoN	✓	Transformer	Loss	0.020 ± 0.047	0.171 ± 0.082
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.110 ± 0.067
GraphSAGE	SBM _{NN}	None	✓	MLP	Loss	0.172 ± 0.108	0.144 ± 0.089
GraphSAGE	SBM _{NN}	DMoN	✓	MLP	Loss	0.184 ± 0.088	0.151 ± 0.063
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.104 ± 0.054
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.149 ± 0.060
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.097 ± 0.064
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.162 ± 0.068
Transformer	DMoN	None	N/A	None	Loss	0.004 ± 0.014	0.112 ± 0.074
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.154 ± 0.072
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.152 ± 0.062
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.072 ± 0.066
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.046 ± 0.002
Transformer	NOCD	L2	N/A	None	Loss	0.071 ± 0.113	0.155 ± 0.110
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.110 ± 0.045
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.114 ± 0.073
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.094 ± 0.029

Continued on next page

Table 87: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.047 ± 0.001
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.047 ± 0.001
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.007 ± 0.019	0.117 ± 0.072
MLP	None	None	N/A	None	Loss	0.680 ± 0.021	0.720 ± 0.020
MLP	None	DMoN	N/A	None	Loss	0.564 ± 0.055	0.603 ± 0.051
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.116 ± 0.071
MLP	DMoN	None	N/A	None	Loss	0.618 ± 0.048	0.652 ± 0.052
MLP	DMoN	DMoN	N/A	None	Loss	0.585 ± 0.053	0.623 ± 0.054
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.078
MLP	NOCD	None	N/A	None	Loss	0.018 ± 0.058	0.058 ± 0.024
MLP	NOCD	DMoN	N/A	None	Loss	0.001 ± 0.003	0.067 ± 0.067
MLP	NOCD	L2	N/A	None	Loss	0.147 ± 0.095	0.185 ± 0.097
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.001	0.148 ± 0.091
MLP	Neuromap	DMoN	N/A	None	Loss	-0.000 ± 0.000	0.181 ± 0.079
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.098 ± 0.053
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.064 ± 0.056
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.067 ± 0.066
MLP	SBM_{NN}	L2	N/A	None	Loss	0.004 ± 0.014	0.055 ± 0.024

Table 88: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.785 ± 0.081	0.803 ± 0.088
GCN	None	DMoN	\times	None	Loss	0.799 ± 0.075	0.819 ± 0.081
GCN	None	L2	\times	None	MCC	0.145 ± 0.133	0.254 ± 0.086
GCN	DMoN	None	\checkmark	MLP	Loss	0.796 ± 0.070	0.817 ± 0.077
GCN	DMoN	DMoN	\checkmark	Transformer	Loss	0.840 ± 0.031	0.862 ± 0.028
GCN	DMoN	L2	\checkmark	Transformer	MCC	0.047 ± 0.098	0.158 ± 0.047
GCN	NOCD	None	\checkmark	MLP	Loss	0.667 ± 0.113	0.683 ± 0.128
GCN	NOCD	DMoN	\checkmark	Transformer	Loss	0.706 ± 0.089	0.722 ± 0.105
GCN	NOCD	L2	\times	None	MCC	0.168 ± 0.065	0.216 ± 0.081
GCN	Neuromap	None	\times	None	MCC	0.151 ± 0.125	0.165 ± 0.134
GCN	Neuromap	DMoN	\times	None	MCC	0.147 ± 0.071	0.147 ± 0.031
GCN	Neuromap	L2	\times	None	MCC	0.068 ± 0.165	0.192 ± 0.123
GCN	SBM_{NN}	None	\checkmark	MLP	Loss	0.504 ± 0.075	0.482 ± 0.086
GCN	SBM_{NN}	DMoN	\checkmark	MLP	Loss	0.534 ± 0.089	0.526 ± 0.093
GCN	SBM_{NN}	L2	\times	None	MCC	0.093 ± 0.087	0.179 ± 0.058
GraphSAGE	None	None	\times	None	Loss	0.830 ± 0.059	0.850 ± 0.062
GraphSAGE	None	DMoN	\times	None	Loss	0.816 ± 0.052	0.838 ± 0.055
GraphSAGE	None	L2	\times	None	MCC	-0.000 ± 0.009	0.122 ± 0.075
GraphSAGE	DMoN	None	\checkmark	MLP	Loss	0.828 ± 0.030	0.851 ± 0.027
GraphSAGE	DMoN	DMoN	\checkmark	Transformer	Loss	0.830 ± 0.026	0.853 ± 0.024
GraphSAGE	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.102 ± 0.059

Continued on next page

Table 88: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.461 ± 0.065	0.439 ± 0.075
GraphSAGE	NOCD	DMoN	✓	Transformer	Loss	0.314 ± 0.054	0.251 ± 0.058
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.141 ± 0.075
GraphSAGE	Neuromap	None	✗	None	MCC	0.054 ± 0.087	0.122 ± 0.100
GraphSAGE	Neuromap	DMoN	✗	None	MCC	0.093 ± 0.102	0.147 ± 0.112
GraphSAGE	Neuromap	L2	✗	None	MCC	-0.004 ± 0.019	0.123 ± 0.087
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.196 ± 0.102	0.172 ± 0.095
GraphSAGE	SBM _{NN}	DMoN	✗	None	MCC	0.324 ± 0.139	0.285 ± 0.143
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.101 ± 0.023
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.149 ± 0.060
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.097 ± 0.064
Transformer	None	L2	N/A	None	MCC	0.093 ± 0.032	0.198 ± 0.054
Transformer	DMoN	None	N/A	None	Loss	0.004 ± 0.014	0.112 ± 0.074
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.154 ± 0.072
Transformer	DMoN	L2	N/A	None	MCC	0.085 ± 0.068	0.244 ± 0.058
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.072 ± 0.066
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.046 ± 0.002
Transformer	NOCD	L2	N/A	None	Loss	0.071 ± 0.113	0.155 ± 0.110
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.110 ± 0.045
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.114 ± 0.073
Transformer	Neuromap	L2	N/A	None	MCC	0.004 ± 0.016	0.120 ± 0.077
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.047 ± 0.001
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.047 ± 0.001
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.007 ± 0.019	0.117 ± 0.072
MLP	None	None	N/A	None	Loss	0.680 ± 0.021	0.720 ± 0.020
MLP	None	DMoN	N/A	None	Loss	0.564 ± 0.055	0.603 ± 0.051
MLP	None	L2	N/A	None	MCC	0.018 ± 0.063	0.192 ± 0.045
MLP	DMoN	None	N/A	None	Loss	0.618 ± 0.048	0.652 ± 0.052
MLP	DMoN	DMoN	N/A	None	Loss	0.585 ± 0.053	0.623 ± 0.054
MLP	DMoN	L2	N/A	None	MCC	-0.008 ± 0.065	0.146 ± 0.057
MLP	NOCD	None	N/A	None	MCC	0.010 ± 0.020	0.069 ± 0.082
MLP	NOCD	DMoN	N/A	None	MCC	0.009 ± 0.018	0.090 ± 0.073
MLP	NOCD	L2	N/A	None	Loss	0.147 ± 0.095	0.185 ± 0.097
MLP	Neuromap	None	N/A	None	MCC	0.013 ± 0.012	0.075 ± 0.057
MLP	Neuromap	DMoN	N/A	None	MCC	0.026 ± 0.043	0.092 ± 0.093
MLP	Neuromap	L2	N/A	None	MCC	0.002 ± 0.026	0.120 ± 0.062
MLP	SBM _{NN}	None	N/A	None	MCC	0.003 ± 0.004	0.091 ± 0.105
MLP	SBM _{NN}	DMoN	N/A	None	MCC	-0.000 ± 0.012	0.067 ± 0.078
MLP	SBM _{NN}	L2	N/A	None	Loss	0.004 ± 0.014	0.055 ± 0.024

Table 89: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	MCC	0.804 ± 0.069	0.826 ± 0.071
GCN	None	DMoN	✗	None	MCC	0.796 ± 0.073	0.815 ± 0.081
GCN	None	L2	✗	None	MCC	0.145 ± 0.133	0.254 ± 0.086
GCN	DMoN	None	✓	MLP	Loss	0.796 ± 0.070	0.817 ± 0.077
GCN	DMoN	DMoN	✓	Transformer	Loss	0.840 ± 0.031	0.862 ± 0.028
GCN	DMoN	L2	✓	Transformer	MCC	0.047 ± 0.098	0.158 ± 0.047
GCN	NOCD	None	✓	MLP	MCC	0.670 ± 0.080	0.695 ± 0.088
GCN	NOCD	DMoN	✓	Transformer	Loss	0.706 ± 0.089	0.722 ± 0.105
GCN	NOCD	L2	✗	None	MCC	0.168 ± 0.065	0.216 ± 0.081
GCN	Neuromap	None	✗	None	MCC	0.151 ± 0.125	0.165 ± 0.134
GCN	Neuromap	DMoN	✗	None	MCC	0.147 ± 0.071	0.147 ± 0.031
GCN	Neuromap	L2	✗	None	MCC	0.068 ± 0.165	0.192 ± 0.123
GCN	SBM _{NN}	None	✓	Transformer	MCC	0.575 ± 0.144	0.582 ± 0.163
GCN	SBM _{NN}	DMoN	✓	MLP	Loss	0.534 ± 0.089	0.526 ± 0.093
GCN	SBM _{NN}	L2	✓	Transformer	MCC	0.088 ± 0.080	0.162 ± 0.088
GraphSAGE	None	None	✗	None	Loss	0.830 ± 0.059	0.850 ± 0.062
GraphSAGE	None	DMoN	✗	None	Loss	0.816 ± 0.052	0.838 ± 0.055
GraphSAGE	None	L2	✗	None	MCC	-0.000 ± 0.009	0.122 ± 0.075
GraphSAGE	DMoN	None	✓	MLP	Loss	0.828 ± 0.030	0.851 ± 0.027
GraphSAGE	DMoN	DMoN	✓	Transformer	Loss	0.830 ± 0.026	0.853 ± 0.024
GraphSAGE	DMoN	L2	✓	Transformer	MCC	0.005 ± 0.014	0.085 ± 0.075
GraphSAGE	NOCD	None	✗	None	MCC	0.641 ± 0.150	0.651 ± 0.163
GraphSAGE	NOCD	DMoN	✓	MLP	MCC	0.694 ± 0.110	0.711 ± 0.116
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.008 ± 0.018	0.129 ± 0.088
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.081 ± 0.096	0.139 ± 0.111
GraphSAGE	Neuromap	DMoN	✗	None	MCC	0.093 ± 0.102	0.147 ± 0.112
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.110 ± 0.067
GraphSAGE	SBM _{NN}	None	✓	MLP	MCC	0.302 ± 0.144	0.258 ± 0.131
GraphSAGE	SBM _{NN}	DMoN	✗	None	MCC	0.324 ± 0.139	0.285 ± 0.143
GraphSAGE	SBM _{NN}	L2	✓	MLP	MCC	-0.001 ± 0.004	0.068 ± 0.079
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.149 ± 0.060
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.097 ± 0.064
Transformer	None	L2	N/A	None	MCC	0.093 ± 0.032	0.198 ± 0.054
Transformer	DMoN	None	N/A	None	Loss	0.004 ± 0.014	0.112 ± 0.074
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.154 ± 0.072
Transformer	DMoN	L2	N/A	None	MCC	0.085 ± 0.068	0.244 ± 0.058
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.072 ± 0.066
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.046 ± 0.002
Transformer	NOCD	L2	N/A	None	Loss	0.071 ± 0.113	0.155 ± 0.110
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.110 ± 0.045
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.114 ± 0.073
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.094 ± 0.029

Continued on next page

Table 89: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.047 ± 0.001
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.047 ± 0.001
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.007 ± 0.019	0.117 ± 0.072
MLP	None	None	N/A	None	MCC	0.698 ± 0.024	0.739 ± 0.023
MLP	None	DMoN	N/A	None	MCC	0.696 ± 0.025	0.741 ± 0.021
MLP	None	L2	N/A	None	MCC	0.018 ± 0.063	0.192 ± 0.045
MLP	DMoN	None	N/A	None	MCC	0.697 ± 0.031	0.740 ± 0.027
MLP	DMoN	DMoN	N/A	None	MCC	0.684 ± 0.023	0.728 ± 0.021
MLP	DMoN	L2	N/A	None	MCC	-0.008 ± 0.065	0.146 ± 0.057
MLP	NOCD	None	N/A	None	Loss	0.018 ± 0.058	0.058 ± 0.024
MLP	NOCD	DMoN	N/A	None	MCC	0.009 ± 0.018	0.090 ± 0.073
MLP	NOCD	L2	N/A	None	Loss	0.147 ± 0.095	0.185 ± 0.097
MLP	Neuromap	None	N/A	None	MCC	0.013 ± 0.012	0.075 ± 0.057
MLP	Neuromap	DMoN	N/A	None	MCC	0.026 ± 0.043	0.092 ± 0.093
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.098 ± 0.053
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.064 ± 0.056
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.067 ± 0.066
MLP	SBM_{NN}	L2	N/A	None	Loss	0.004 ± 0.014	0.055 ± 0.024

F.4.8 AMAZON PHOTO DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 90: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.616 ± 0.099	0.647 ± 0.112
GCN	None	DMoN	\times	None	Loss	0.644 ± 0.088	0.679 ± 0.088
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.048 ± 0.000
GCN	DMoN	None	\times	None	Loss	0.608 ± 0.098	0.635 ± 0.104
GCN	DMoN	DMoN	\times	None	Loss	0.700 ± 0.075	0.736 ± 0.068
GCN	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.048 ± 0.000
GCN	NOCD	None	\checkmark	MLP	Loss	0.361 ± 0.124	0.355 ± 0.150
GCN	NOCD	DMoN	\checkmark	MLP	Loss	0.374 ± 0.044	0.365 ± 0.039
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.141 ± 0.073
GCN	Neuromap	None	\checkmark	Transformer	Loss	-0.000 ± 0.004	0.159 ± 0.076
GCN	Neuromap	DMoN	\checkmark	Transformer	Loss	0.023 ± 0.072	0.157 ± 0.083
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.048 ± 0.000
GCN	SBM_{NN}	None	\times	None	Loss	0.218 ± 0.073	0.203 ± 0.060
GCN	SBM_{NN}	DMoN	\times	None	Loss	0.258 ± 0.109	0.241 ± 0.114
GCN	SBM_{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.134 ± 0.080
GraphSAGE	None	None	\times	None	Loss	0.640 ± 0.068	0.675 ± 0.072
GraphSAGE	None	DMoN	\times	None	Loss	0.637 ± 0.089	0.676 ± 0.083
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.105 ± 0.075

Continued on next page

Table 90: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	DMoN	None	✗	None	Loss	0.667 ± 0.068	0.702 ± 0.070
GraphSAGE	DMoN	DMoN	✗	None	Loss	0.642 ± 0.082	0.681 ± 0.080
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.117 ± 0.053
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.389 ± 0.107	0.402 ± 0.110
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.252 ± 0.082	0.205 ± 0.088
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.098 ± 0.048
GraphSAGE	Neuromap	None	✓	MLP	Loss	0.000 ± 0.000	0.184 ± 0.070
GraphSAGE	Neuromap	DMoN	✗	None	Loss	-0.002 ± 0.006	0.135 ± 0.044
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.180 ± 0.091
GraphSAGE	SBM _{NN}	None	✓	MLP	Loss	0.119 ± 0.101	0.130 ± 0.064
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.098 ± 0.112	0.108 ± 0.089
GraphSAGE	SBM _{NN}	L2	✓	MLP	Loss	0.000 ± 0.000	0.130 ± 0.069
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.159 ± 0.075
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.105 ± 0.078
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.145 ± 0.078
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.170 ± 0.070
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.094 ± 0.033
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.128 ± 0.077
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.067 ± 0.030
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.096 ± 0.086
Transformer	NOCD	L2	N/A	None	Loss	0.155 ± 0.138	0.232 ± 0.137
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.157 ± 0.086
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.096 ± 0.054
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.116 ± 0.042
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.075 ± 0.066
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.093 ± 0.086
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.001 ± 0.004	0.125 ± 0.066
MLP	None	None	N/A	None	Loss	0.205 ± 0.099	0.276 ± 0.114
MLP	None	DMoN	N/A	None	Loss	0.235 ± 0.063	0.288 ± 0.062
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.089 ± 0.037
MLP	DMoN	None	N/A	None	Loss	0.235 ± 0.049	0.301 ± 0.061
MLP	DMoN	DMoN	N/A	None	MCC	0.396 ± 0.038	0.469 ± 0.034
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.154 ± 0.075
MLP	NOCD	None	N/A	None	Loss	0.037 ± 0.078	0.123 ± 0.106
MLP	NOCD	DMoN	N/A	None	Loss	0.012 ± 0.024	0.084 ± 0.066
MLP	NOCD	L2	N/A	None	Loss	0.202 ± 0.115	0.250 ± 0.115
MLP	Neuromap	None	N/A	None	Loss	0.004 ± 0.013	0.143 ± 0.066
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.134 ± 0.079
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.103 ± 0.052
MLP	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.048 ± 0.000
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.052 ± 0.014
MLP	SBM _{NN}	L2	N/A	None	Loss	0.005 ± 0.017	0.073 ± 0.034

Table 91: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.616 ± 0.099	0.647 ± 0.112
GCN	None	DMoN	\times	None	Loss	0.644 ± 0.088	0.679 ± 0.088
GCN	None	L2	\times	None	MCC	0.039 ± 0.169	0.163 ± 0.130
GCN	DMoN	None	\times	None	Loss	0.608 ± 0.098	0.635 ± 0.104
GCN	DMoN	DMoN	\times	None	Loss	0.700 ± 0.075	0.736 ± 0.068
GCN	DMoN	L2	\times	None	MCC	0.033 ± 0.174	0.184 ± 0.112
GCN	NOCD	None	\checkmark	MLP	Loss	0.361 ± 0.124	0.355 ± 0.150
GCN	NOCD	DMoN	\checkmark	MLP	Loss	0.374 ± 0.044	0.365 ± 0.039
GCN	NOCD	L2	\times	None	MCC	0.107 ± 0.115	0.201 ± 0.101
GCN	Neuromap	None	\times	None	MCC	0.102 ± 0.110	0.129 ± 0.142
GCN	Neuromap	DMoN	\checkmark	Transformer	MCC	0.055 ± 0.049	0.143 ± 0.084
GCN	Neuromap	L2	\times	None	MCC	0.054 ± 0.154	0.208 ± 0.105
GCN	SBM _{NN}	None	\checkmark	MLP	Loss	0.263 ± 0.086	0.251 ± 0.088
GCN	SBM _{NN}	DMoN	\times	None	Loss	0.234 ± 0.097	0.203 ± 0.094
GCN	SBM _{NN}	L2	\times	None	MCC	0.045 ± 0.053	0.144 ± 0.050
GraphSAGE	None	None	\times	None	Loss	0.640 ± 0.068	0.675 ± 0.072
GraphSAGE	None	DMoN	\times	None	Loss	0.637 ± 0.089	0.676 ± 0.083
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.105 ± 0.075
GraphSAGE	DMoN	None	\times	None	Loss	0.667 ± 0.068	0.702 ± 0.070
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.642 ± 0.082	0.681 ± 0.080
GraphSAGE	DMoN	L2	\checkmark	Transformer	MCC	0.001 ± 0.002	0.058 ± 0.083
GraphSAGE	NOCD	None	\checkmark	Transformer	Loss	0.389 ± 0.107	0.402 ± 0.110
GraphSAGE	NOCD	DMoN	\checkmark	MLP	MCC	0.496 ± 0.103	0.521 ± 0.116
GraphSAGE	NOCD	L2	\checkmark	MLP	MCC	-0.006 ± 0.018	0.064 ± 0.089
GraphSAGE	Neuromap	None	\checkmark	MLP	MCC	0.072 ± 0.077	0.097 ± 0.093
GraphSAGE	Neuromap	DMoN	\checkmark	MLP	MCC	0.040 ± 0.071	0.148 ± 0.124
GraphSAGE	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.180 ± 0.091
GraphSAGE	SBM _{NN}	None	\checkmark	MLP	Loss	0.119 ± 0.101	0.130 ± 0.064
GraphSAGE	SBM _{NN}	DMoN	\checkmark	MLP	MCC	0.185 ± 0.149	0.187 ± 0.133
GraphSAGE	SBM _{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.121 ± 0.060
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.159 ± 0.075
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.105 ± 0.078
Transformer	None	L2	N/A	None	MCC	0.016 ± 0.024	0.130 ± 0.079
Transformer	DMoN	None	N/A	None	MCC	0.001 ± 0.004	0.011 ± 0.035
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.094 ± 0.033
Transformer	DMoN	L2	N/A	None	MCC	0.012 ± 0.023	0.138 ± 0.057
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.067 ± 0.030
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.096 ± 0.086
Transformer	NOCD	L2	N/A	None	Loss	0.155 ± 0.138	0.232 ± 0.137
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.157 ± 0.086
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.096 ± 0.054
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.116 ± 0.042

Continued on next page

Table 91: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.075 ± 0.066
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.093 ± 0.086
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.001 ± 0.004	0.125 ± 0.066
MLP	None	None	N/A	None	Loss	0.205 ± 0.099	0.276 ± 0.114
MLP	None	DMoN	N/A	None	Loss	0.235 ± 0.063	0.288 ± 0.062
MLP	None	L2	N/A	None	MCC	0.008 ± 0.027	0.157 ± 0.070
MLP	DMoN	None	N/A	None	Loss	0.235 ± 0.049	0.301 ± 0.061
MLP	DMoN	DMoN	N/A	None	MCC	0.396 ± 0.038	0.469 ± 0.034
MLP	DMoN	L2	N/A	None	MCC	0.013 ± 0.034	0.176 ± 0.077
MLP	NOCD	None	N/A	None	MCC	0.011 ± 0.020	0.063 ± 0.099
MLP	NOCD	DMoN	N/A	None	Loss	0.012 ± 0.024	0.084 ± 0.066
MLP	NOCD	L2	N/A	None	Loss	0.202 ± 0.115	0.250 ± 0.115
MLP	Neuromap	None	N/A	None	MCC	0.017 ± 0.016	0.099 ± 0.085
MLP	Neuromap	DMoN	N/A	None	MCC	0.002 ± 0.011	0.031 ± 0.043
MLP	Neuromap	L2	N/A	None	MCC	0.008 ± 0.040	0.136 ± 0.072
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.048 ± 0.000
MLP	SBM_{NN}	DMoN	N/A	None	MCC	0.001 ± 0.018	0.063 ± 0.089
MLP	SBM_{NN}	L2	N/A	None	Loss	0.005 ± 0.017	0.073 ± 0.034

Table 92: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	MCC	0.620 ± 0.078	0.653 ± 0.087
GCN	None	DMoN	✗	None	Loss	0.644 ± 0.088	0.679 ± 0.088
GCN	None	L2	✗	None	MCC	0.039 ± 0.169	0.163 ± 0.130
GCN	DMoN	None	✗	None	MCC	0.618 ± 0.093	0.659 ± 0.098
GCN	DMoN	DMoN	✗	None	MCC	0.640 ± 0.074	0.680 ± 0.076
GCN	DMoN	L2	✗	None	MCC	0.033 ± 0.174	0.184 ± 0.112
GCN	NOCD	None	✓	MLP	MCC	0.489 ± 0.126	0.522 ± 0.127
GCN	NOCD	DMoN	✓	Transformer	MCC	0.502 ± 0.169	0.541 ± 0.176
GCN	NOCD	L2	✗	None	MCC	0.107 ± 0.115	0.201 ± 0.101
GCN	Neuromap	None	✗	None	MCC	0.102 ± 0.110	0.129 ± 0.142
GCN	Neuromap	DMoN	✓	Transformer	MCC	0.055 ± 0.049	0.143 ± 0.084
GCN	Neuromap	L2	✗	None	MCC	0.054 ± 0.154	0.208 ± 0.105
GCN	SBM_{NN}	None	✓	MLP	MCC	0.437 ± 0.202	0.471 ± 0.199
GCN	SBM_{NN}	DMoN	✗	None	MCC	0.433 ± 0.106	0.453 ± 0.118
GCN	SBM_{NN}	L2	✓	Transformer	MCC	0.075 ± 0.054	0.178 ± 0.068
GraphSAGE	None	None	✗	None	MCC	0.603 ± 0.104	0.641 ± 0.110
GraphSAGE	None	DMoN	✗	None	MCC	0.626 ± 0.089	0.671 ± 0.086
GraphSAGE	None	L2	✗	None	Loss	0.000 ± 0.000	0.105 ± 0.075
GraphSAGE	DMoN	None	✗	None	MCC	0.581 ± 0.048	0.617 ± 0.052
GraphSAGE	DMoN	DMoN	✗	None	MCC	0.571 ± 0.102	0.618 ± 0.105
GraphSAGE	DMoN	L2	✓	Transformer	MCC	0.001 ± 0.002	0.058 ± 0.083

Continued on next page

Table 92: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✓	MLP	MCC	0.480 ± 0.091	0.505 ± 0.091
GraphSAGE	NOCD	DMoN	✓	MLP	MCC	0.496 ± 0.103	0.521 ± 0.116
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.098 ± 0.048
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.072 ± 0.077	0.097 ± 0.093
GraphSAGE	Neuromap	DMoN	✗	None	MCC	0.054 ± 0.087	0.129 ± 0.131
GraphSAGE	Neuromap	L2	✓	MLP	MCC	0.002 ± 0.009	0.041 ± 0.071
GraphSAGE	SBM _{NN}	None	✗	None	MCC	0.138 ± 0.116	0.163 ± 0.102
GraphSAGE	SBM _{NN}	DMoN	✓	MLP	MCC	0.185 ± 0.149	0.187 ± 0.133
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.121 ± 0.060
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.159 ± 0.075
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.105 ± 0.078
Transformer	None	L2	N/A	None	MCC	0.016 ± 0.024	0.130 ± 0.079
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.170 ± 0.070
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.094 ± 0.033
Transformer	DMoN	L2	N/A	None	MCC	0.012 ± 0.023	0.138 ± 0.057
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.067 ± 0.030
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.096 ± 0.086
Transformer	NOCD	L2	N/A	None	Loss	0.155 ± 0.138	0.232 ± 0.137
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.157 ± 0.086
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.096 ± 0.054
Transformer	Neuromap	L2	N/A	None	MCC	-0.001 ± 0.022	0.097 ± 0.077
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.075 ± 0.066
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.093 ± 0.086
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.001 ± 0.004	0.125 ± 0.066
MLP	None	None	N/A	None	MCC	0.269 ± 0.073	0.362 ± 0.089
MLP	None	DMoN	N/A	None	MCC	0.399 ± 0.052	0.474 ± 0.044
MLP	None	L2	N/A	None	MCC	0.008 ± 0.027	0.157 ± 0.070
MLP	DMoN	None	N/A	None	MCC	0.387 ± 0.019	0.461 ± 0.016
MLP	DMoN	DMoN	N/A	None	MCC	0.396 ± 0.038	0.469 ± 0.034
MLP	DMoN	L2	N/A	None	MCC	0.013 ± 0.034	0.176 ± 0.077
MLP	NOCD	None	N/A	None	Loss	0.037 ± 0.078	0.123 ± 0.106
MLP	NOCD	DMoN	N/A	None	Loss	0.012 ± 0.024	0.084 ± 0.066
MLP	NOCD	L2	N/A	None	Loss	0.202 ± 0.115	0.250 ± 0.115
MLP	Neuromap	None	N/A	None	Loss	0.004 ± 0.013	0.143 ± 0.066
MLP	Neuromap	DMoN	N/A	None	MCC	0.002 ± 0.011	0.031 ± 0.043
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.103 ± 0.052
MLP	SBM _{NN}	None	N/A	None	MCC	-0.005 ± 0.020	0.077 ± 0.092
MLP	SBM _{NN}	DMoN	N/A	None	MCC	0.001 ± 0.018	0.063 ± 0.089
MLP	SBM _{NN}	L2	N/A	None	Loss	0.005 ± 0.017	0.073 ± 0.034

F.4.9 COAUTHOR PHYSICS DATASET, DEFAULT SPLIT WITH 20 TRAIN NODES PER CLASS AND 500 VALIDATION NODES.

Table 93: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	Loss	0.891 ± 0.015	0.925 ± 0.012
GCN	None	DMoN	✗	None	Loss	0.706 ± 0.019	0.752 ± 0.019
GCN	None	L2	✗	None	Loss	0.000 ± 0.000	0.167 ± 0.000
GCN	DMoN	None	✗	None	Loss	0.703 ± 0.031	0.750 ± 0.029
GCN	DMoN	DMoN	✗	None	Loss	0.700 ± 0.039	0.747 ± 0.038
GCN	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.167 ± 0.000
GCN	NOCD	None	✓	Transformer	Loss	0.436 ± 0.162	0.486 ± 0.184
GCN	NOCD	DMoN	✓	Transformer	Loss	0.298 ± 0.114	0.363 ± 0.138
GCN	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.185 ± 0.171
GCN	Neuromap	None	✗	None	Loss	0.361 ± 0.322	0.420 ± 0.333
GCN	Neuromap	DMoN	✓	MLP	Loss	0.134 ± 0.231	0.310 ± 0.201
GCN	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.167 ± 0.000
GCN	SBM _{NN}	None	✗	None	Loss	0.350 ± 0.148	0.418 ± 0.176
GCN	SBM _{NN}	DMoN	✗	None	Loss	0.336 ± 0.132	0.426 ± 0.137
GCN	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.211 ± 0.159
GraphSAGE	None	None	✗	None	Loss	0.899 ± 0.011	0.931 ± 0.008
GraphSAGE	None	DMoN	✗	None	Loss	0.520 ± 0.064	0.582 ± 0.048
GraphSAGE	None	L2	✗	None	Loss	0.000 ± 0.000	0.238 ± 0.187
GraphSAGE	DMoN	None	✓	MLP	Loss	0.580 ± 0.027	0.628 ± 0.025
GraphSAGE	DMoN	DMoN	✗	None	Loss	0.548 ± 0.041	0.605 ± 0.031
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.283 ± 0.194
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.508 ± 0.108	0.582 ± 0.130
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.150 ± 0.053	0.212 ± 0.024
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.209 ± 0.159
GraphSAGE	Neuromap	None	✓	MLP	Loss	0.123 ± 0.099	0.174 ± 0.068
GraphSAGE	Neuromap	DMoN	✓	MLP	Loss	0.124 ± 0.108	0.167 ± 0.071
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.175 ± 0.121
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.277 ± 0.141	0.349 ± 0.158
GraphSAGE	SBM _{NN}	DMoN	✓	MLP	Loss	0.332 ± 0.152	0.432 ± 0.140
GraphSAGE	SBM _{NN}	L2	✓	MLP	Loss	0.000 ± 0.000	0.170 ± 0.124
Transformer	None	None	N/A	None	Loss	0.032 ± 0.100	0.243 ± 0.200
Transformer	None	DMoN	N/A	None	Loss	0.021 ± 0.067	0.284 ± 0.200
Transformer	None	L2	N/A	None	Loss	0.590 ± 0.106	0.697 ± 0.095
Transformer	DMoN	None	N/A	None	Loss	0.034 ± 0.108	0.244 ± 0.194
Transformer	DMoN	DMoN	N/A	None	Loss	0.016 ± 0.051	0.162 ± 0.134
Transformer	DMoN	L2	N/A	None	Loss	0.642 ± 0.073	0.744 ± 0.062
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.194 ± 0.111
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.158 ± 0.028
Transformer	NOCD	L2	N/A	None	Loss	0.689 ± 0.246	0.745 ± 0.237
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.221 ± 0.198
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.259 ± 0.213
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.200 ± 0.164

Continued on next page

Table 93: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.411 ± 0.173	0.490 ± 0.205
MLP	None	None	N/A	None	Loss	0.820 ± 0.026	0.872 ± 0.023
MLP	None	DMoN	N/A	None	Loss	0.152 ± 0.082	0.238 ± 0.096
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.236 ± 0.189
MLP	DMoN	None	N/A	None	Loss	0.228 ± 0.097	0.248 ± 0.047
MLP	DMoN	DMoN	N/A	None	Loss	0.230 ± 0.072	0.263 ± 0.076
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.151 ± 0.130
MLP	NOCD	None	N/A	None	Loss	0.152 ± 0.138	0.254 ± 0.101
MLP	NOCD	DMoN	N/A	None	Loss	0.121 ± 0.153	0.244 ± 0.108
MLP	NOCD	L2	N/A	None	Loss	0.744 ± 0.198	0.826 ± 0.103
MLP	Neuromap	None	N/A	None	Loss	0.019 ± 0.053	0.206 ± 0.161
MLP	Neuromap	DMoN	N/A	None	Loss	0.049 ± 0.089	0.245 ± 0.198
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.219 ± 0.199
MLP	SBM_{NN}	None	N/A	None	Loss	0.025 ± 0.070	0.174 ± 0.021
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
MLP	SBM_{NN}	L2	N/A	None	Loss	0.016 ± 0.047	0.182 ± 0.123

Table 94: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.891 ± 0.015	0.925 ± 0.012
GCN	None	DMoN	\times	None	Loss	0.706 ± 0.019	0.752 ± 0.019
GCN	None	L2	\times	None	MCC	0.411 ± 0.078	0.599 ± 0.046
GCN	DMoN	None	\times	None	Loss	0.703 ± 0.031	0.750 ± 0.029
GCN	DMoN	DMoN	\times	None	Loss	0.700 ± 0.039	0.747 ± 0.038
GCN	DMoN	L2	\times	None	MCC	0.081 ± 0.333	0.275 ± 0.205
GCN	NOCD	None	\checkmark	Transformer	Loss	0.436 ± 0.162	0.486 ± 0.184
GCN	NOCD	DMoN	\checkmark	Transformer	Loss	0.298 ± 0.114	0.363 ± 0.138
GCN	NOCD	L2	\checkmark	MLP	MCC	0.248 ± 0.125	0.367 ± 0.167
GCN	Neuromap	None	\times	None	MCC	0.728 ± 0.135	0.774 ± 0.167
GCN	Neuromap	DMoN	\times	None	MCC	0.734 ± 0.105	0.781 ± 0.142
GCN	Neuromap	L2	\times	None	MCC	0.228 ± 0.316	0.425 ± 0.222
GCN	SBM_{NN}	None	\times	None	MCC	0.784 ± 0.156	0.808 ± 0.193
GCN	SBM_{NN}	DMoN	\times	None	Loss	0.336 ± 0.132	0.426 ± 0.137
GCN	SBM_{NN}	L2	\checkmark	MLP	MCC	0.234 ± 0.177	0.355 ± 0.210
GraphSAGE	None	None	\times	None	Loss	0.899 ± 0.011	0.931 ± 0.008
GraphSAGE	None	DMoN	\times	None	Loss	0.520 ± 0.064	0.582 ± 0.048
GraphSAGE	None	L2	\times	None	MCC	0.002 ± 0.014	0.067 ± 0.073
GraphSAGE	DMoN	None	\times	None	Loss	0.565 ± 0.043	0.615 ± 0.030
GraphSAGE	DMoN	DMoN	\times	None	Loss	0.548 ± 0.041	0.605 ± 0.031
GraphSAGE	DMoN	L2	\checkmark	Transformer	MCC	-0.003 ± 0.009	0.046 ± 0.146

Continued on next page

Table 94: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✓	Transformer	Loss	0.508 ± 0.108	0.582 ± 0.130
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.150 ± 0.053	0.212 ± 0.024
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.020 ± 0.037	0.045 ± 0.077
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.515 ± 0.238	0.537 ± 0.291
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.496 ± 0.199	0.489 ± 0.249
GraphSAGE	Neuromap	L2	✓	Transformer	MCC	0.003 ± 0.026	0.028 ± 0.045
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.277 ± 0.141	0.349 ± 0.158
GraphSAGE	SBM _{NN}	DMoN	✓	MLP	Loss	0.332 ± 0.152	0.432 ± 0.140
GraphSAGE	SBM _{NN}	L2	✗	None	MCC	0.019 ± 0.042	0.059 ± 0.131
Transformer	None	None	N/A	None	MCC	0.143 ± 0.176	0.201 ± 0.248
Transformer	None	DMoN	N/A	None	MCC	0.082 ± 0.119	0.078 ± 0.106
Transformer	None	L2	N/A	None	MCC	0.751 ± 0.092	0.821 ± 0.074
Transformer	DMoN	None	N/A	None	MCC	0.090 ± 0.126	0.160 ± 0.169
Transformer	DMoN	DMoN	N/A	None	MCC	0.184 ± 0.176	0.243 ± 0.234
Transformer	DMoN	L2	N/A	None	MCC	0.750 ± 0.073	0.829 ± 0.051
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.194 ± 0.111
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.158 ± 0.028
Transformer	NOCD	L2	N/A	None	Loss	0.689 ± 0.246	0.745 ± 0.237
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.221 ± 0.198
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.259 ± 0.213
Transformer	Neuromap	L2	N/A	None	MCC	0.013 ± 0.036	0.097 ± 0.060
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.411 ± 0.173	0.490 ± 0.205
MLP	None	None	N/A	None	Loss	0.820 ± 0.026	0.872 ± 0.023
MLP	None	DMoN	N/A	None	Loss	0.152 ± 0.082	0.238 ± 0.096
MLP	None	L2	N/A	None	MCC	0.364 ± 0.085	0.531 ± 0.115
MLP	DMoN	None	N/A	None	Loss	0.228 ± 0.097	0.248 ± 0.047
MLP	DMoN	DMoN	N/A	None	Loss	0.230 ± 0.072	0.263 ± 0.076
MLP	DMoN	L2	N/A	None	MCC	0.277 ± 0.139	0.490 ± 0.138
MLP	NOCD	None	N/A	None	Loss	0.152 ± 0.138	0.254 ± 0.101
MLP	NOCD	DMoN	N/A	None	MCC	0.303 ± 0.211	0.343 ± 0.256
MLP	NOCD	L2	N/A	None	Loss	0.744 ± 0.198	0.826 ± 0.103
MLP	Neuromap	None	N/A	None	MCC	0.247 ± 0.154	0.318 ± 0.239
MLP	Neuromap	DMoN	N/A	None	MCC	0.262 ± 0.159	0.298 ± 0.232
MLP	Neuromap	L2	N/A	None	MCC	-0.003 ± 0.051	0.215 ± 0.161
MLP	SBM _{NN}	None	N/A	None	Loss	0.025 ± 0.070	0.174 ± 0.021
MLP	SBM _{NN}	DMoN	N/A	None	MCC	0.020 ± 0.094	0.046 ± 0.091
MLP	SBM _{NN}	L2	N/A	None	Loss	0.016 ± 0.047	0.182 ± 0.123

Table 95: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	MCC	0.907 ± 0.009	0.936 ± 0.007
GCN	None	DMoN	✗	None	MCC	0.909 ± 0.012	0.938 ± 0.008
GCN	None	L2	✗	None	MCC	0.411 ± 0.078	0.599 ± 0.046
GCN	DMoN	None	✗	None	MCC	0.907 ± 0.008	0.936 ± 0.006
GCN	DMoN	DMoN	✗	None	MCC	0.902 ± 0.014	0.933 ± 0.009
GCN	DMoN	L2	✗	None	MCC	0.081 ± 0.333	0.275 ± 0.205
GCN	NOCD	None	✓	Transformer	MCC	0.716 ± 0.188	0.735 ± 0.231
GCN	NOCD	DMoN	✓	Transformer	MCC	0.738 ± 0.225	0.757 ± 0.253
GCN	NOCD	L2	✓	Transformer	MCC	0.254 ± 0.135	0.373 ± 0.175
GCN	Neuromap	None	✗	None	MCC	0.728 ± 0.135	0.774 ± 0.167
GCN	Neuromap	DMoN	✗	None	MCC	0.734 ± 0.105	0.781 ± 0.142
GCN	Neuromap	L2	✗	None	MCC	0.228 ± 0.316	0.425 ± 0.222
GCN	SBM _{NN}	None	✗	None	MCC	0.784 ± 0.156	0.808 ± 0.193
GCN	SBM _{NN}	DMoN	✓	Transformer	MCC	0.740 ± 0.213	0.769 ± 0.252
GCN	SBM _{NN}	L2	✓	MLP	MCC	0.234 ± 0.177	0.355 ± 0.210
GraphSAGE	None	None	✗	None	Loss	0.899 ± 0.011	0.931 ± 0.008
GraphSAGE	None	DMoN	✗	None	MCC	0.905 ± 0.006	0.935 ± 0.005
GraphSAGE	None	L2	✗	None	MCC	0.002 ± 0.014	0.067 ± 0.073
GraphSAGE	DMoN	None	✗	None	MCC	0.904 ± 0.007	0.934 ± 0.005
GraphSAGE	DMoN	DMoN	✓	MLP	MCC	0.899 ± 0.006	0.931 ± 0.004
GraphSAGE	DMoN	L2	✓	Transformer	MCC	-0.003 ± 0.009	0.046 ± 0.146
GraphSAGE	NOCD	None	✓	Transformer	MCC	0.756 ± 0.168	0.805 ± 0.183
GraphSAGE	NOCD	DMoN	✓	Transformer	MCC	0.766 ± 0.121	0.816 ± 0.139
GraphSAGE	NOCD	L2	✗	None	MCC	0.029 ± 0.056	0.087 ± 0.163
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.515 ± 0.238	0.537 ± 0.291
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.496 ± 0.199	0.489 ± 0.249
GraphSAGE	Neuromap	L2	✓	Transformer	MCC	0.003 ± 0.026	0.028 ± 0.045
GraphSAGE	SBM _{NN}	None	✓	Transformer	MCC	0.576 ± 0.180	0.622 ± 0.235
GraphSAGE	SBM _{NN}	DMoN	✗	None	MCC	0.514 ± 0.179	0.541 ± 0.233
GraphSAGE	SBM _{NN}	L2	✗	None	MCC	0.025 ± 0.069	0.060 ± 0.119
Transformer	None	None	N/A	None	MCC	0.143 ± 0.176	0.201 ± 0.248
Transformer	None	DMoN	N/A	None	MCC	0.082 ± 0.119	0.078 ± 0.106
Transformer	None	L2	N/A	None	MCC	0.751 ± 0.092	0.821 ± 0.074
Transformer	DMoN	None	N/A	None	MCC	0.090 ± 0.126	0.160 ± 0.169
Transformer	DMoN	DMoN	N/A	None	MCC	0.184 ± 0.176	0.243 ± 0.234
Transformer	DMoN	L2	N/A	None	MCC	0.750 ± 0.073	0.829 ± 0.051
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.194 ± 0.111
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.158 ± 0.028
Transformer	NOCD	L2	N/A	None	Loss	0.689 ± 0.246	0.745 ± 0.237
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.221 ± 0.198
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.259 ± 0.213
Transformer	Neuromap	L2	N/A	None	MCC	0.013 ± 0.036	0.097 ± 0.060

Continued on next page

Table 95: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.411 ± 0.173	0.490 ± 0.205
MLP	None	None	N/A	None	MCC	0.819 ± 0.023	0.874 ± 0.016
MLP	None	DMoN	N/A	None	MCC	0.721 ± 0.036	0.791 ± 0.033
MLP	None	L2	N/A	None	MCC	0.364 ± 0.085	0.531 ± 0.115
MLP	DMoN	None	N/A	None	MCC	0.732 ± 0.051	0.800 ± 0.047
MLP	DMoN	DMoN	N/A	None	MCC	0.732 ± 0.064	0.799 ± 0.056
MLP	DMoN	L2	N/A	None	MCC	0.277 ± 0.139	0.490 ± 0.138
MLP	NOCD	None	N/A	None	MCC	0.239 ± 0.229	0.286 ± 0.276
MLP	NOCD	DMoN	N/A	None	MCC	0.303 ± 0.211	0.343 ± 0.256
MLP	NOCD	L2	N/A	None	Loss	0.744 ± 0.198	0.826 ± 0.103
MLP	Neuromap	None	N/A	None	MCC	0.247 ± 0.154	0.318 ± 0.239
MLP	Neuromap	DMoN	N/A	None	MCC	0.262 ± 0.159	0.298 ± 0.232
MLP	Neuromap	L2	N/A	None	MCC	-0.003 ± 0.051	0.215 ± 0.161
MLP	SBM_{NN}	None	N/A	None	MCC	0.042 ± 0.168	0.123 ± 0.203
MLP	SBM_{NN}	DMoN	N/A	None	MCC	0.020 ± 0.094	0.046 ± 0.091
MLP	SBM_{NN}	L2	N/A	None	Loss	0.016 ± 0.047	0.182 ± 0.123

F.4.10 COAUTHOR PHYSICS DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 96: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.796 ± 0.058	0.859 ± 0.041
GCN	None	DMoN	\times	None	Loss	0.502 ± 0.072	0.567 ± 0.060
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.167 ± 0.000
GCN	DMoN	None	\times	None	Loss	0.487 ± 0.062	0.552 ± 0.060
GCN	DMoN	DMoN	\times	None	Loss	0.497 ± 0.066	0.564 ± 0.055
GCN	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.167 ± 0.000
GCN	NOCD	None	\times	None	Loss	0.391 ± 0.104	0.434 ± 0.122
GCN	NOCD	DMoN	\times	None	Loss	0.217 ± 0.124	0.275 ± 0.073
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.196 ± 0.166
GCN	Neuromap	None	\times	None	Loss	0.451 ± 0.364	0.524 ± 0.344
GCN	Neuromap	DMoN	\checkmark	MLP	Loss	0.290 ± 0.180	0.336 ± 0.142
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.167 ± 0.000
GCN	SBM_{NN}	None	\checkmark	MLP	Loss	0.397 ± 0.117	0.511 ± 0.126
GCN	SBM_{NN}	DMoN	\checkmark	Transformer	Loss	0.431 ± 0.126	0.539 ± 0.131
GCN	SBM_{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.215 ± 0.156
GraphSAGE	None	None	\times	None	Loss	0.802 ± 0.047	0.864 ± 0.034
GraphSAGE	None	DMoN	\times	None	Loss	0.214 ± 0.085	0.380 ± 0.097
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.196 ± 0.165

Continued on next page

Table 96: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	DMoN	None	✓	MLP	Loss	0.298 ± 0.077	0.422 ± 0.075
GraphSAGE	DMoN	DMoN	✓	MLP	Loss	0.280 ± 0.075	0.386 ± 0.048
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.122 ± 0.032
GraphSAGE	NOCD	None	✗	None	Loss	0.256 ± 0.150	0.349 ± 0.120
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.213 ± 0.060	0.278 ± 0.030
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.229 ± 0.193
GraphSAGE	Neuromap	None	✓	MLP	Loss	0.281 ± 0.147	0.307 ± 0.096
GraphSAGE	Neuromap	DMoN	✓	MLP	Loss	0.222 ± 0.113	0.247 ± 0.098
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.187 ± 0.170
GraphSAGE	SBM _{NN}	None	✓	Transformer	Loss	0.232 ± 0.178	0.335 ± 0.175
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.294 ± 0.154	0.394 ± 0.157
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.181 ± 0.118
Transformer	None	None	N/A	None	Loss	0.023 ± 0.081	0.159 ± 0.124
Transformer	None	DMoN	N/A	None	Loss	0.001 ± 0.002	0.128 ± 0.034
Transformer	None	L2	N/A	None	Loss	0.356 ± 0.119	0.443 ± 0.164
Transformer	DMoN	None	N/A	None	Loss	0.023 ± 0.074	0.199 ± 0.154
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.001	0.191 ± 0.169
Transformer	DMoN	L2	N/A	None	Loss	0.379 ± 0.116	0.490 ± 0.143
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.152 ± 0.032
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.194 ± 0.111
Transformer	NOCD	L2	N/A	None	Loss	0.542 ± 0.124	0.632 ± 0.126
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.187 ± 0.169
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.238 ± 0.187
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.227 ± 0.194
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.325 ± 0.181	0.445 ± 0.200
MLP	None	None	N/A	None	Loss	0.488 ± 0.060	0.604 ± 0.089
MLP	None	DMoN	N/A	None	Loss	0.279 ± 0.141	0.383 ± 0.118
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.240 ± 0.186
MLP	DMoN	None	N/A	None	Loss	0.302 ± 0.114	0.396 ± 0.145
MLP	DMoN	DMoN	N/A	None	Loss	0.293 ± 0.106	0.408 ± 0.163
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.198 ± 0.166
MLP	NOCD	None	N/A	None	Loss	0.058 ± 0.128	0.195 ± 0.102
MLP	NOCD	DMoN	N/A	None	Loss	0.114 ± 0.140	0.258 ± 0.121
MLP	NOCD	L2	N/A	None	Loss	0.622 ± 0.090	0.725 ± 0.064
MLP	Neuromap	None	N/A	None	Loss	0.020 ± 0.028	0.308 ± 0.210
MLP	Neuromap	DMoN	N/A	None	Loss	0.055 ± 0.090	0.223 ± 0.159
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.164 ± 0.124
MLP	SBM _{NN}	None	N/A	None	Loss	0.033 ± 0.077	0.183 ± 0.047
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
MLP	SBM _{NN}	L2	N/A	None	Loss	0.099 ± 0.193	0.243 ± 0.189

Table 97: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	Loss	0.796 ± 0.058	0.859 ± 0.041
GCN	None	DMoN	✗	None	Loss	0.502 ± 0.072	0.567 ± 0.060
GCN	None	L2	✗	None	MCC	0.010 ± 0.199	0.309 ± 0.203
GCN	DMoN	None	✗	None	Loss	0.487 ± 0.062	0.552 ± 0.060
GCN	DMoN	DMoN	✗	None	Loss	0.497 ± 0.066	0.564 ± 0.055
GCN	DMoN	L2	✗	None	MCC	0.168 ± 0.260	0.397 ± 0.250
GCN	NOCD	None	✗	None	Loss	0.391 ± 0.104	0.434 ± 0.122
GCN	NOCD	DMoN	✓	MLP	Loss	0.241 ± 0.108	0.293 ± 0.068
GCN	NOCD	L2	✗	None	MCC	0.176 ± 0.147	0.313 ± 0.187
GCN	Neuromap	None	✗	None	MCC	0.341 ± 0.204	0.413 ± 0.213
GCN	Neuromap	DMoN	✗	None	MCC	0.510 ± 0.141	0.615 ± 0.155
GCN	Neuromap	L2	✗	None	MCC	0.096 ± 0.183	0.308 ± 0.185
GCN	SBM _{NN}	None	✓	MLP	Loss	0.397 ± 0.117	0.511 ± 0.126
GCN	SBM _{NN}	DMoN	✓	Transformer	Loss	0.431 ± 0.126	0.539 ± 0.131
GCN	SBM _{NN}	L2	✗	None	MCC	0.089 ± 0.128	0.283 ± 0.192
GraphSAGE	None	None	✗	None	Loss	0.802 ± 0.047	0.864 ± 0.034
GraphSAGE	None	DMoN	✗	None	MCC	0.757 ± 0.068	0.820 ± 0.064
GraphSAGE	None	L2	✗	None	MCC	-0.001 ± 0.004	0.012 ± 0.039
GraphSAGE	DMoN	None	✗	None	Loss	0.209 ± 0.073	0.350 ± 0.054
GraphSAGE	DMoN	DMoN	✗	None	Loss	0.239 ± 0.119	0.368 ± 0.088
GraphSAGE	DMoN	L2	✗	None	MCC	0.018 ± 0.042	0.082 ± 0.166
GraphSAGE	NOCD	None	✓	MLP	Loss	0.367 ± 0.121	0.454 ± 0.119
GraphSAGE	NOCD	DMoN	✗	None	Loss	0.213 ± 0.060	0.278 ± 0.030
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.026 ± 0.082	0.041 ± 0.129
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.350 ± 0.122	0.354 ± 0.154
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.394 ± 0.129	0.449 ± 0.190
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.187 ± 0.170
GraphSAGE	SBM _{NN}	None	✓	MLP	MCC	0.415 ± 0.246	0.458 ± 0.287
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.294 ± 0.154	0.394 ± 0.157
GraphSAGE	SBM _{NN}	L2	✗	None	MCC	0.016 ± 0.050	0.052 ± 0.165
Transformer	None	None	N/A	None	MCC	0.027 ± 0.047	0.083 ± 0.166
Transformer	None	DMoN	N/A	None	Loss	0.001 ± 0.002	0.128 ± 0.034
Transformer	None	L2	N/A	None	Loss	0.356 ± 0.119	0.443 ± 0.164
Transformer	DMoN	None	N/A	None	MCC	0.043 ± 0.074	0.091 ± 0.172
Transformer	DMoN	DMoN	N/A	None	MCC	0.033 ± 0.085	0.072 ± 0.179
Transformer	DMoN	L2	N/A	None	Loss	0.379 ± 0.116	0.490 ± 0.143
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.152 ± 0.032
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.194 ± 0.111
Transformer	NOCD	L2	N/A	None	Loss	0.542 ± 0.124	0.632 ± 0.126
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.187 ± 0.169
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.238 ± 0.187
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.227 ± 0.194

Continued on next page

Table 97: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.325 ± 0.181	0.445 ± 0.200
MLP	None	None	N/A	None	Loss	0.488 ± 0.060	0.604 ± 0.089
MLP	None	DMoN	N/A	None	Loss	0.279 ± 0.141	0.383 ± 0.118
MLP	None	L2	N/A	None	MCC	0.222 ± 0.110	0.324 ± 0.192
MLP	DMoN	None	N/A	None	Loss	0.302 ± 0.114	0.396 ± 0.145
MLP	DMoN	DMoN	N/A	None	Loss	0.293 ± 0.106	0.408 ± 0.163
MLP	DMoN	L2	N/A	None	MCC	0.182 ± 0.142	0.315 ± 0.192
MLP	NOCD	None	N/A	None	Loss	0.058 ± 0.128	0.195 ± 0.102
MLP	NOCD	DMoN	N/A	None	Loss	0.114 ± 0.140	0.258 ± 0.121
MLP	NOCD	L2	N/A	None	Loss	0.622 ± 0.090	0.725 ± 0.064
MLP	Neuromap	None	N/A	None	MCC	0.149 ± 0.090	0.209 ± 0.125
MLP	Neuromap	DMoN	N/A	None	MCC	0.178 ± 0.123	0.280 ± 0.206
MLP	Neuromap	L2	N/A	None	MCC	-0.011 ± 0.046	0.134 ± 0.050
MLP	SBM_{NN}	None	N/A	None	MCC	0.036 ± 0.062	0.087 ± 0.096
MLP	SBM_{NN}	DMoN	N/A	None	MCC	0.029 ± 0.058	0.071 ± 0.117
MLP	SBM_{NN}	L2	N/A	None	Loss	0.099 ± 0.193	0.243 ± 0.189

Table 98: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	MCC	0.816 ± 0.046	0.870 ± 0.039
GCN	None	DMoN	\times	None	MCC	0.757 ± 0.049	0.827 ± 0.040
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.167 ± 0.000
GCN	DMoN	None	\times	None	MCC	0.830 ± 0.042	0.882 ± 0.030
GCN	DMoN	DMoN	\checkmark	Transformer	MCC	0.817 ± 0.054	0.873 ± 0.038
GCN	DMoN	L2	\times	None	MCC	0.168 ± 0.260	0.397 ± 0.250
GCN	NOCD	None	\checkmark	MLP	MCC	0.673 ± 0.193	0.726 ± 0.222
GCN	NOCD	DMoN	\checkmark	Transformer	MCC	0.616 ± 0.170	0.679 ± 0.199
GCN	NOCD	L2	\times	None	MCC	0.176 ± 0.147	0.313 ± 0.187
GCN	Neuromap	None	\times	None	Loss	0.451 ± 0.364	0.524 ± 0.344
GCN	Neuromap	DMoN	\times	None	MCC	0.510 ± 0.141	0.615 ± 0.155
GCN	Neuromap	L2	\times	None	MCC	0.096 ± 0.183	0.308 ± 0.185
GCN	SBM_{NN}	None	\times	None	MCC	0.641 ± 0.162	0.698 ± 0.203
GCN	SBM_{NN}	DMoN	\checkmark	Transformer	MCC	0.698 ± 0.113	0.768 ± 0.130
GCN	SBM_{NN}	L2	\checkmark	MLP	MCC	0.103 ± 0.152	0.265 ± 0.174
GraphSAGE	None	None	\times	None	MCC	0.795 ± 0.061	0.856 ± 0.048
GraphSAGE	None	DMoN	\times	None	MCC	0.757 ± 0.068	0.820 ± 0.064
GraphSAGE	None	L2	\times	None	MCC	-0.001 ± 0.004	0.012 ± 0.039
GraphSAGE	DMoN	None	\times	None	MCC	0.792 ± 0.047	0.855 ± 0.034
GraphSAGE	DMoN	DMoN	\checkmark	Transformer	MCC	0.791 ± 0.040	0.853 ± 0.026
GraphSAGE	DMoN	L2	\times	None	MCC	0.018 ± 0.042	0.082 ± 0.166

Continued on next page

Table 98: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✗	None	MCC	0.635 ± 0.156	0.690 ± 0.190
GraphSAGE	NOCD	DMoN	✓	Transformer	MCC	0.672 ± 0.171	0.743 ± 0.180
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.026 ± 0.082	0.041 ± 0.129
GraphSAGE	Neuromap	None	✗	None	MCC	0.387 ± 0.120	0.501 ± 0.171
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.394 ± 0.129	0.449 ± 0.190
GraphSAGE	Neuromap	L2	✓	MLP	MCC	0.005 ± 0.017	0.014 ± 0.045
GraphSAGE	SBM _{NN}	None	✗	None	MCC	0.462 ± 0.174	0.504 ± 0.235
GraphSAGE	SBM _{NN}	DMoN	✗	None	MCC	0.474 ± 0.208	0.546 ± 0.242
GraphSAGE	SBM _{NN}	L2	✗	None	MCC	0.016 ± 0.050	0.052 ± 0.165
Transformer	None	None	N/A	None	Loss	0.023 ± 0.081	0.159 ± 0.124
Transformer	None	DMoN	N/A	None	Loss	0.001 ± 0.002	0.128 ± 0.034
Transformer	None	L2	N/A	None	MCC	0.490 ± 0.107	0.592 ± 0.134
Transformer	DMoN	None	N/A	None	MCC	0.043 ± 0.074	0.091 ± 0.172
Transformer	DMoN	DMoN	N/A	None	MCC	0.033 ± 0.085	0.072 ± 0.179
Transformer	DMoN	L2	N/A	None	MCC	0.497 ± 0.095	0.614 ± 0.115
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.152 ± 0.032
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.194 ± 0.111
Transformer	NOCD	L2	N/A	None	MCC	0.495 ± 0.265	0.586 ± 0.310
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.187 ± 0.169
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.238 ± 0.187
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.227 ± 0.194
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.167 ± 0.000
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.325 ± 0.181	0.445 ± 0.200
MLP	None	None	N/A	None	MCC	0.493 ± 0.120	0.613 ± 0.146
MLP	None	DMoN	N/A	None	MCC	0.446 ± 0.126	0.551 ± 0.144
MLP	None	L2	N/A	None	MCC	0.222 ± 0.110	0.324 ± 0.192
MLP	DMoN	None	N/A	None	MCC	0.533 ± 0.093	0.641 ± 0.089
MLP	DMoN	DMoN	N/A	None	MCC	0.548 ± 0.097	0.655 ± 0.087
MLP	DMoN	L2	N/A	None	MCC	0.182 ± 0.142	0.315 ± 0.192
MLP	NOCD	None	N/A	None	MCC	0.197 ± 0.185	0.240 ± 0.230
MLP	NOCD	DMoN	N/A	None	MCC	0.223 ± 0.185	0.295 ± 0.236
MLP	NOCD	L2	N/A	None	Loss	0.622 ± 0.090	0.725 ± 0.064
MLP	Neuromap	None	N/A	None	MCC	0.149 ± 0.090	0.209 ± 0.125
MLP	Neuromap	DMoN	N/A	None	MCC	0.178 ± 0.123	0.280 ± 0.206
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.164 ± 0.124
MLP	SBM _{NN}	None	N/A	None	Loss	0.033 ± 0.077	0.183 ± 0.047
MLP	SBM _{NN}	DMoN	N/A	None	MCC	0.029 ± 0.058	0.071 ± 0.117
MLP	SBM _{NN}	L2	N/A	None	Loss	0.099 ± 0.193	0.243 ± 0.189

F.4.11 ROMAN-EMPIRE DATASET, DEFAULT SPLIT WITH DEFAULT TRAIN NODES PER CLASS AND DEFAULT VALIDATION NODES.

Table 99: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.087 ± 0.005	0.191 ± 0.006
GCN	None	DMoN	\times	None	Loss	0.087 ± 0.005	0.192 ± 0.005
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	DMoN	None	\times	None	Loss	0.089 ± 0.005	0.193 ± 0.004
GCN	DMoN	DMoN	\times	None	Loss	0.088 ± 0.007	0.193 ± 0.007
GCN	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	NOCD	None	\times	None	Loss	0.056 ± 0.018	0.163 ± 0.016
GCN	NOCD	DMoN	\times	None	Loss	0.065 ± 0.008	0.173 ± 0.009
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	Neuromap	None	\checkmark	MLP	Loss	0.000 ± 0.000	0.139 ± 0.001
GCN	Neuromap	DMoN	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	SBM _{NN}	None	\times	None	Loss	0.088 ± 0.018	0.196 ± 0.017
GCN	SBM _{NN}	DMoN	\checkmark	MLP	Loss	0.074 ± 0.032	0.183 ± 0.021
GCN	SBM _{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GraphSAGE	None	None	\times	None	Loss	0.285 ± 0.150	0.355 ± 0.113
GraphSAGE	None	DMoN	\times	None	Loss	0.143 ± 0.184	0.248 ± 0.140
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.139 ± 0.000
GraphSAGE	DMoN	None	\checkmark	MLP	Loss	0.318 ± 0.034	0.381 ± 0.027
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.323 ± 0.019	0.385 ± 0.014
GraphSAGE	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GraphSAGE	NOCD	None	\checkmark	MLP	Loss	0.280 ± 0.037	0.317 ± 0.050
GraphSAGE	NOCD	DMoN	\checkmark	MLP	Loss	0.248 ± 0.065	0.278 ± 0.061
GraphSAGE	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.139 ± 0.001
GraphSAGE	Neuromap	None	\times	None	Loss	0.000 ± 0.000	0.130 ± 0.014
GraphSAGE	Neuromap	DMoN	\times	None	Loss	0.000 ± 0.000	0.136 ± 0.009
GraphSAGE	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.139 ± 0.001
GraphSAGE	SBM _{NN}	None	\times	None	Loss	0.167 ± 0.123	0.248 ± 0.084
GraphSAGE	SBM _{NN}	DMoN	\times	None	Loss	0.201 ± 0.112	0.274 ± 0.081
GraphSAGE	SBM _{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.139 ± 0.001
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.126 ± 0.031
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.132 ± 0.015
Transformer	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.066 ± 0.058
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.098 ± 0.040
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.109 ± 0.046

Continued on next page

Table 99: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.100 ± 0.050
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.086 ± 0.048
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
MLP	None	None	N/A	None	Loss	0.056 ± 0.090	0.155 ± 0.027
MLP	None	DMoN	N/A	None	Loss	0.081 ± 0.102	0.183 ± 0.062
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	DMoN	None	N/A	None	Loss	0.061 ± 0.099	0.173 ± 0.065
MLP	DMoN	DMoN	N/A	None	Loss	0.120 ± 0.112	0.203 ± 0.071
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	NOCD	None	N/A	None	Loss	0.021 ± 0.067	0.141 ± 0.029
MLP	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.111 ± 0.039
MLP	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.059 ± 0.041
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.074 ± 0.053
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.129 ± 0.032
MLP	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.098 ± 0.043
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.116 ± 0.031
MLP	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.126 ± 0.031

Table 100: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.087 ± 0.005	0.191 ± 0.006
GCN	None	DMoN	\times	None	Loss	0.087 ± 0.005	0.192 ± 0.005
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	DMoN	None	\checkmark	MLP	Loss	0.089 ± 0.005	0.198 ± 0.004
GCN	DMoN	DMoN	\checkmark	MLP	Loss	0.089 ± 0.002	0.199 ± 0.002
GCN	DMoN	L2	\checkmark	MLP	MCC	0.006 ± 0.066	0.097 ± 0.063
GCN	NOCD	None	\times	None	Loss	0.056 ± 0.018	0.163 ± 0.016
GCN	NOCD	DMoN	\times	None	Loss	0.065 ± 0.008	0.173 ± 0.009
GCN	NOCD	L2	\checkmark	Transformer	MCC	0.022 ± 0.035	0.108 ± 0.078
GCN	Neuromap	None	\checkmark	MLP	MCC	0.032 ± 0.052	0.122 ± 0.060
GCN	Neuromap	DMoN	\checkmark	Transformer	MCC	0.010 ± 0.030	0.113 ± 0.053
GCN	Neuromap	L2	\checkmark	Transformer	MCC	0.020 ± 0.053	0.128 ± 0.043
GCN	SBM_{NN}	None	\times	None	Loss	0.088 ± 0.018	0.196 ± 0.017
GCN	SBM_{NN}	DMoN	\times	None	Loss	0.073 ± 0.032	0.187 ± 0.023
GCN	SBM_{NN}	L2	\checkmark	Transformer	MCC	0.036 ± 0.028	0.112 ± 0.079
GraphSAGE	None	None	\times	None	Loss	0.285 ± 0.150	0.355 ± 0.113
GraphSAGE	None	DMoN	\times	None	Loss	0.143 ± 0.184	0.248 ± 0.140
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.139 ± 0.000
GraphSAGE	DMoN	None	\checkmark	MLP	Loss	0.318 ± 0.034	0.381 ± 0.027
GraphSAGE	DMoN	DMoN	\checkmark	MLP	Loss	0.323 ± 0.019	0.385 ± 0.014
GraphSAGE	DMoN	L2	\checkmark	MLP	MCC	-0.000 ± 0.000	0.010 ± 0.018

Continued on next page

Table 100: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✓	MLP	Loss	0.280 ± 0.037	0.317 ± 0.050
GraphSAGE	NOCD	DMoN	✓	MLP	Loss	0.248 ± 0.065	0.278 ± 0.061
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.001 ± 0.006	0.015 ± 0.043
GraphSAGE	Neuromap	None	✗	None	Loss	0.000 ± 0.000	0.130 ± 0.014
GraphSAGE	Neuromap	DMoN	✗	None	MCC	0.013 ± 0.041	0.023 ± 0.071
GraphSAGE	Neuromap	L2	✓	Transformer	MCC	0.000 ± 0.000	0.002 ± 0.006
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.167 ± 0.123	0.248 ± 0.084
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.201 ± 0.112	0.274 ± 0.081
GraphSAGE	SBM _{NN}	L2	✓	MLP	MCC	0.007 ± 0.021	0.022 ± 0.041
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	None	L2	N/A	None	MCC	0.012 ± 0.038	0.022 ± 0.068
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	DMoN	L2	N/A	None	MCC	0.000 ± 0.019	0.020 ± 0.043
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.126 ± 0.031
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.132 ± 0.015
Transformer	NOCD	L2	N/A	None	MCC	0.013 ± 0.039	0.030 ± 0.065
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.066 ± 0.058
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.098 ± 0.040
Transformer	Neuromap	L2	N/A	None	MCC	0.008 ± 0.014	0.042 ± 0.048
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.100 ± 0.050
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.086 ± 0.048
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
MLP	None	None	N/A	None	MCC	0.226 ± 0.037	0.251 ± 0.038
MLP	None	DMoN	N/A	None	MCC	0.216 ± 0.026	0.245 ± 0.025
MLP	None	L2	N/A	None	MCC	0.004 ± 0.029	0.049 ± 0.069
MLP	DMoN	None	N/A	None	MCC	0.239 ± 0.046	0.261 ± 0.051
MLP	DMoN	DMoN	N/A	None	MCC	0.245 ± 0.043	0.262 ± 0.051
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	NOCD	None	N/A	None	MCC	0.055 ± 0.090	0.065 ± 0.105
MLP	NOCD	DMoN	N/A	None	MCC	0.062 ± 0.095	0.083 ± 0.111
MLP	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.059 ± 0.041
MLP	Neuromap	DMoN	N/A	None	MCC	0.001 ± 0.002	0.037 ± 0.057
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.129 ± 0.032
MLP	SBM _{NN}	None	N/A	None	MCC	0.002 ± 0.007	0.019 ± 0.044
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.116 ± 0.031
MLP	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.126 ± 0.031

Table 101: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	Loss	0.087 ± 0.005	0.191 ± 0.006
GCN	None	DMoN	✗	None	Loss	0.087 ± 0.005	0.192 ± 0.005
GCN	None	L2	✗	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	DMoN	None	✗	None	Loss	0.089 ± 0.005	0.193 ± 0.004
GCN	DMoN	DMoN	✓	MLP	Loss	0.089 ± 0.002	0.199 ± 0.002
GCN	DMoN	L2	✓	MLP	MCC	0.006 ± 0.066	0.097 ± 0.063
GCN	NOCD	None	✗	None	Loss	0.056 ± 0.018	0.163 ± 0.016
GCN	NOCD	DMoN	✗	None	Loss	0.065 ± 0.008	0.173 ± 0.009
GCN	NOCD	L2	✓	Transformer	MCC	0.022 ± 0.035	0.108 ± 0.078
GCN	Neuromap	None	✓	MLP	MCC	0.032 ± 0.052	0.122 ± 0.060
GCN	Neuromap	DMoN	✓	Transformer	MCC	0.010 ± 0.030	0.113 ± 0.053
GCN	Neuromap	L2	✓	Transformer	MCC	0.020 ± 0.053	0.128 ± 0.043
GCN	SBM _{NN}	None	✗	None	Loss	0.088 ± 0.018	0.196 ± 0.017
GCN	SBM _{NN}	DMoN	✓	MLP	Loss	0.074 ± 0.032	0.183 ± 0.021
GCN	SBM _{NN}	L2	✓	Transformer	MCC	0.036 ± 0.028	0.112 ± 0.079
GraphSAGE	None	None	✗	None	Loss	0.285 ± 0.150	0.355 ± 0.113
GraphSAGE	None	DMoN	✗	None	Loss	0.143 ± 0.184	0.248 ± 0.140
GraphSAGE	None	L2	✗	None	Loss	0.000 ± 0.000	0.139 ± 0.000
GraphSAGE	DMoN	None	✓	MLP	Loss	0.318 ± 0.034	0.381 ± 0.027
GraphSAGE	DMoN	DMoN	✓	MLP	Loss	0.323 ± 0.019	0.385 ± 0.014
GraphSAGE	DMoN	L2	✓	Transformer	MCC	-0.000 ± 0.007	0.021 ± 0.025
GraphSAGE	NOCD	None	✓	MLP	Loss	0.280 ± 0.037	0.317 ± 0.050
GraphSAGE	NOCD	DMoN	✓	MLP	Loss	0.248 ± 0.065	0.278 ± 0.061
GraphSAGE	NOCD	L2	✓	MLP	MCC	0.001 ± 0.006	0.015 ± 0.043
GraphSAGE	Neuromap	None	✗	None	Loss	0.000 ± 0.000	0.130 ± 0.014
GraphSAGE	Neuromap	DMoN	✗	None	MCC	0.013 ± 0.041	0.023 ± 0.071
GraphSAGE	Neuromap	L2	✓	Transformer	MCC	0.000 ± 0.000	0.002 ± 0.006
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.167 ± 0.123	0.248 ± 0.084
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.201 ± 0.112	0.274 ± 0.081
GraphSAGE	SBM _{NN}	L2	✓	MLP	MCC	0.007 ± 0.021	0.022 ± 0.041
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	None	L2	N/A	None	MCC	0.012 ± 0.038	0.022 ± 0.068
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
Transformer	DMoN	L2	N/A	None	MCC	0.000 ± 0.019	0.020 ± 0.043
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.126 ± 0.031
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.132 ± 0.015
Transformer	NOCD	L2	N/A	None	MCC	0.013 ± 0.039	0.030 ± 0.065
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.066 ± 0.058
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.098 ± 0.040
Transformer	Neuromap	L2	N/A	None	MCC	0.008 ± 0.014	0.042 ± 0.048

Continued on next page

Table 101: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.100 ± 0.050
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.086 ± 0.048
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
MLP	None	None	N/A	None	MCC	0.226 ± 0.037	0.251 ± 0.038
MLP	None	DMoN	N/A	None	MCC	0.216 ± 0.026	0.245 ± 0.025
MLP	None	L2	N/A	None	MCC	0.004 ± 0.029	0.049 ± 0.069
MLP	DMoN	None	N/A	None	MCC	0.239 ± 0.046	0.261 ± 0.051
MLP	DMoN	DMoN	N/A	None	MCC	0.245 ± 0.043	0.262 ± 0.051
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	NOCD	None	N/A	None	MCC	0.055 ± 0.090	0.065 ± 0.105
MLP	NOCD	DMoN	N/A	None	MCC	0.062 ± 0.095	0.083 ± 0.111
MLP	NOCD	L2	N/A	None	MCC	-0.000 ± 0.071	0.059 ± 0.077
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.059 ± 0.041
MLP	Neuromap	DMoN	N/A	None	MCC	0.001 ± 0.002	0.037 ± 0.057
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.129 ± 0.032
MLP	SBM_{NN}	None	N/A	None	MCC	0.002 ± 0.007	0.019 ± 0.044
MLP	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.116 ± 0.031
MLP	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.126 ± 0.031

F.4.12 ROMAN-EMPIRE DATASET, SPARSE SPLIT WITH 2 TRAIN NODES PER CLASS AND 50 VALIDATION NODES.

Table 102: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training loss.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	\times	None	Loss	0.089 ± 0.003	0.195 ± 0.005
GCN	None	DMoN	\times	None	Loss	0.090 ± 0.005	0.196 ± 0.006
GCN	None	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	DMoN	None	\times	None	Loss	0.090 ± 0.005	0.195 ± 0.005
GCN	DMoN	DMoN	\times	None	Loss	0.090 ± 0.006	0.195 ± 0.006
GCN	DMoN	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	NOCD	None	\times	None	Loss	0.057 ± 0.019	0.166 ± 0.017
GCN	NOCD	DMoN	\times	None	Loss	0.049 ± 0.025	0.161 ± 0.018
GCN	NOCD	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	Neuromap	None	\times	None	Loss	0.000 ± 0.000	0.139 ± 0.001
GCN	Neuromap	DMoN	\times	None	Loss	0.000 ± 0.000	0.139 ± 0.001
GCN	Neuromap	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	SBM_{NN}	None	\checkmark	MLP	Loss	0.069 ± 0.023	0.178 ± 0.015
GCN	SBM_{NN}	DMoN	\times	None	Loss	0.072 ± 0.023	0.185 ± 0.018
GCN	SBM_{NN}	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GraphSAGE	None	None	\times	None	Loss	0.141 ± 0.183	0.247 ± 0.139
GraphSAGE	None	DMoN	\times	None	Loss	0.214 ± 0.184	0.301 ± 0.139
GraphSAGE	None	L2	\times	None	Loss	0.000 ± 0.000	0.140 ± 0.000

Continued on next page

Table 102: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	DMoN	None	✓	MLP	Loss	0.324 ± 0.035	0.385 ± 0.026
GraphSAGE	DMoN	DMoN	✓	MLP	Loss	0.327 ± 0.019	0.388 ± 0.014
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.139 ± 0.000
GraphSAGE	NOCD	None	✓	MLP	Loss	0.272 ± 0.077	0.316 ± 0.073
GraphSAGE	NOCD	DMoN	✓	MLP	Loss	0.280 ± 0.050	0.312 ± 0.065
GraphSAGE	NOCD	L2	✗	None	Loss	0.000 ± 0.000	0.139 ± 0.001
GraphSAGE	Neuromap	None	✗	None	Loss	0.000 ± 0.000	0.137 ± 0.009
GraphSAGE	Neuromap	DMoN	✓	MLP	Loss	0.000 ± 0.000	0.129 ± 0.017
GraphSAGE	Neuromap	L2	✗	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.183 ± 0.101	0.256 ± 0.066
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.183 ± 0.130	0.269 ± 0.093
GraphSAGE	SBM _{NN}	L2	✗	None	Loss	0.000 ± 0.000	0.139 ± 0.001
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
Transformer	DMoN	DMoN	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
Transformer	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.117 ± 0.040
Transformer	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.112 ± 0.040
Transformer	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.108 ± 0.046
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.095 ± 0.031
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.128 ± 0.027
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.116 ± 0.031
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.101 ± 0.043
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
MLP	None	None	N/A	None	Loss	0.040 ± 0.086	0.166 ± 0.060
MLP	None	DMoN	N/A	None	Loss	0.109 ± 0.098	0.196 ± 0.058
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	DMoN	None	N/A	None	Loss	0.112 ± 0.081	0.202 ± 0.049
MLP	DMoN	DMoN	N/A	None	Loss	0.071 ± 0.092	0.186 ± 0.061
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.120 ± 0.041
MLP	NOCD	DMoN	N/A	None	Loss	0.000 ± 0.000	0.123 ± 0.031
MLP	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.098 ± 0.046
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.071 ± 0.046
MLP	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.128 ± 0.026
MLP	SBM _{NN}	None	N/A	None	Loss	0.018 ± 0.057	0.109 ± 0.060
MLP	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.114 ± 0.040
MLP	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.129 ± 0.031

Table 103: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on training set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	Loss	0.089 ± 0.003	0.195 ± 0.005
GCN	None	DMoN	✗	None	Loss	0.090 ± 0.005	0.196 ± 0.006
GCN	None	L2	✗	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	DMoN	None	✓	MLP	Loss	0.089 ± 0.005	0.198 ± 0.004
GCN	DMoN	DMoN	✓	MLP	Loss	0.089 ± 0.003	0.198 ± 0.002
GCN	DMoN	L2	✓	Transformer	MCC	0.033 ± 0.074	0.149 ± 0.047
GCN	NOCD	None	✗	None	Loss	0.057 ± 0.019	0.166 ± 0.017
GCN	NOCD	DMoN	✓	Transformer	Loss	0.050 ± 0.038	0.164 ± 0.026
GCN	NOCD	L2	✓	Transformer	MCC	0.002 ± 0.046	0.094 ± 0.070
GCN	Neuromap	None	✗	None	MCC	0.018 ± 0.034	0.066 ± 0.086
GCN	Neuromap	DMoN	✓	Transformer	MCC	0.015 ± 0.033	0.076 ± 0.063
GCN	Neuromap	L2	✓	MLP	MCC	0.022 ± 0.066	0.123 ± 0.040
GCN	SBM _{NN}	None	✓	MLP	Loss	0.069 ± 0.023	0.178 ± 0.015
GCN	SBM _{NN}	DMoN	✓	MLP	Loss	0.072 ± 0.035	0.183 ± 0.028
GCN	SBM _{NN}	L2	✓	MLP	MCC	0.030 ± 0.045	0.131 ± 0.058
GraphSAGE	None	None	✗	None	Loss	0.141 ± 0.183	0.247 ± 0.139
GraphSAGE	None	DMoN	✗	None	Loss	0.214 ± 0.184	0.301 ± 0.139
GraphSAGE	None	L2	✗	None	MCC	0.001 ± 0.005	0.023 ± 0.045
GraphSAGE	DMoN	None	✓	MLP	Loss	0.324 ± 0.035	0.385 ± 0.026
GraphSAGE	DMoN	DMoN	✓	MLP	Loss	0.327 ± 0.019	0.388 ± 0.014
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.139 ± 0.000
GraphSAGE	NOCD	None	✓	MLP	Loss	0.272 ± 0.077	0.316 ± 0.073
GraphSAGE	NOCD	DMoN	✓	MLP	Loss	0.280 ± 0.050	0.312 ± 0.065
GraphSAGE	NOCD	L2	✗	None	MCC	0.001 ± 0.002	0.011 ± 0.035
GraphSAGE	Neuromap	None	✗	None	Loss	0.000 ± 0.000	0.137 ± 0.009
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.000 ± 0.001	0.011 ± 0.035
GraphSAGE	Neuromap	L2	✓	Transformer	MCC	0.004 ± 0.019	0.026 ± 0.044
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.183 ± 0.101	0.256 ± 0.066
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.183 ± 0.130	0.269 ± 0.093
GraphSAGE	SBM _{NN}	L2	✓	Transformer	MCC	0.000 ± 0.001	0.003 ± 0.011
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	None	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
Transformer	DMoN	DMoN	N/A	None	MCC	0.004 ± 0.014	0.017 ± 0.053
Transformer	DMoN	L2	N/A	None	MCC	0.001 ± 0.002	0.028 ± 0.059
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.117 ± 0.040
Transformer	NOCD	DMoN	N/A	None	MCC	0.004 ± 0.012	0.013 ± 0.041
Transformer	NOCD	L2	N/A	None	MCC	0.000 ± 0.000	0.025 ± 0.053
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.108 ± 0.046
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.095 ± 0.031
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.128 ± 0.027

Continued on next page

Table 103: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
Transformer	SBM_{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.116 ± 0.031
Transformer	SBM_{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.101 ± 0.043
Transformer	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
MLP	None	None	N/A	None	MCC	0.207 ± 0.004	0.233 ± 0.011
MLP	None	DMoN	N/A	None	MCC	0.222 ± 0.034	0.240 ± 0.035
MLP	None	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	DMoN	None	N/A	None	MCC	0.225 ± 0.039	0.251 ± 0.039
MLP	DMoN	DMoN	N/A	None	MCC	0.234 ± 0.043	0.261 ± 0.041
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	NOCD	None	N/A	None	MCC	0.039 ± 0.082	0.046 ± 0.096
MLP	NOCD	DMoN	N/A	None	MCC	0.042 ± 0.088	0.062 ± 0.103
MLP	NOCD	L2	N/A	None	MCC	0.000 ± 0.001	0.025 ± 0.053
MLP	Neuromap	None	N/A	None	MCC	0.001 ± 0.002	0.020 ± 0.042
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.071 ± 0.046
MLP	Neuromap	L2	N/A	None	MCC	0.005 ± 0.012	0.039 ± 0.041
MLP	SBM_{NN}	None	N/A	None	Loss	0.018 ± 0.057	0.109 ± 0.060
MLP	SBM_{NN}	DMoN	N/A	None	MCC	0.021 ± 0.064	0.035 ± 0.077
MLP	SBM_{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.129 ± 0.031

Table 104: Comparing graph clustering and regularization objectives with an ablation study of regularization. The ablation compares clustering with clustering with regularization, clustering without regularization, and regularization without clustering for each neural network. Model selection based on validation set MCC.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GCN	None	None	✗	None	Loss	0.089 ± 0.003	0.195 ± 0.005
GCN	None	DMoN	✗	None	Loss	0.090 ± 0.005	0.196 ± 0.006
GCN	None	L2	✗	None	Loss	0.000 ± 0.000	0.140 ± 0.000
GCN	DMoN	None	✗	None	Loss	0.090 ± 0.005	0.195 ± 0.005
GCN	DMoN	DMoN	✗	None	Loss	0.090 ± 0.006	0.195 ± 0.006
GCN	DMoN	L2	✓	Transformer	MCC	0.033 ± 0.074	0.149 ± 0.047
GCN	NOCD	None	✗	None	Loss	0.057 ± 0.019	0.166 ± 0.017
GCN	NOCD	DMoN	✗	None	Loss	0.049 ± 0.025	0.161 ± 0.018
GCN	NOCD	L2	✗	None	MCC	0.003 ± 0.035	0.094 ± 0.068
GCN	Neuromap	None	✗	None	MCC	0.018 ± 0.034	0.066 ± 0.086
GCN	Neuromap	DMoN	✓	Transformer	MCC	0.015 ± 0.033	0.076 ± 0.063
GCN	Neuromap	L2	✓	MLP	MCC	0.022 ± 0.066	0.123 ± 0.040
GCN	SBM_{NN}	None	✓	MLP	Loss	0.069 ± 0.023	0.178 ± 0.015
GCN	SBM_{NN}	DMoN	✓	MLP	Loss	0.072 ± 0.035	0.183 ± 0.028
GCN	SBM_{NN}	L2	✓	MLP	MCC	0.030 ± 0.045	0.131 ± 0.058
GraphSAGE	None	None	✗	None	Loss	0.141 ± 0.183	0.247 ± 0.139
GraphSAGE	None	DMoN	✗	None	Loss	0.214 ± 0.184	0.301 ± 0.139
GraphSAGE	None	L2	✗	None	MCC	0.001 ± 0.005	0.023 ± 0.045
GraphSAGE	DMoN	None	✓	MLP	Loss	0.324 ± 0.035	0.385 ± 0.026
GraphSAGE	DMoN	DMoN	✓	MLP	Loss	0.327 ± 0.019	0.388 ± 0.014
GraphSAGE	DMoN	L2	✗	None	Loss	0.000 ± 0.000	0.139 ± 0.000

Continued on next page

Table 104: Comparing graph clustering and regularization objectives for each neural network.

Model	f	$L_{\text{regularization}}$	L_V	$\text{NN}_{S \rightarrow X}$	ES	MCC	Accuracy
GraphSAGE	NOCD	None	✓	MLP	Loss	0.272 ± 0.077	0.316 ± 0.073
GraphSAGE	NOCD	DMoN	✓	MLP	Loss	0.280 ± 0.050	0.312 ± 0.065
GraphSAGE	NOCD	L2	✗	None	MCC	0.001 ± 0.002	0.011 ± 0.035
GraphSAGE	Neuromap	None	✓	MLP	MCC	0.000 ± 0.001	0.012 ± 0.034
GraphSAGE	Neuromap	DMoN	✓	MLP	MCC	0.000 ± 0.001	0.011 ± 0.035
GraphSAGE	Neuromap	L2	✓	Transformer	MCC	0.004 ± 0.019	0.026 ± 0.044
GraphSAGE	SBM _{NN}	None	✗	None	Loss	0.183 ± 0.101	0.256 ± 0.066
GraphSAGE	SBM _{NN}	DMoN	✗	None	Loss	0.183 ± 0.130	0.269 ± 0.093
GraphSAGE	SBM _{NN}	L2	✓	Transformer	MCC	0.000 ± 0.001	0.003 ± 0.011
Transformer	None	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	None	DMoN	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.000
Transformer	None	L2	N/A	None	MCC	-0.000 ± 0.026	0.022 ± 0.050
Transformer	DMoN	None	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
Transformer	DMoN	DMoN	N/A	None	MCC	0.004 ± 0.014	0.017 ± 0.053
Transformer	DMoN	L2	N/A	None	MCC	0.001 ± 0.002	0.028 ± 0.059
Transformer	NOCD	None	N/A	None	Loss	0.000 ± 0.000	0.117 ± 0.040
Transformer	NOCD	DMoN	N/A	None	MCC	0.004 ± 0.012	0.013 ± 0.041
Transformer	NOCD	L2	N/A	None	MCC	0.000 ± 0.000	0.025 ± 0.053
Transformer	Neuromap	None	N/A	None	Loss	0.000 ± 0.000	0.108 ± 0.046
Transformer	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.095 ± 0.031
Transformer	Neuromap	L2	N/A	None	Loss	0.000 ± 0.000	0.128 ± 0.027
Transformer	SBM _{NN}	None	N/A	None	Loss	0.000 ± 0.000	0.116 ± 0.031
Transformer	SBM _{NN}	DMoN	N/A	None	Loss	0.000 ± 0.000	0.101 ± 0.043
Transformer	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.139 ± 0.001
MLP	None	None	N/A	None	MCC	0.207 ± 0.004	0.233 ± 0.011
MLP	None	DMoN	N/A	None	MCC	0.222 ± 0.034	0.240 ± 0.035
MLP	None	L2	N/A	None	MCC	-0.001 ± 0.015	0.050 ± 0.063
MLP	DMoN	None	N/A	None	MCC	0.225 ± 0.039	0.251 ± 0.039
MLP	DMoN	DMoN	N/A	None	MCC	0.234 ± 0.043	0.261 ± 0.041
MLP	DMoN	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	NOCD	None	N/A	None	MCC	0.039 ± 0.082	0.046 ± 0.096
MLP	NOCD	DMoN	N/A	None	MCC	0.042 ± 0.088	0.062 ± 0.103
MLP	NOCD	L2	N/A	None	Loss	0.000 ± 0.000	0.140 ± 0.000
MLP	Neuromap	None	N/A	None	MCC	0.001 ± 0.002	0.020 ± 0.042
MLP	Neuromap	DMoN	N/A	None	Loss	0.000 ± 0.000	0.071 ± 0.046
MLP	Neuromap	L2	N/A	None	MCC	0.005 ± 0.012	0.039 ± 0.041
MLP	SBM _{NN}	None	N/A	None	Loss	0.018 ± 0.057	0.109 ± 0.060
MLP	SBM _{NN}	DMoN	N/A	None	MCC	0.021 ± 0.064	0.035 ± 0.077
MLP	SBM _{NN}	L2	N/A	None	Loss	0.000 ± 0.000	0.129 ± 0.031